

Practical Parallel Computing (実践的並列コンピューティング) 2021 No. 8

Part2: GPU (2)
May 10, 2021

Toshio Endo
School of Computing & GSIC
endo@is.titech.ac.jp

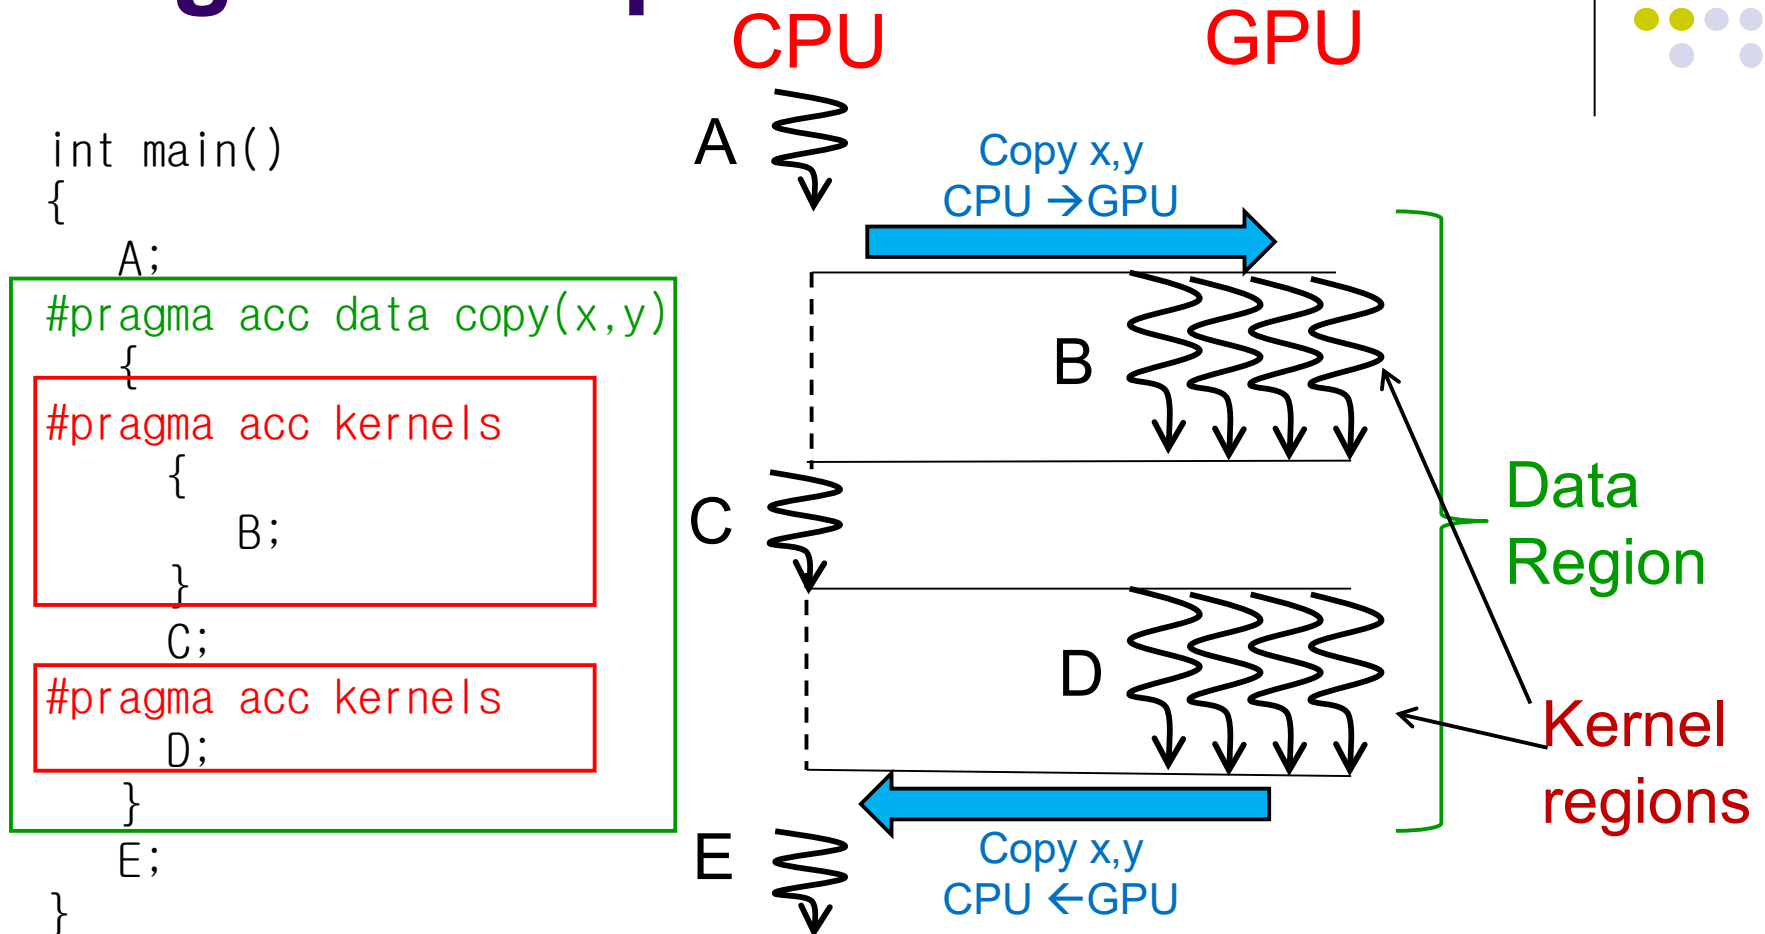
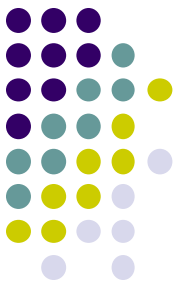




Overview of This Course

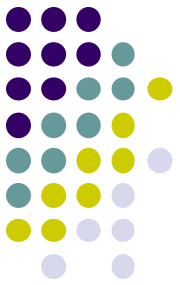
- Part 0: Introduction
 - 2 classes
- Part 1: OpenMP for shared memory programming
 - 4 classes
- Part 2: **GPU** programming
 - 4 classes **← We are here (2/4)**
 - OpenACC (1.5 classes) and CUDA (2.5 classes)
- Part 3: **MPI** for distributed memory programming
 - 3 classes

Data Region and Kernel Region in OpenACC

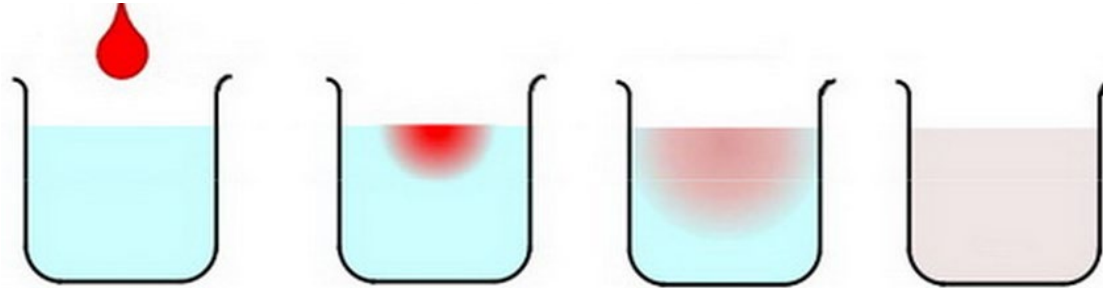


- Data movement occurs at beginning and end of data region
- Data region may contain 1 or more kernel regions

“diffusion” Sample Program related to [G1]



An example of diffusion phenomena:

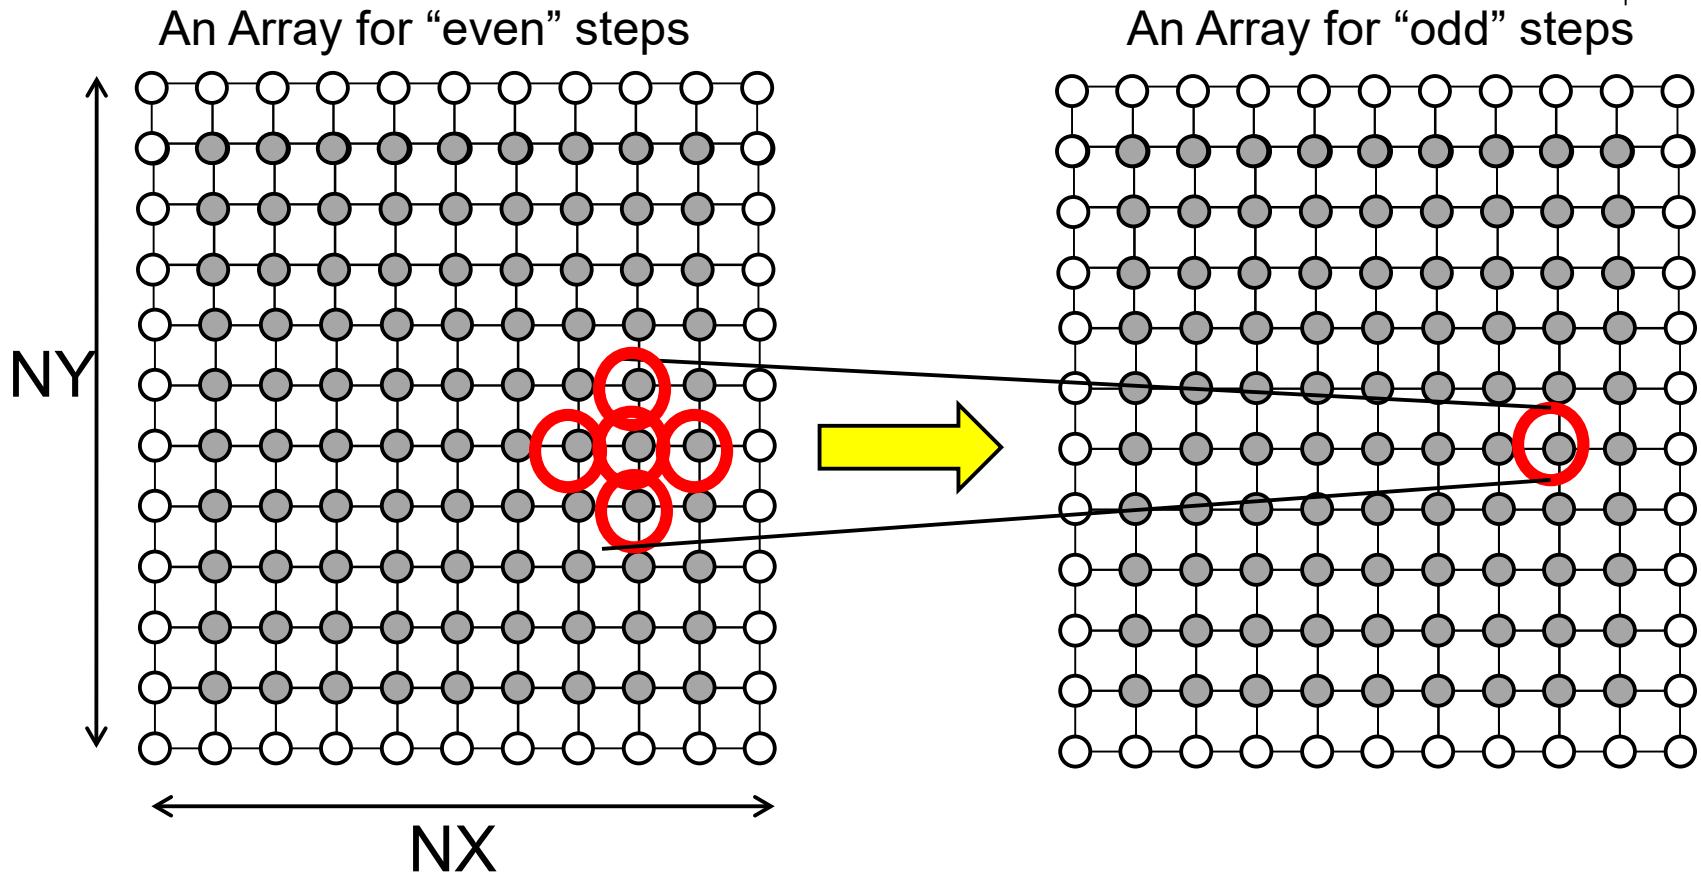
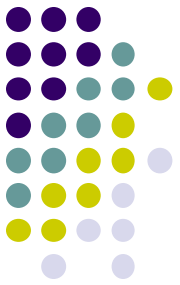


The ink spreads gradually, and finally the density becomes uniform (Figure by Prof. T. Aoki)

Available at [/gs/hs1/tga-ppcomp/21/diffusion/](https://github.com/tga-ppcomp/21/diffusion/)

- Execution: `./diffusion [nt]`
 - nt: Number of time steps

Data Structure in “diffusion”



Consideration of Parallelizing Diffusion with OpenACC related to [G1]



- x, y loops can be parallelized
 - We can use “#pragma acc loop” twice
- t loop cannot be parallelized

[Data transfer from CPU to GPU]

```
for (t = 0; t < nt; t++) {
```

```
    for (y = 1; y < NY-1; y++) {  
        for (x = 1; x < NX-1; x++) {  
            :  
        }  
    }
```

```
}
```

[Data transfer from GPU to CPU]

Kernel region on GPU
Parallel x, y loops

It's better to transfer
data *out of* t-loop

data Clause for Multi-Dimensional arrays



`float A[2000][1000];` → an example of a 2-dimension array

.... `data copy(A)`

→ **OK**, all elements of A are copied

.... `data copy(A[0:2000][0:1000])`

→ **OK**, all elements of A are copied

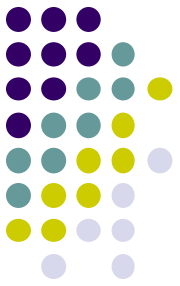
.... `data copy(A[500:600][0:1000])`

→ **OK**, rows[500,1100) are copied

.... `data copy(A[0:2000][300:400])`

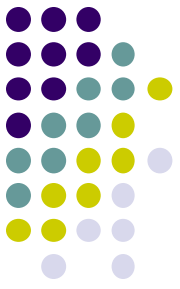
→ **NG** in current OpenACC

✂ Currently, OpenACC does not support non-consecutive transfer



Notes on Assignment [G1] (1)

- You will need compiler options different from the [diffusion](#) directory for OpenACC
- You can use files in [diffusion-acc](#) directory as basis
 - [/gs/hs1/tga-ppcomp/21/diffusion-acc/](#)
 - “Makefile” in this directory supports compiler options for OpenACC
 - Don’t forget “[module load nvhpc](#)” before “make”



Notes on Assignment [G1] (2)

If you see compile error messages like:

```
50, Accelerator restriction: call to 'fflush' with no acc  
routine information  
NVC++/x86-64 Linux 21.2-0: compilation completed with severe  
errors  
Makefile:16: recipe for target 'diffusion.o' failed
```

I/O functions (fflush(0) in this case) cannot be executed inside a kernel region

- Exceptionally, printf(“...”) is ok

In this case, please do either of

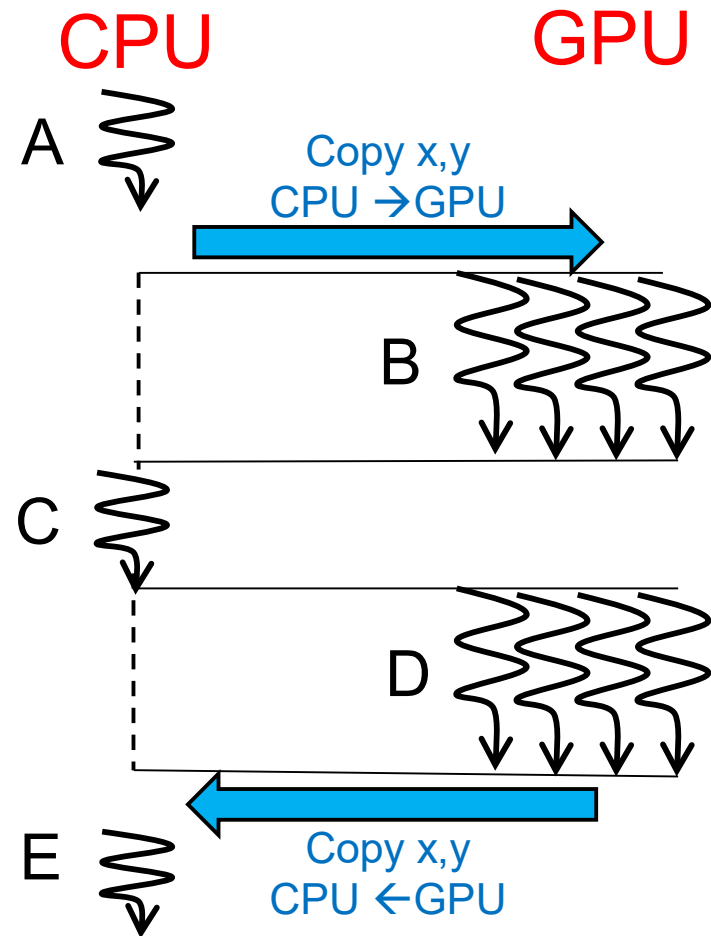
- Delete fflush(0) simply
- Consider to shorten the length of a kernel region

Data Transfer Costs in GPU Programming

Related to [G2]



- In GPU programming, **data transfer costs between CPU and GPU** have impacts on speed
 - Program speed may be slower than expected ☹️



Speed of GPU Programs: case of mm-acc



In mm-acc, speed in Gflops is computed by

$$S = 2mnk / T_{\text{total}}$$

T_{total} includes both computation time and transfer

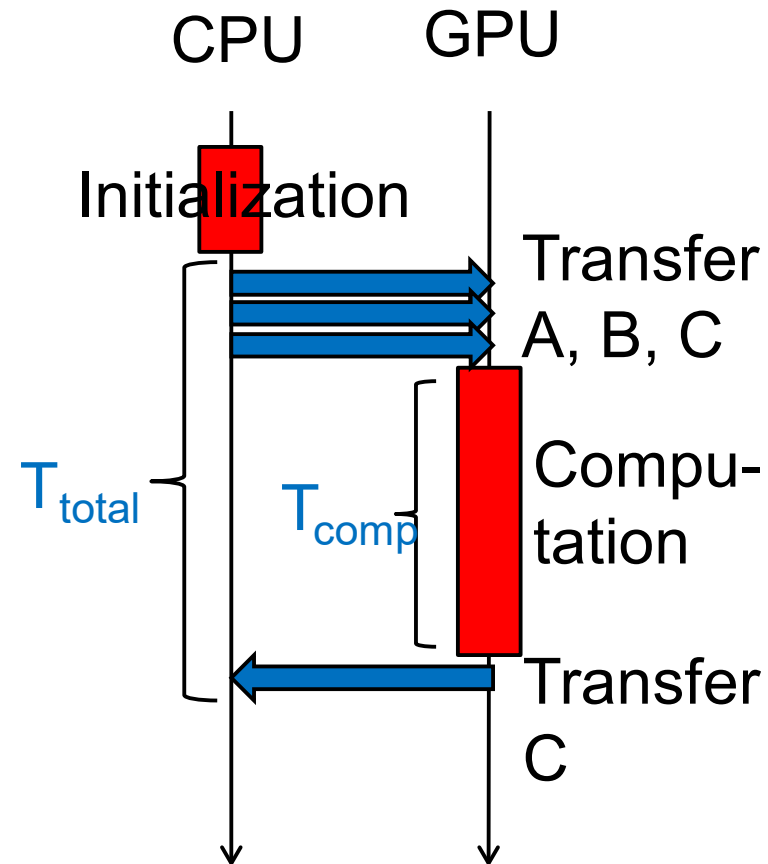
→ S counts slow-down by transfer

To see the effects, let's try another sample

[/gs/hs1/tga-ppcomp/21/mm-meas-acc](https://gs/hs1/tga-ppcomp/21/mm-meas-acc)

which outputs time for

- copyin (transfer A, B, C)
- computation
- copyout (transfer C)



In [G2], please evaluate effects of transfer costs

Another Description Way for Copying Data



How can we measure transfer time?

- With “data” directive, “when we can copy” is restricted
→ We can copy data anytime by “**acc enter data**”, “**acc exit data**” directives

// x,y are on CPU

```
#pragma acc data copy(x,y)
{
    // x,y are on GPU
}
```

// x,y are on CPU



// x,y are on CPU

```
#pragma acc enter data copyin(x,y)
    // x,y are on GPU
#pragma acc exit data copyout(x,y)
```

// x,y are on CPU

See differences between mm-acc/mm.c and mm-meas-acc/mm.c

Discussion on Data Transfer Costs



- Time for data transfer $T_{\text{trans}} \doteq M / B + L$
 - M : Data size in bytes
 - B : “Bandwidth” (speed)
 - L : “Latency” (if M is sufficiently large, we can ignore it)
 - In a P100 GPU,
 - Theoretical computation speed is 5.3TFlops
 - Theoretical bandwidth B is 16GB/s (2G double values per second)
- Transfer of values is much slower than computation ☹️

Discussion on Computation and Transfer Costs

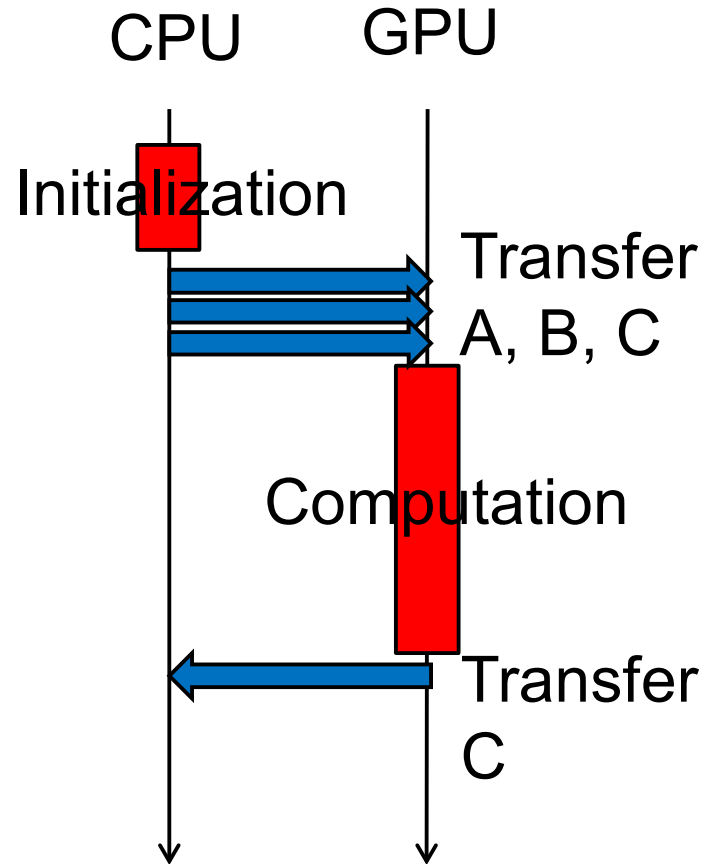


In mm-acc,

- Computation amount: $O(mnk)$
- Data transfer amount:
 - A, B, C: CPU \rightarrow GPU: $O(mk+kn+mn)$
 - C: GPU \rightarrow CPU: $O(mn)$

Transfer costs are relatively smaller with larger m , n , k

In diffusion-acc [G1], how can we reduce transfer costs?





Function Calls from GPU

- Calling functions in kernel region is ok, but we need to be careful
 - “**acc routine**” directive is required by compiler to generate GPU code

```
int main()
{
    #pragma acc kernels
    {
        ... func(A[i]) ...
    }
}

#pragma acc routine
int func(int arg)
{
    :
    :
    return ...;
}
```

How about Library Functions?



- Available library functions is very limited 😞
- We cannot use `strlen()`, `memcpy()`, `fopen()`, `fflush()`... 😞
- Exceptionally, some mathematical functions are ok 😊
 - `fabs`, `sqrt`, `fmax`...
 - `#include <math.h>` is needed
- Recently, `printf()` in kernel regions is ok! 😊



Now explanation of OpenACC is finished; we will go to CUDA



OpenACC and CUDA for GPUs

- **OpenACC**

- C/Fortran + directives (`#pragma acc ...`), Easier programming
 - NVIDIA HPC SDK compiler works
 - `module load nvhpc`
 - `pgcc -acc ... XXX.c`
 - Basically for data parallel programs with for-loops
- Only for limited types of algorithms ☹️

- **CUDA**

- Most popular and suitable for higher performance
- Use “nvcc” command for compile
 - `module load cuda`
 - `nvcc ... XXX.cu`

Programming is harder, but more general



An OpenACC Program Look Like

```
int A[100], B[100];  
int i;  
#pragma acc data copy(A,B)  
#pragma acc kernels  
#pragma acc loop independent  
for (i = 0; i < 100; i++) {  
    A[i] += B[i];  
}
```

Executed on GPU
in parallel

// After kernel region finishes,
CPU can access to A[i],B[i]



A CUDA Program Look Like

Sample:

</gs/hs1/tga-ppcomp/21/add-cuda/>

```
int A[100], B[100];
int *DA, *DB;
int i;
cudaMalloc(&DA, sizeof(int)*100);
cudaMalloc(&DB, sizeof(int)*100);
cudaMemcpy(DA,A,sizeof(int)*100,
           cudaMemcpyHostToDevice);
cudaMemcpy(DB,B,sizeof(int)*100,
           cudaMemcpyHostToDevice);
```

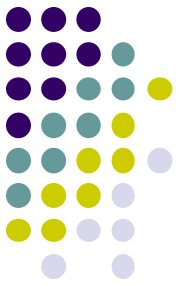
```
add<<<20, 5>>>(DA, DB);
```

```
cudaMemcpy(A,DA,sizeof(int)*100,
           cudaMemcpyDeviceToHost);
```

```
__global__ void add
(int *DA, int *DB)
{
    int i = blockIdx.x*blockDim.x
          + threadIdx.x;
    DA[i] += DB[i];
}
```

Executed on GPU
(called a *kernel function*)

We have to separate code regions executed on CPU and GPU



Using add-cuda Sample

```
[make sure that you are at a interactive node (r7i7nX) ]  
module purge    [If you have loaded nvhpc, delete it]  
module load cuda    [Do once after login]  
cd ~/t3workspace    [Example in web-only route]  
cp -r /gs/hs1/tga-ppcomp/21/add-cuda .  
cd add-cuda  
make  
[An executable file “add” is created]  
./add
```

- ※ **[Standard route]** A log-in node does not have a GPU
→ You can compile the sample there, but the program does not work!

Preparing Data on Device Memory

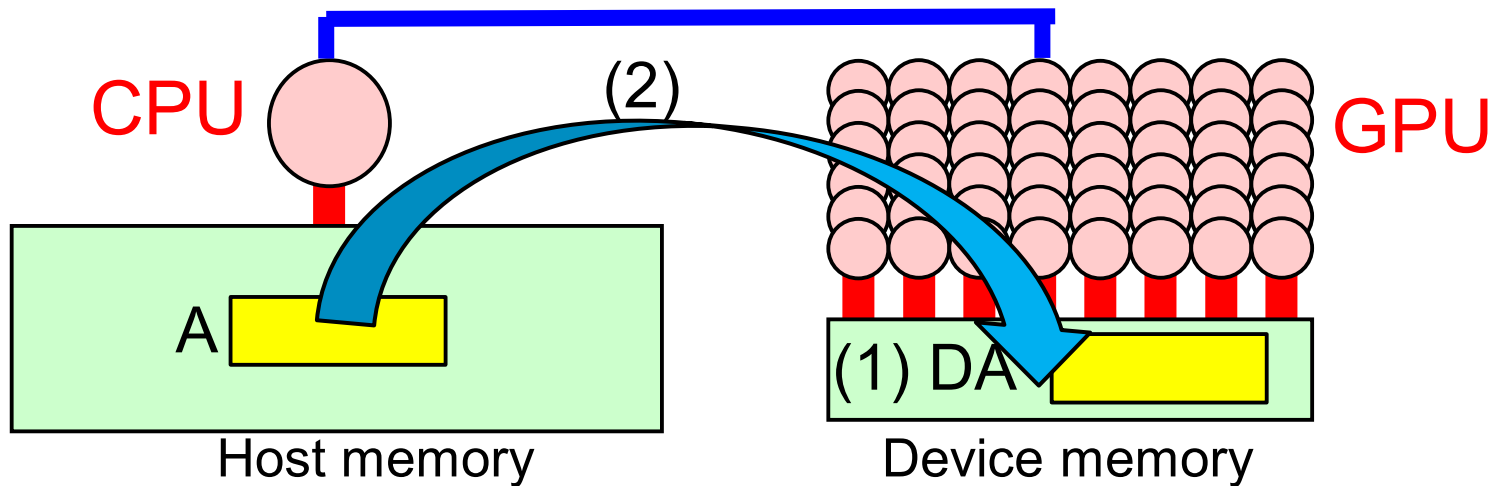


(1) Allocate a region on device memory

cf) `cudaMalloc((void**)&DA, size);`

(2) Copy data from host to device

cf) `cudaMemcpy(DA, A, size, cudaMemcpyDefault);`



Note: `cudaMalloc` and `cudaMemcpy` must be called on CPU, NOT on GPU

Comparing OpenACC and CUDA



OpenACC

Both allocation and copy are done by **acc data copyin**

One variable name A may represent both

- A on host memory
- A on device memory

```
int A[100]; ← on CPU
#pragma acc data copy(A)
#pragma acc kernels
```

```
{
  ... A[i] ...
}
```

← on GPU

CUDA

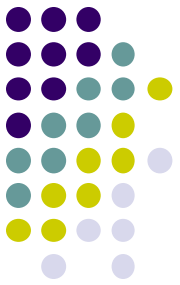
cudaMalloc and **cudaMemcpy** are separated

Programmer have to prepare two pointers, such as A and DA

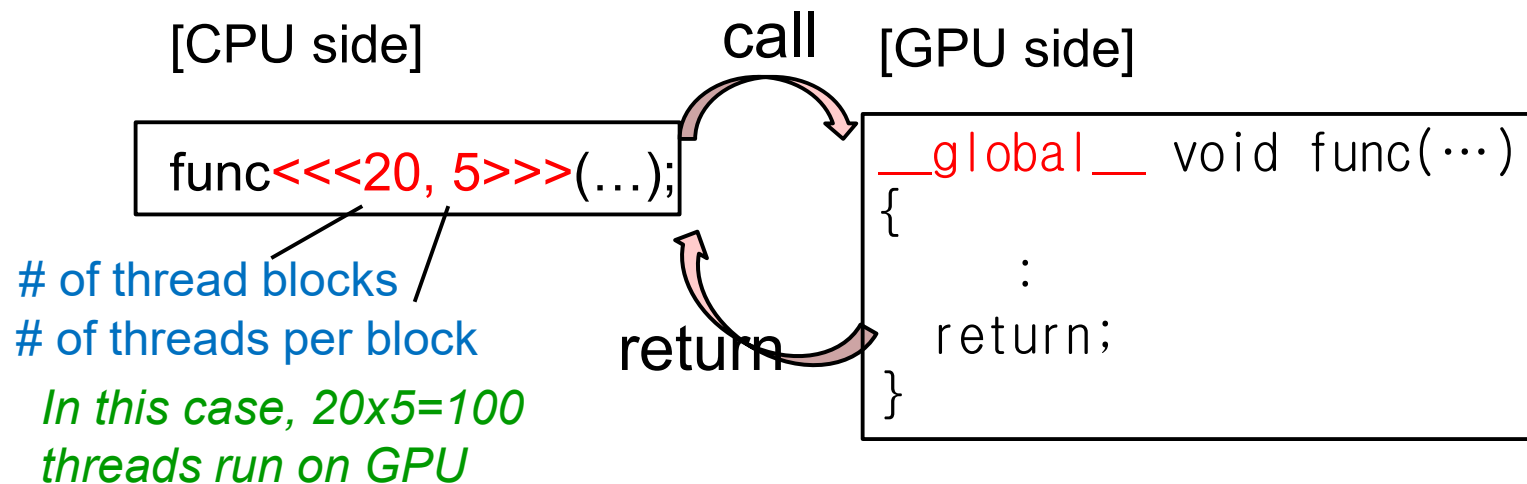
```
int A[100];
int *DA;
cudaMalloc(&DA, ...);
cudaMemcpy(DA, A, ..., ...);
// Here CPU cannot access DA[i]

func<<<..., ...>>>(DA, ...);
```

Calling A GPU Kernel Function from CPU



- A region executed by GPU must be a distinct function
 - called a GPU kernel function

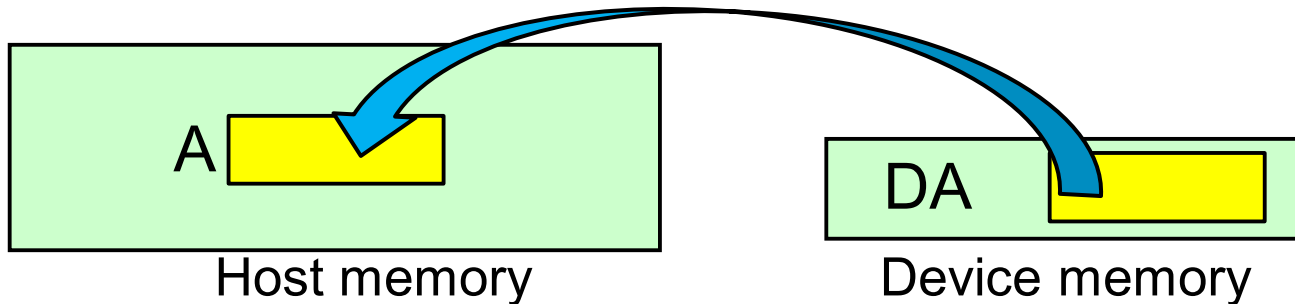


A GPU kernel function (called from CPU)

- needs **__global__** keyword
- can take parameters
- can **NOT** return value; return type must be void



Copying Back Data from GPU



- Copy data using `cudaMemcpy`
 - cf) `cudaMemcpy(A, DA, size, cudaMemcpyDefault);`
 - 4th argument is one of
 - `cudaMemcpyHostToDevice`, `cudaMemcpyDeviceToHost`
 - `cudaMemcpyDeviceToDevice`, `cudaMemcpyHostToHost`
 - `cudaMemcpyDefault` ← Detect memory type automatically 😊
- When a memory area is unnecessary, free it
 - cf) `cudaFree(DA);`

Assignments in GPU Part (Abstract)



Choose one of [G1]—[G3], and submit a report

Due date: May 27 (Thursday)

[G1] Parallelize “diffusion” sample program by OpenACC or CUDA

[G2] Evaluate speed of “mm-acc” or “mm-cuda” in detail

[G3] (Freestyle) Parallelize *any* program by OpenACC or CUDA.

Notes in Report Submission (1)



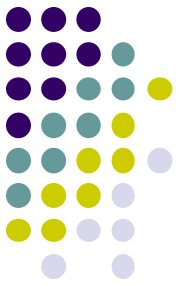
- Submit the followings via **T2SCHOLA**
 - (1) **A report document**
 - PDF, MS-Word or text file
 - 2 pages or more
 - in English or Japanese (日本語もok)
 - (2) **Source code files** of your program
 - Try “zip” to submit multiple files

Notes in Report Submission (2)



The report document should include:

- Which problem you have chosen
- How you parallelized
 - It is even better if you mention efforts for high performance or new functions
- Performance evaluation on TSUBAME
 - With varying number of threads
 - With varying problem sizes
 - Discussion with your findings
 - Other machines than TSUBAME are ok, if available



Next Class:

- GPU Programming (3) on May 13
 - Multi-threads on CUDA
- Also please note due date of OpenMP assignment is May 13