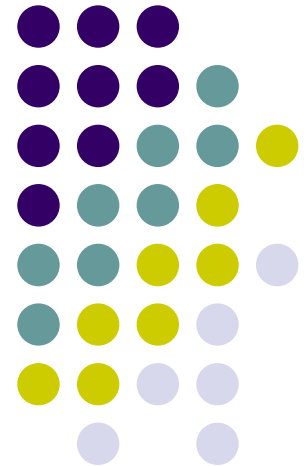


# Practical Parallel Computing (実践的並列コンピューティング) 2021 No. 10

Part2: GPU (4)  
May 17, 2021

Toshio Endo  
School of Computing & GSIC  
[endo@is.titech.ac.jp](mailto:endo@is.titech.ac.jp)





# Overview of This Course

- Part 0: Introduction
  - 2 classes
- Part 1: OpenMP for shared memory programming
  - 4 classes
- Part 2: **GPU** programming
  - 4 classes **← We are here (4/4)**
  - OpenACC (1.5 classes) and **CUDA (2.5 classes)**
- Part 3: **MPI** for distributed memory programming
  - 3 classes

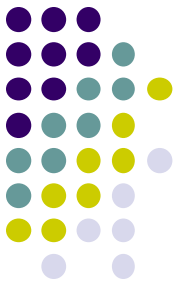
# Comparing OpenMP/OpenACC/CUDA



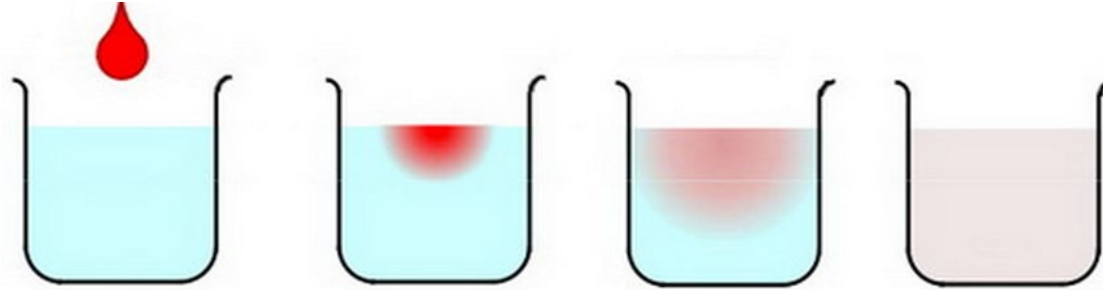
	OpenMP	OpenACC	CUDA
Processors	CPU	CPU+GPU	
File extension	.c, .cc		.cu
To start parallel (GPU) region	#pragma omp parallel	#pragma acc kernels	func<<<..., ...>>>()
To specify # of threads	export OMP_NUM_THREADS=...	(num_gangs, vector_length etc)	
Desirable # of threads	# of CPU cores or less	# of GPU cores or “more”	
To get thread ID	omp_thread_num()	-	blockIdx, threadIdx
Parallel for loop	#pragma omp for	#pragma acc loop	-
Task parallel	#pragma omp task	-	-
To allocate device memory	-	#pragma acc data	cudaMalloc()
To copy to/from device memory	-	#pragma acc data #pragma acc update	cudaMemcpy()
Function on GPU	-	#pragma acc routine	__global__, __device__

※ “# of XXX” = “The number of XXX”

# “diffusion” Sample Program related to [G1]



An example of diffusion phenomena:



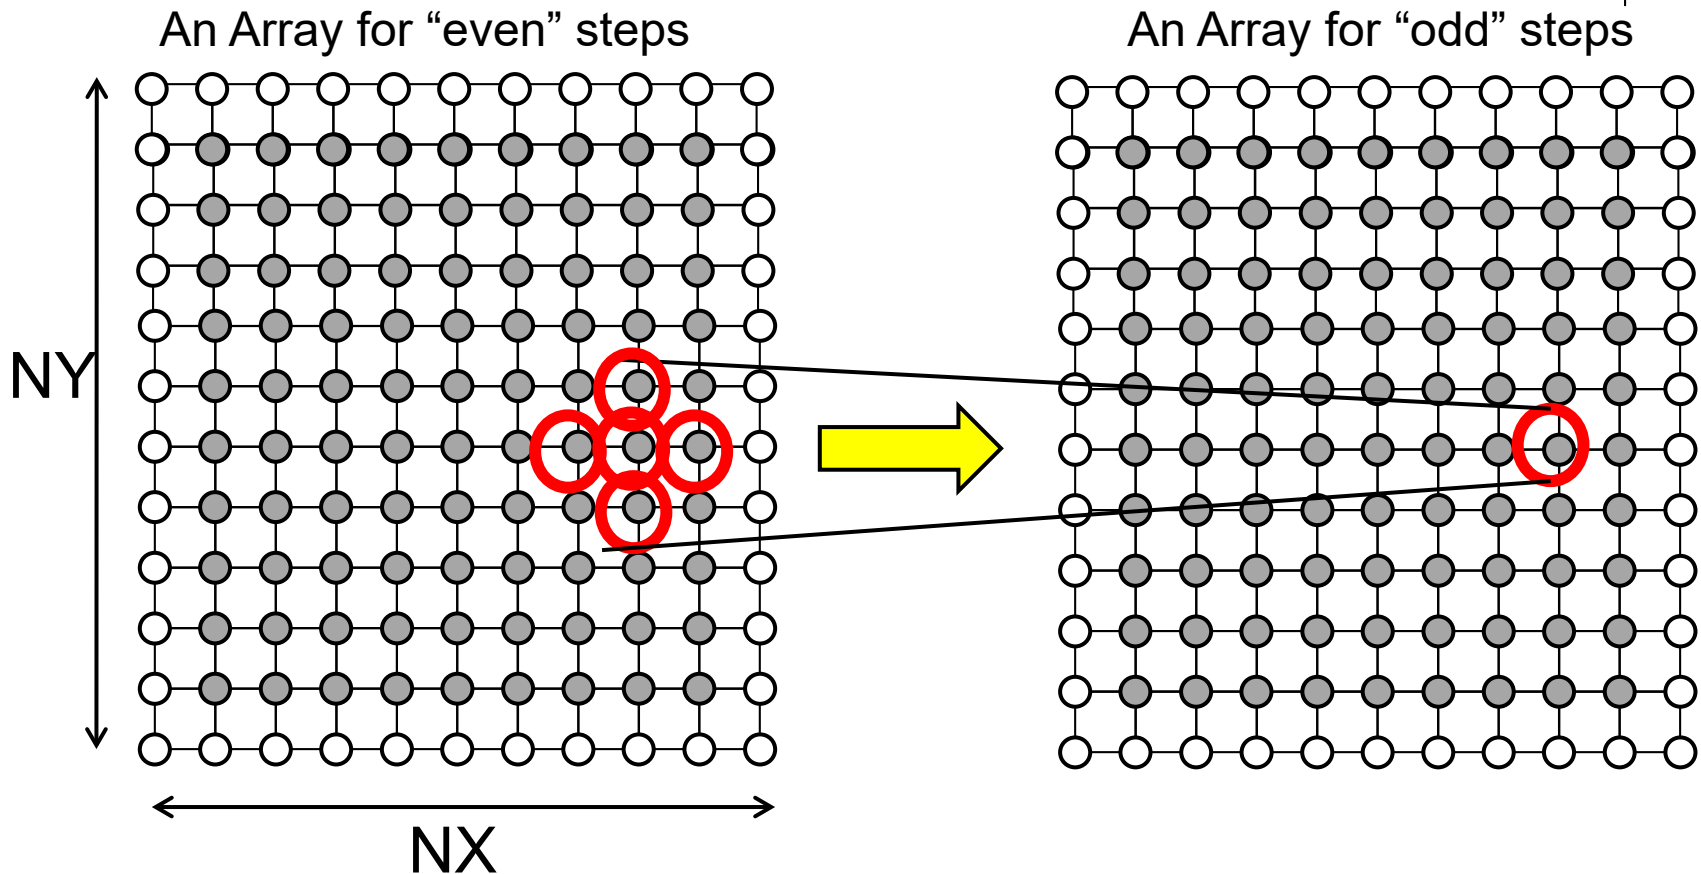
The ink spreads gradually, and finally the density becomes uniform (Figure by Prof. T. Aoki)

Available at </gs/hs1/tga-ppcomp/21/diffusion/>  
You can use </gs/hs1/tga-ppcomp/21/diffusion-cuda/>

- Execution: `./diffusion [nt]`
  - nt: Number of time steps



# Discussion on diffusion sample



Both arrays have to be on GPU device memory when computations are done

# Consideration of Parallelizing Diffusion with CUDA

## related to [G1]



- x, y loops can be parallelized
- t loop cannot be parallelized

[Data transfer from CPU to GPU]

```
for (t = 0; t < nt; t++) {
```

```
    for (y = 1; y < NY-1; y++) {  
        for (x = 1; x < NX-1; x++) {  
            :  
        }  
    }
```

```
}
```

[Data transfer from GPU to CPU]

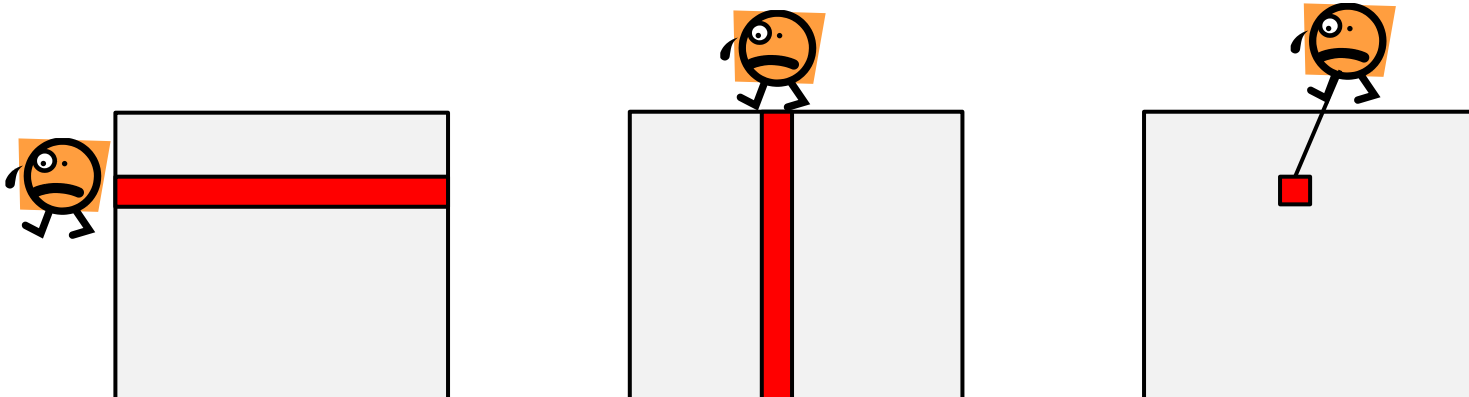
GPU computation must be a distinct function (GPU kernel function)

It's better to transfer data *out of* t-loop

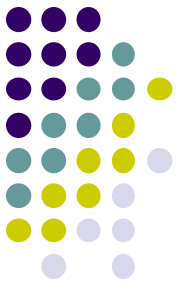


# Considering CUDA Threads

- How do we design threads on CUDA?
- There are several choices in [G1]
  - 1thread = 1row
    - We use NY threads in total → only x-loop in kernel function
  - 1thread = 1column
    - We use NX threads in total → only y-loop in kernel function
  - 1thread = 1element
    - We use NX x NY threads in total → No loop in kernel function!
    - This looks fast since the number of threads is very large

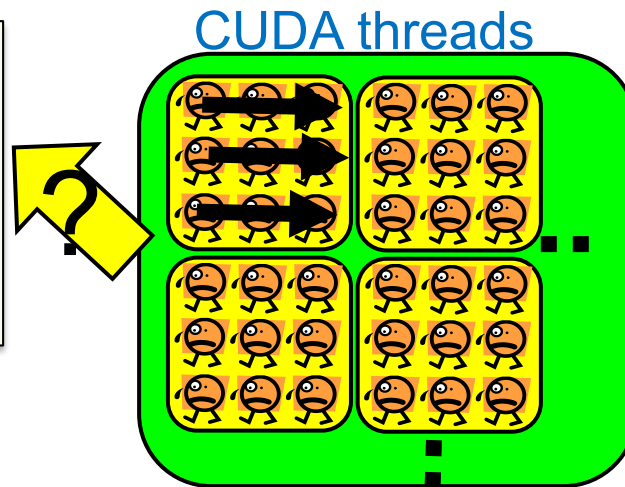
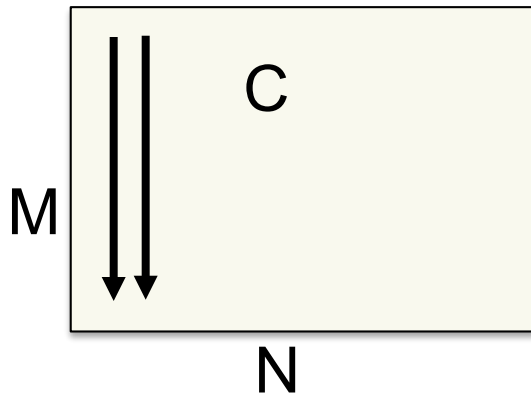


# Mapping between Threads and Data



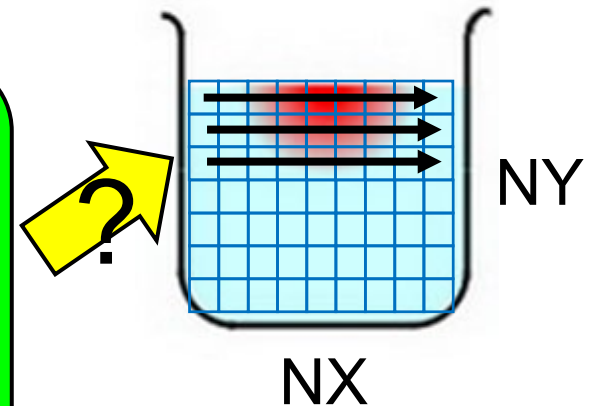
mm-cuda:

Matrices has  
column-major format



diffusion:

2D array has  
row-major format



```
j = blockIdx.y * blockDim.y +  
threadIdx.y;  
i = blockIdx.x * blockDim.x +  
threadIdx.x;  
: This thread computes Cij
```

```
y = blockIdx.y * blockDim.y +  
threadIdx.y;  
x = blockIdx.x * blockDim.x +  
threadIdx.x;  
: This thread computes [y][x]
```

[Q] What if the dimensions are exchanged?

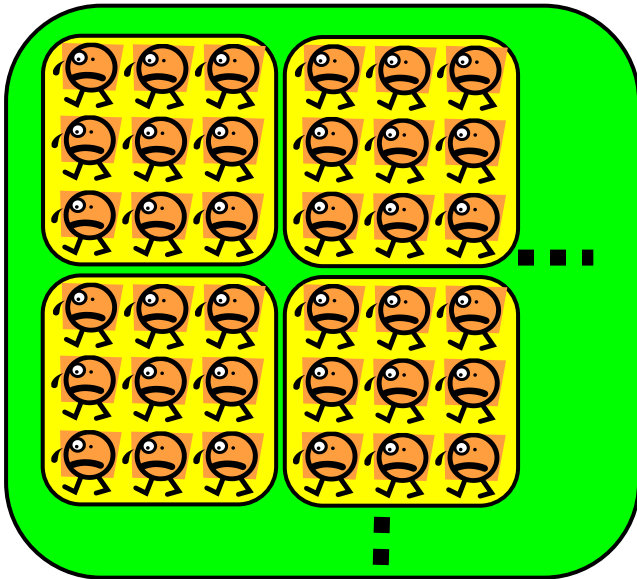




# Considering gridDim/blockDim (1)

```
func <<< dim3 ( ?, ?, ? ), dim3 ( ?, ?, ? ) >>> (...);
```

gridDim                      blockDim



(1) We decide total number of threads

→ (NX, NY, 1) threads

- See notes on the next page

(2) We tune each block size (blockDim)

→ Good candidates are (4, 4, 1), (8, 8, 1), (16, 16, 1), (32, 32, 1)

- The number must be  $\leq 1024$
- How about non-square blocks?

(3) Then block number (gridDim) is determined

We should consider indivisible cases

# Considering gridDim/blockDim (2)



- In diffusion, Points  $[1, NX-1) \times [1, NY-1)$ , excluded boundary, should be computed

There are choices:

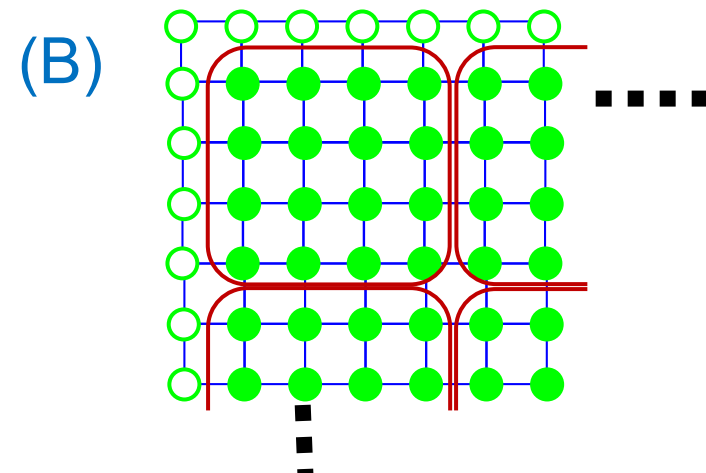
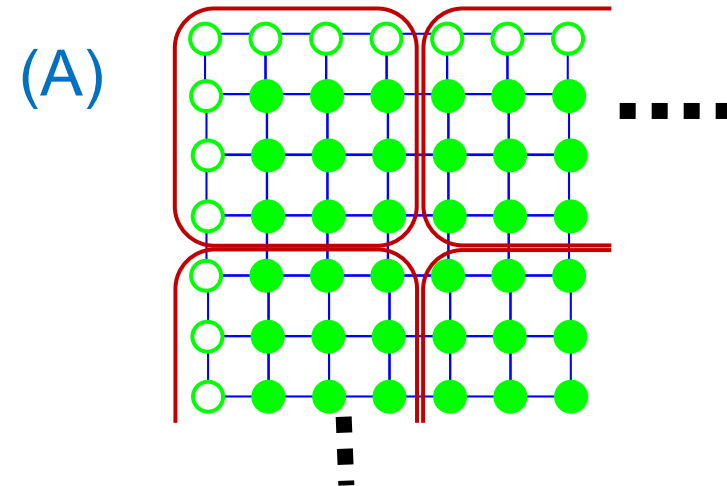
(A) Create  $NX \times NY$  threads

- Thread  $(x,y)$  computes  $(x,y)$
- Threads with below IDs do nothing
  - $x == 0$  or  $y == 0$  or  $x \geq NX-1$  or  $y \geq NY-1$

(B) Create  $(NX-2) \times (NY-2)$  threads

- Thread  $(x,y)$  computes  $(x+1,y+1)$
- Threads with below IDs do nothing
  - $x \geq NX-2$  or  $y \geq NY-2$

Either is ok 😊



# Discussion on Data Transfer of Diffusion



Both codes will work, but how about speeds?

[Data transfer from CPU to GPU]

```
for (t = 0; t < nt; t++) {  
    :
```

```
    for (y = 1; y < NY-1; y++) {  
        for (x = 1; x < NX-1; x++) {  
            :  
        }  
    }  
}
```

```
}
```

[Data transfer from GPU to CPU]

Computation:  $O(NX \ NY \ nt)$   
Transfer:  $O(NX \ NY)$

```
for (t = 0; t < nt; t++) {  
    :
```

[Data transfer from CPU to GPU]

```
for (y = 1; y < NY-1; y++) {  
    for (x = 1; x < NX-1; x++) {  
        :  
    }  
}
```

```
}
```

[Data transfer from GPU to CPU]

Computation:  $O(NX \ NY \ nt)$   
Transfer:  $O(\underline{NX \ NY \ nt})$

# Speed of GPU Programs and GPU Architecture



**Advanced topic**

Case 1: How should block-size be determined?

When creating 1,000,000 threads,

- `<<<1, 1000000>>>` causes an error
  - blockDim must be  $\leq 1024$
- `<<<1000000, 1>>>` can work, but slow
- `<<<1000, 1000>>>` is faster  $\rightarrow$  Why?

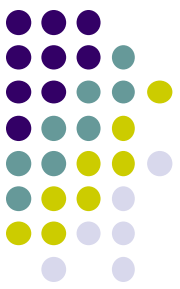


Case 2: How should each thread access memory?

- In mm-cuda,  $(x = \text{row}, y = \text{col})$  and  $(x = \text{col}, y = \text{row})$  shows different speed

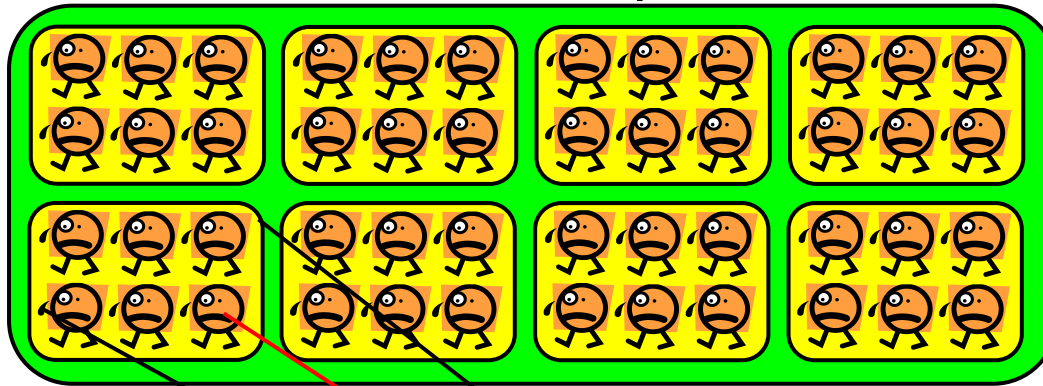
Knowledge of GPU architecture helps understanding of speeds

# Why Do We Have to Specify both `gridDim` and `blockDim`?



- and why did NVIDIA decide so?

→ Hierarchical structure of GPU processor is considered



Structure of P100 GPU  
(16nm, 15Billion transistors)

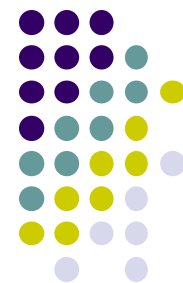
1 GPU = 56 **SMXs**

1 **SMX** = 64 CUDA cores  
(16 cores x 4 groups)

→ 1GPU=3,584 CUDA cores



# Mapping between Threads and Cores

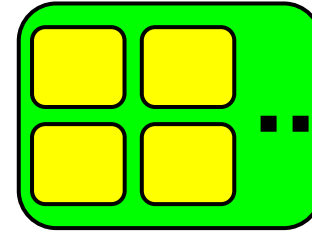
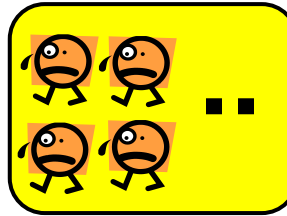


- 1 thread blocks (or more) run on 1 SMX
  - At least 56 blocks are needed to use all SMXs on P100
  - `gridDim (gx*gy*gz)` should be  $\geq 56$
- 1 thread (or more) run on a CUDA core
  - At least  $56 \times 64 = 3584$  threads in total are needed to use all CUDA cores on P100
  - `Total threads (gx*gy*gz * bx*by*bz)` should be  $\geq 3584$
- 32 consecutive threads (in a block) are batched (called a warp) and scheduled
  - At least 32 threads per block are needed for performance
  - `blockDim (bx*by*bz)` should be  $\geq 32$

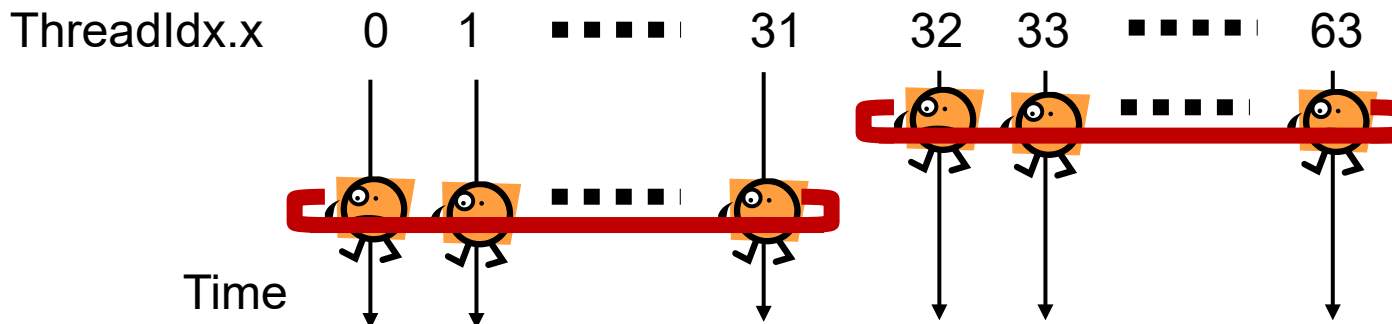


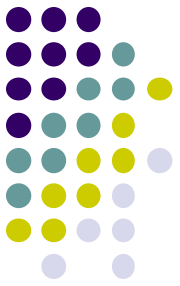
# Warp: Internal Execution Unit

thread < **warp** < thread block < grid



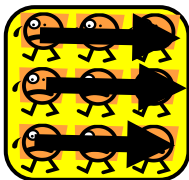
- Threads in a thread block are internally divided into “**warp**”, a group of contiguous 32 threads
- 32 threads in a warp always are executed synchronously
  - They execute the same instruction simultaneously
  - Only 1 program counter for 32 threads → GPU hardware is simplified
  - Actually 32 threads are executed on 16 CUDA cores





# Observations due to Warps

- If number of threads per block (blockDim) is not  $32 \times n$ , it is inefficient
  - Even if blockDim=1, the system creates a warp for it
- Characteristics in memory addresses accessed by threads in a warp affect the performance
  - Coalesced accesses are fast



⌘ In multi-dimensional cases (blockDim.y>1 or blockDim.z>1), “neighborhood” is defined by x-dimension

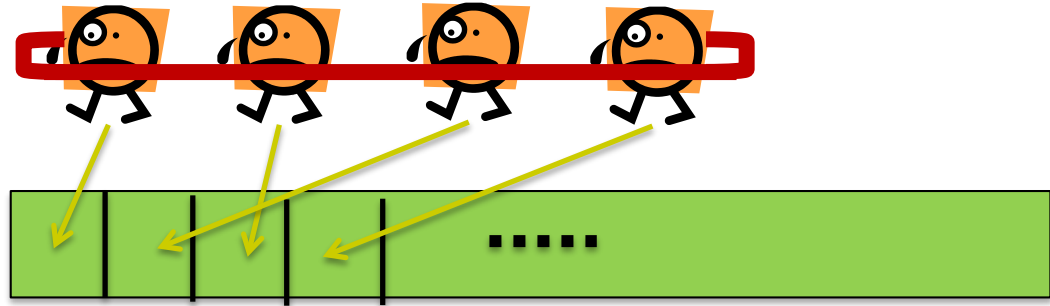




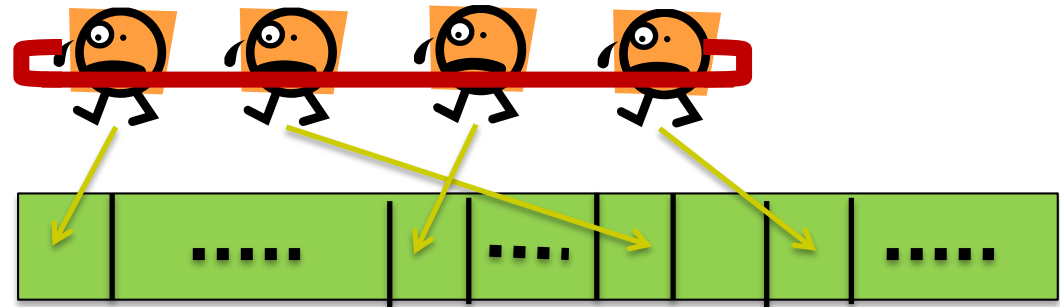
# Coalesced Memory Access

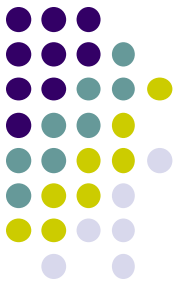
- When threads in a warp access “neighbor” address on memory (**coalesced access**), it is more efficient

Coalesced access  
→ **Faster**



Non-coalesced access  
→ **Slower**

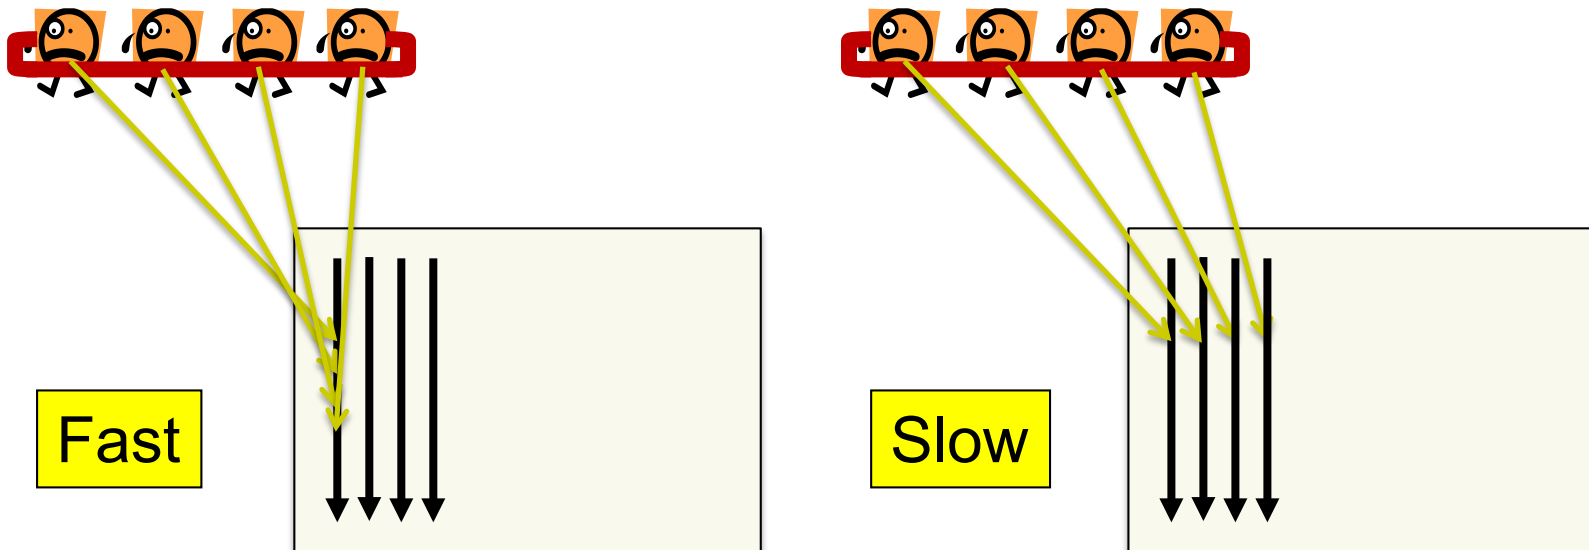




# Accesses in mm-cuda Sample

- **mm-cuda**: ( $x = \text{row}, y = \text{col}$ )  $\rightarrow$  coalesced and fast
- **mm-nc-cuda**: ( $x = \text{col}, y = \text{row}$ )  $\rightarrow$  non-coalesced and slow

We should see “what data are accessed by threads in a warp simultaneously”



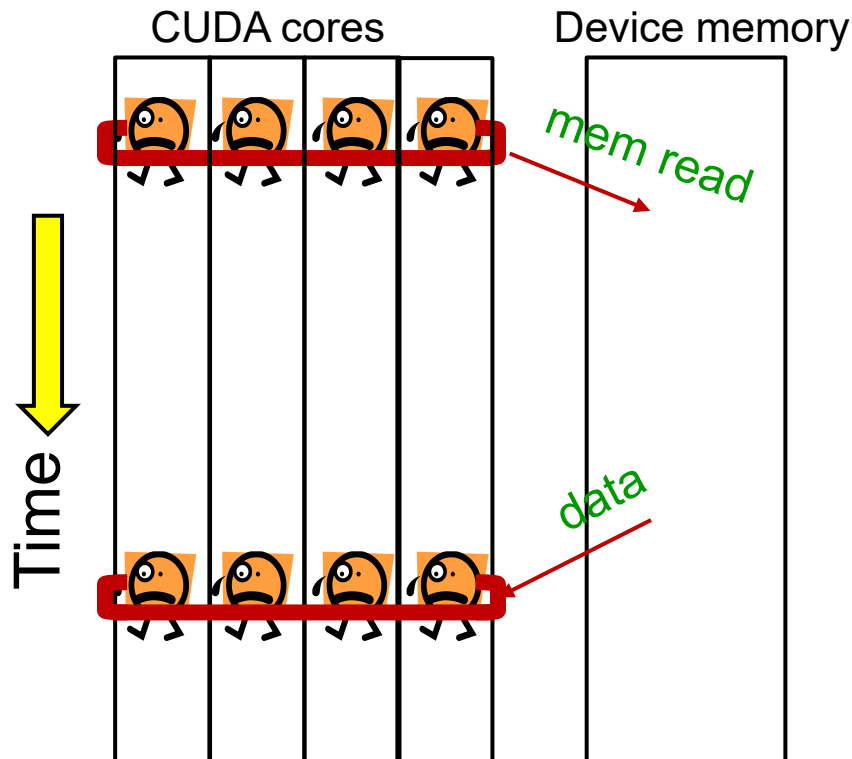
matrices in column-major format

# Why $\#threads \gg \#cores$ Works Well on GPUs?

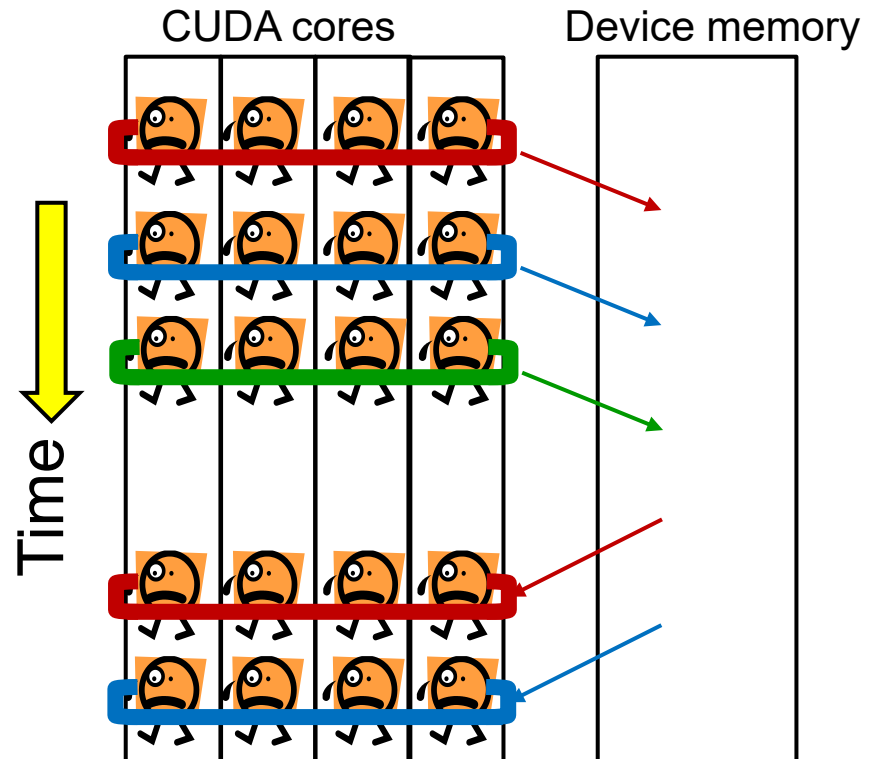


- GPU supports very fast ( $\sim 1$  clock) context switches  
→ With many threads, memory access latency can be hidden

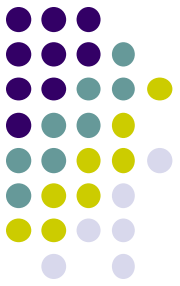
$\#threads == \#cores$



$\#threads > \#cores$



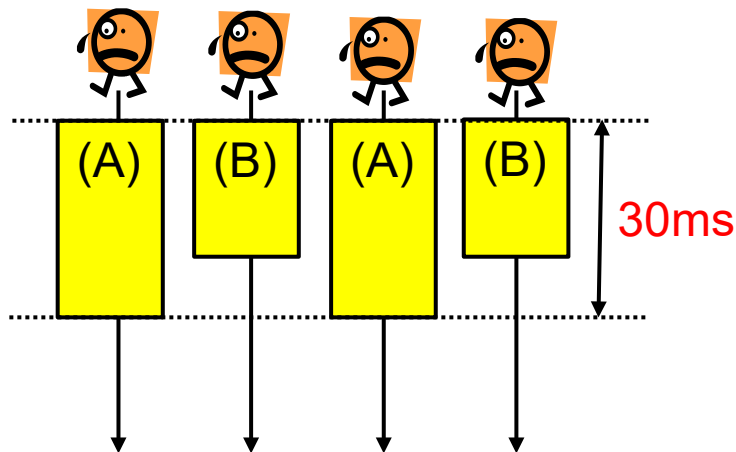
# Considering Branches in Parallel Programs



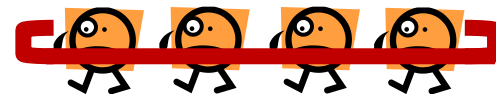
Consider this code. How long is execution time?

```
if (thread-id % 2 == 0) {  
    : // (A) 30msec  
} else {  
    : // (B) 20msec  
}
```

On CPU (OpenMP)

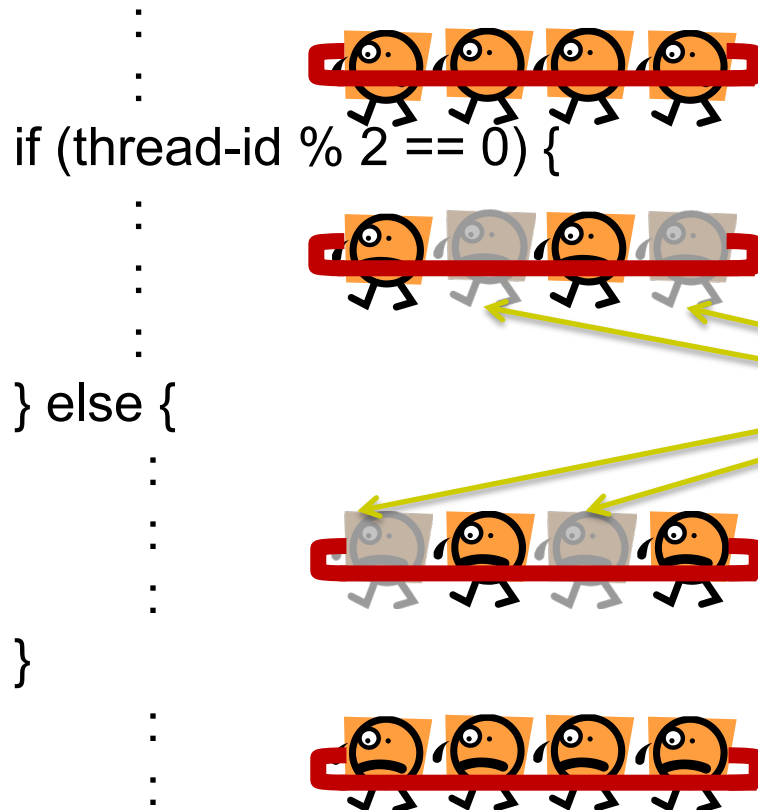


On GPU, threads in a warp must execute the same instruction. What happens?





# Branches on GPU (1)



Some threads are made sleep  
Both "then" and "else" are executed!

→ Answer to previous question is **50ms** !

⌘ Similar cases happen in for, while...



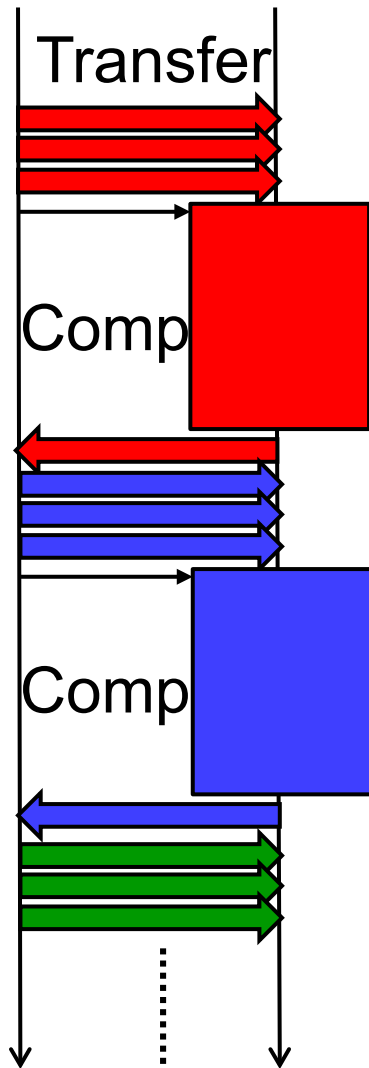
# Branches on GPU (2)

- As exceptional cases, if threads in a warp “agree” in branch condition, either “then” part or “else” part is executed → Efficient!
  - If there is difference of opinion (previous page), it is called a **divergent branch**
- Agreement among buddies (threads in a warp) is important for speed



# Considering Data Transfer Costs

CPU GPU



*Example case:* We are going to compute multiply for different matrices

- Input data are on host memory

- $C1 = A1 \times B1$

- $C2 = A2 \times B2$

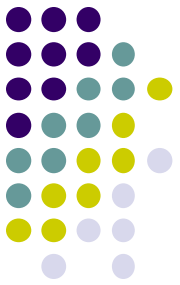
....

- $Cn = An \times Bn$

- In default, GPU cannot compute during transfer  
→ **cudaStream** is useful for hiding transfer costs

This is also useful for speed-up of mm-cuda, by dividing matrices into pieces

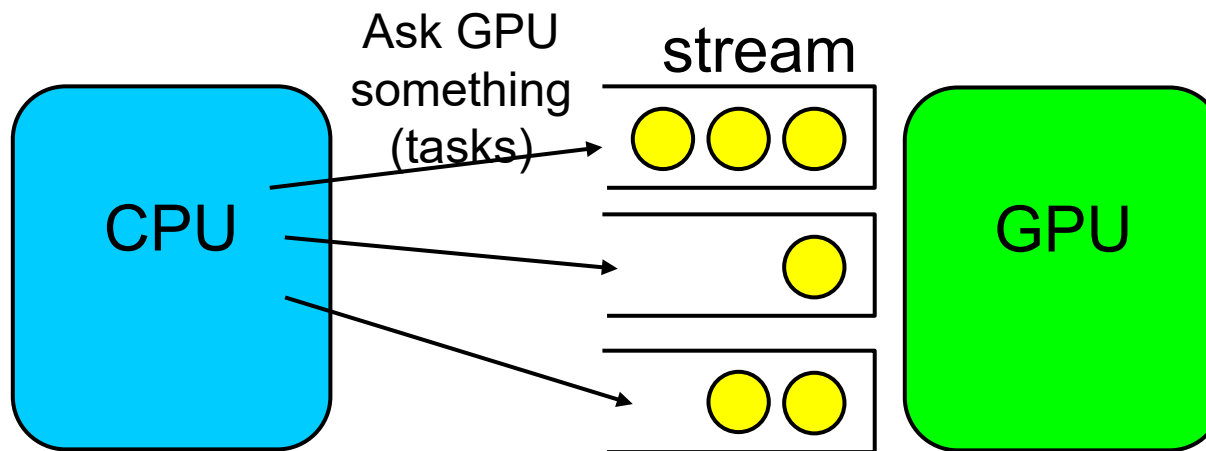
# Asynchronous Executions with `cudaStream` (1)



What are `streams`?

- GPU's "service counters" that accept tasks from CPU
  - Each stream looks like a queue
- "Tasks" from CPU to GPU include
  - Data transfer (Host → Device)
  - GPU kernel function call
  - Data transfer (Device → Host)

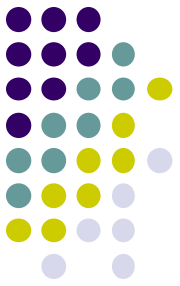
*"Tasks" here are a bit different from "omp task"*



All of sample programs are using the "default stream"



# Asynchronous Executions with cudaStream (2)



Create a stream

```
cudaStream_t str;  
cudaStreamCreate(&str); // Create a stream
```

Data transfer using a specific stream

```
cudaMemcpyAsync(dst, src, size, type, str);
```

Call GPU kernel function using a stream

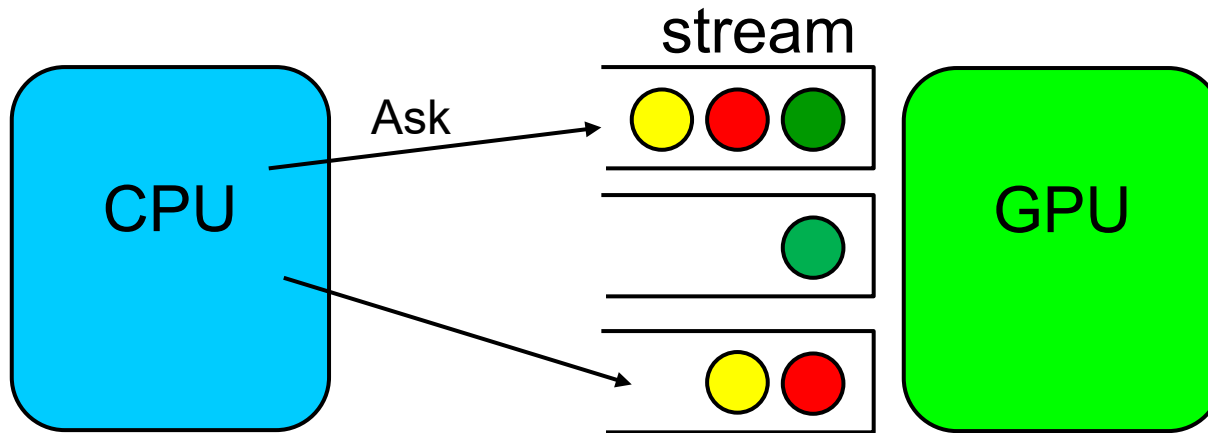
```
func<<<gs, bs, 0, str>>>( ... );  
// 3rd parameter is related to for “shared memory”
```

Wait until all tasks on a stream are finished

```
cudaStreamSynchronize(str);
```



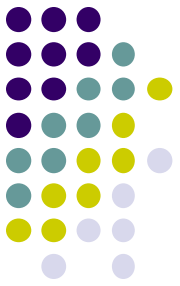
# How GPU Executes Tasks



- Tasks on the same stream is done in FIFO
- If tasks are in different streams, and have different kinds, they may be done simultaneously
  - Kinds:  $H \rightarrow D$ , kernel,  $D \rightarrow H$
  - Note: If tasks are in the same kind, no speed up

This is not a unique solution;  
Use 2 or 3 streams repeatedly → we can save  
memory and stream resources

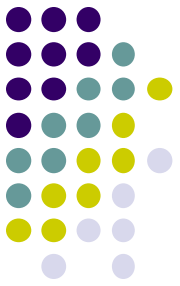




# More Things to Study

- Using CUDA shared memory
  - fast and small memory than device memory
- Unified memory in recent CUDA
  - `cudaMemcpy` can be omitted for automatic data transfer
- Using Tensor-core to accelerate deep learning
  - Only on V100 GPUs or later
  - Unfortunately, TSUBAME3 has older P100 ☹
- Using multiple GPUs towards petascale computation
  - MPI+CUDA, MPI+OpenACC
- More and more...

# Assignments in GPU Part (Abstract)



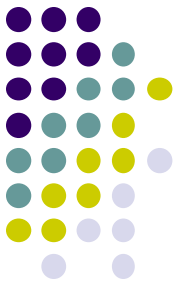
Choose one of [G1]—[G3], and submit a report

Due date: May 27 (Thursday)

[G1] Parallelize “diffusion” sample program by OpenACC or CUDA

[G2] Evaluate speed of “mm-acc” or “mm-cuda” in detail

[G3] (Freestyle) Parallelize *any* program by OpenACC or CUDA.



# Next Class:

- Part 3: MPI Programming (1)
  - Introduction to distributed memory parallel programming