Practical Parallel Computing
(実践的並列コンピューティング)

Part2: GPU (4)
May 16, 2022

Toshio Endo

School of Computing & GSIC

endo@is.titech.ac.jp

Overview of This Course

- Part 0: Introduction
  - 2 classes
- Part 1: OpenMP for shared memory programming
  - 4 classes
- Part 2: GPU programming
  - 4 classes   ← We are here (4/4)
  - OpenACC (1.5 classes) and CUDA (2.5 classes)
- Part 3: MPI for distributed memory programming
  - 3 classes

Comparing OpenMP/OpenACC/CUDA

	OpenMP	OpenACC	CUDA
Processors	CPU	CPU+GPU	
File extension	.c, .cc		.cu
To start parallel (GPU) region	#pragma omp parallel	#pragma acc kernels	func<<<…, …>>>()
To specify # of threads	export OMP_NUM _THREADS=…	(num_gangs, vector_length etc)	
Desirable # of threads	# of CPU cores or less	# of GPU cores or "more"	
To get thread ID	omp_thread_num()	-	blockIdx, threadIdx
Parallel for loop	#pragma omp for	#pragma acc loop	-
Task parallel	#pragma omp task	-	-
To allocate device memory	-	#pragma acc data	cudaMalloc()
To copy to/from device memory	-	#pragma acc data #pragma acc update	cudaMemcpy()
Function on GPU	-	#pragma acc routine	__global__, __device__

※ "# of XXX" = "The number of XXX"

Speed of GPU Programs and GPU Architecture

Case 1: How should block-size be determined?

When creating 1,000,000 threads,
- <<<1, 1000000>>> causes an error
  - blockDim must be <= 1024
- <<<1000000, 1>>> can work, but slow
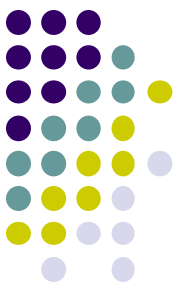- <<<1000, 1000>>> is faster → Why?

Case 2: How should each thread access memory?

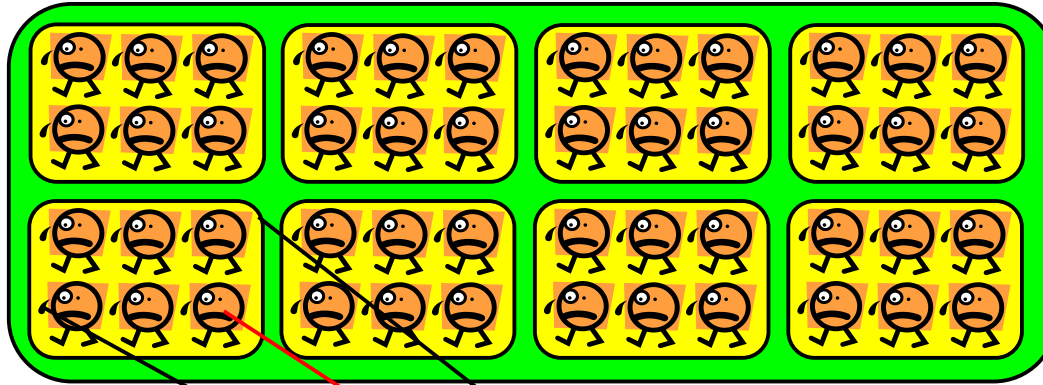- In mm-cuda, (x = row, y = col) and (x = col, y = row) shows different speed

Knowledge of GPU architecture helps understanding of speeds

4

Why Do We Have to Specify both gridDim and blockDim?

- and why did NVIDIA decide so?

→ Hierarchical structure of GPU processor is considered

Structure of P100 GPU
(16nm, 15Billion transistors)

1 GPU = 56 SMXs
1 SMX = 64 CUDA cores
   (16 cores x 4 groups)

→ 1GPU=3,584 CUDA cores

Mapping between Threads and Cores

- 1 thread blocks (or more) run on 1 SMX
  - → At least 56 blocks are needed to use all SMXs on P100
  - → gridDim (gx*gy*gz) should be $\geqq$ 56
- 1 thread (or more) run on a CUDA core
  - → At least 56*64=3584 threads in total are needed to use all CUDA cores on P100
  - → Total threads (gx*gy*gz * bx*by*bz) should be $\geqq$ 3584
- 32 consective threads (in a block) are batched (called a _warp_) and scheduled
  - → At least 32 threads per block are needed for performance
  - → blockDim (bx*by*bz) should be $\geqq$ 32

Warp: Internal Execution Unit

thread < warp < thread block < grid
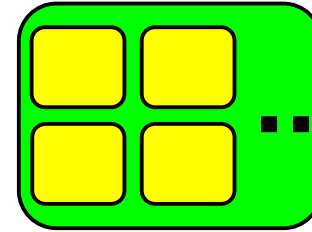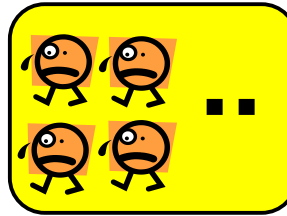
- Threads in a thread block are internally divided into "warp", a group of contiguous 32 threads
- 32 threads in a warp always are executed synchronously
  - They execute the same instruction simultaneously
  - Only 1 program counter for 32 threads → GPU hardware is simplified
 - Actually 32 threads are executed on 16 CUDA cores

ThreadIdx.x   0   1   ■ ■ ■ ■   31   32   33   ■ ■ ■ ■   63

Time
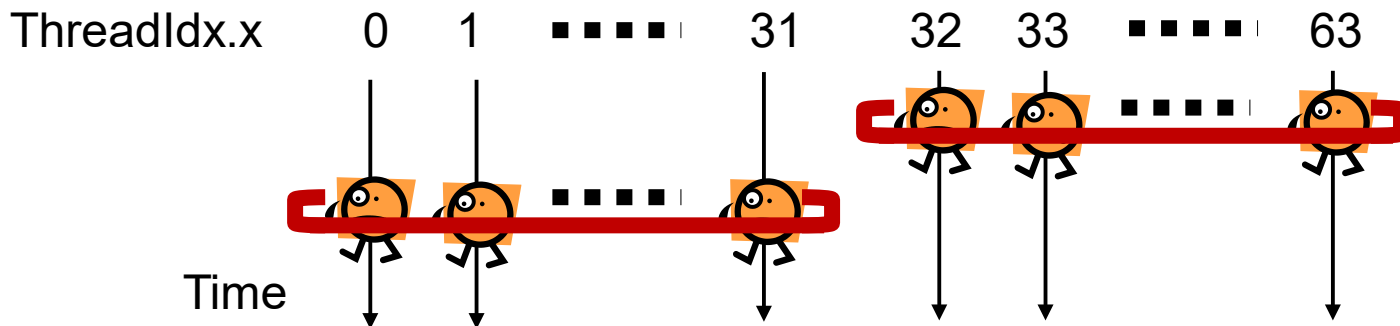
Observations due to Warps

- If number of threads per block (blockDim) is not *32 x n*, it is inefficient
  - Even if blockDim=1, the system creates a warp for it
- Characteristics in memory addresses accessed by threads in a warp affect the performance
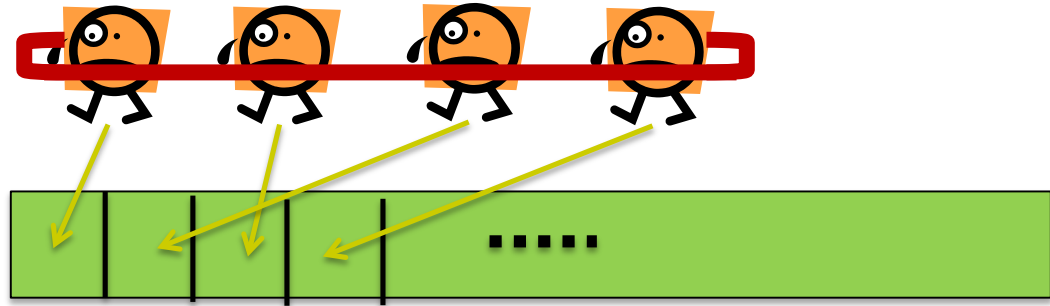 - Coalesced accesses are fast
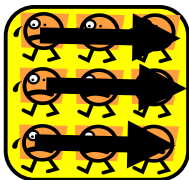
 ※ In multi-dimensional cases (blockDim.y>1 or blockDim.z>1), "neighborhood" is defined by x-dimension
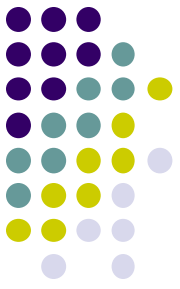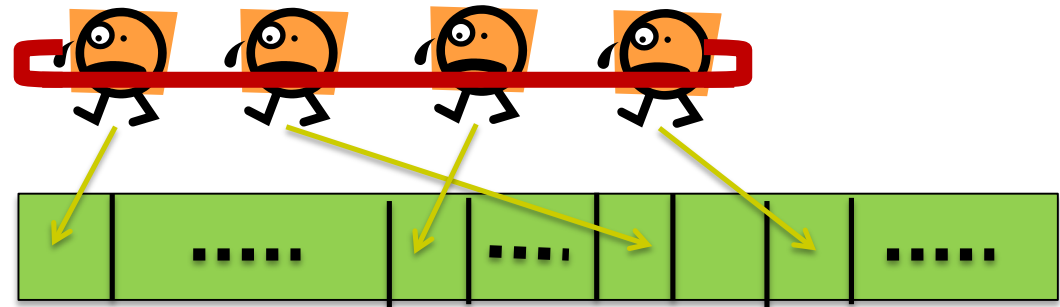
Coalesced Memory Access

- When threads in a warp access "neighbor" address on memory (coalesced access), it is more efficient

Coalesced access
→ **Faster**

Non-coalesced access
→ **Slower**

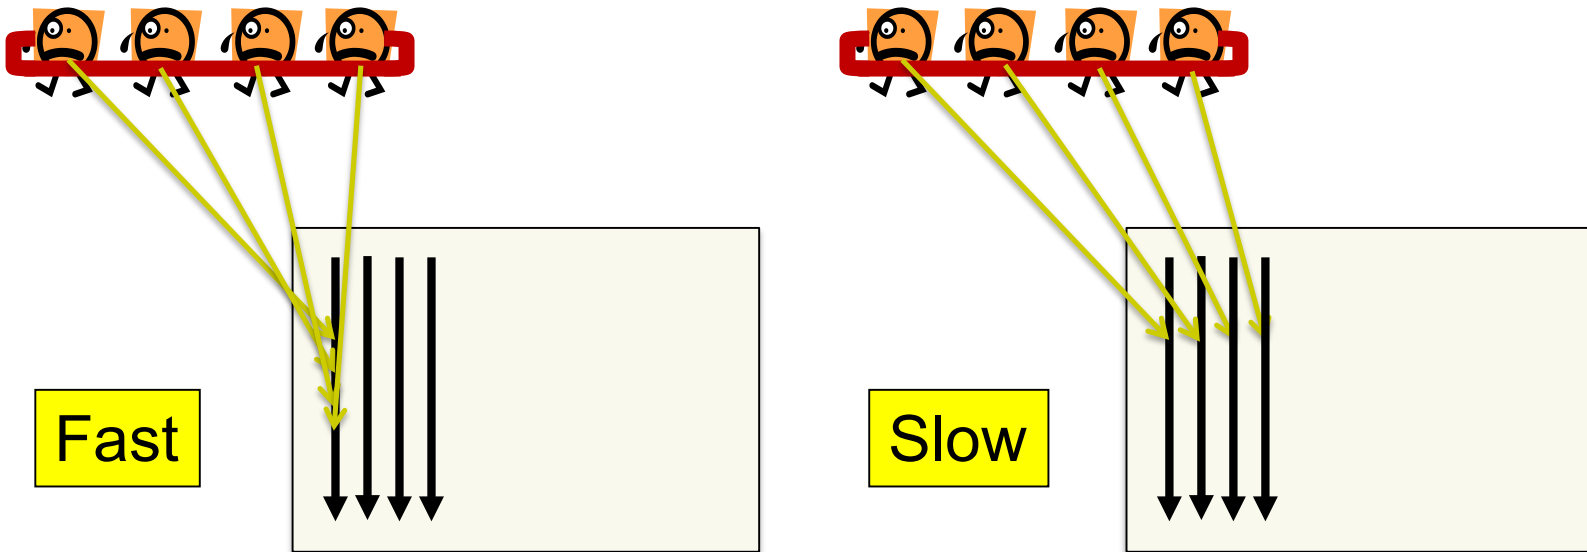# Accesses in mm-cuda Sample

- mm-cuda: (x = row,y = col) → coalesced and fast
- mm-nc-cuda: (x = col, y = row) → non-coalesced and slow
  - /gs/hs1/tga-ppcomp/22/mm-nc-cuda

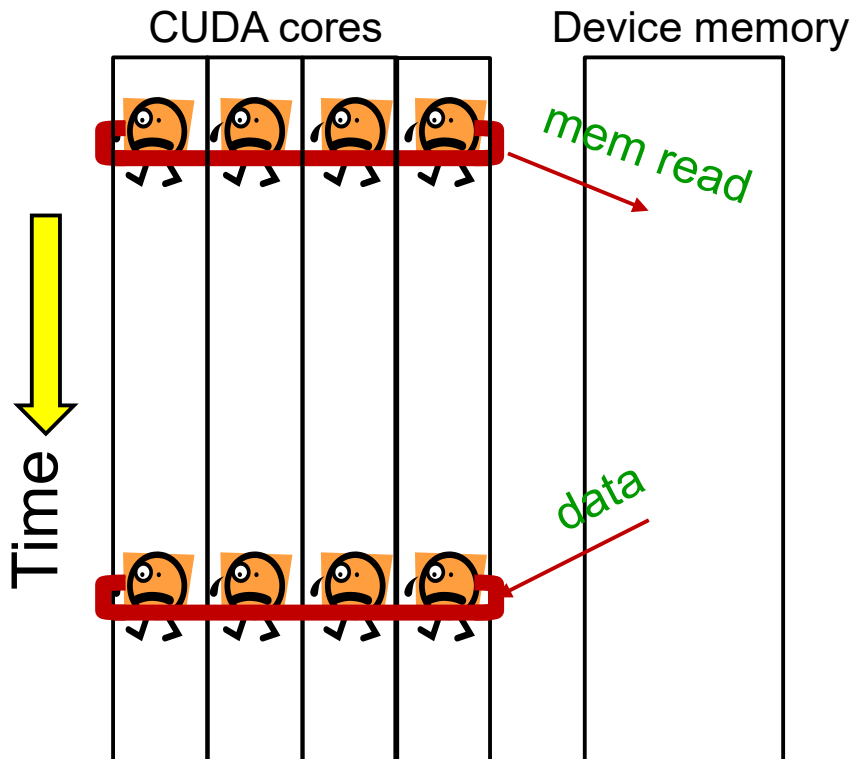We should see "what data are accessed by threads in a warp simultaneously"

Fast

Slow

matrices in column-major format

Why #threads >> #cores Works Well on GPUs?
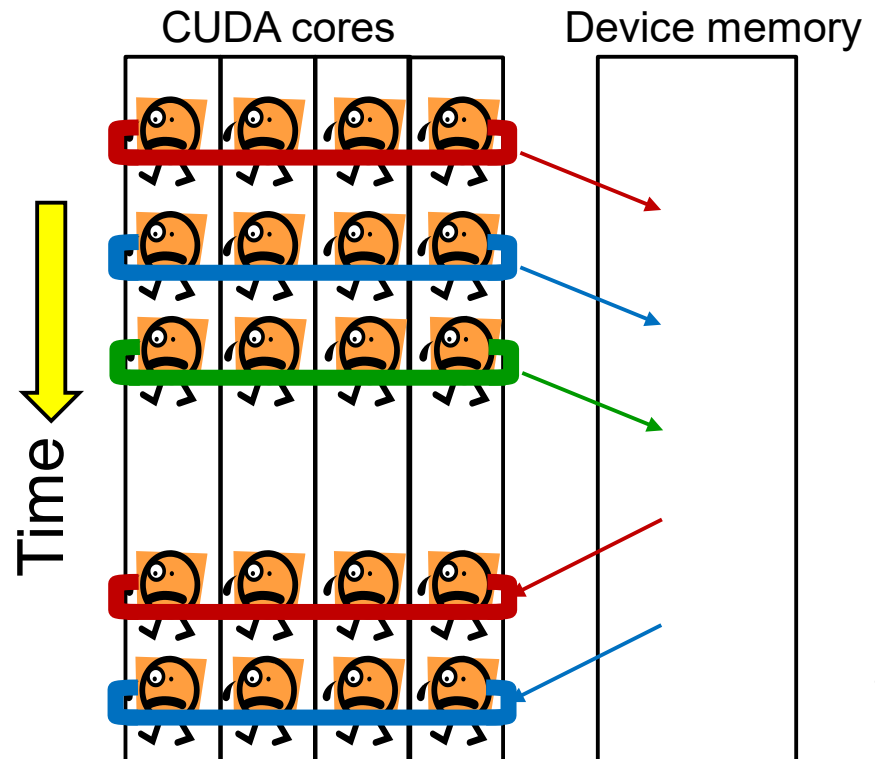
- GPU supports very fast (~1 clock) context switches
  - ➔ With many threads, memory access latency can be hidden

#threads == #cores   #threads > #cores
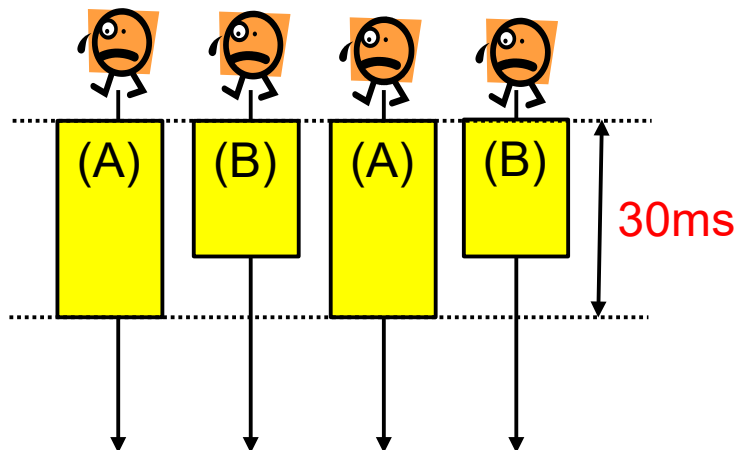
Considering Branches in Parallel Programs

Consider this code. How long is execution time?

```
if (thread-id % 2 == 0) {
        :    // (A) 30msec
} else {
        :    // (B) 20msec
}
```
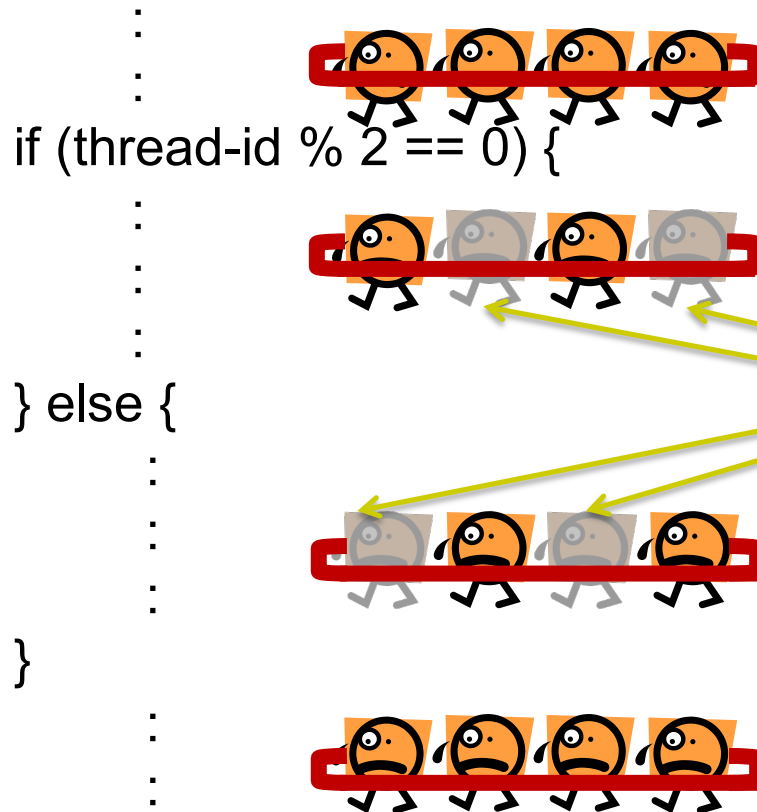
On CPU (OpenMP)

(A) (B) (A) (B)

30ms

On GPU, threads in a warp must execute the same instruction. What happens?

Branches on GPU (1)

if (thread-id % 2 == 0) {

} else {

}

Some threads are made sleep
Both "then" and "else" are executed!

→ Answer to previous question is 50ms !

※ Similar cases happen in for, while…

Branches on GPU (2)

- As exceptional cases, if threads in a warp "agree" in branch condition, either "then" part or "else" part is executed → Efficient!

- If there is difference of opinion (previous page), it is called a divergent branch

→ Agreement among buddies (threads in a warp) is important for speed

Considering Data Transfer Costs

CPU   GPU
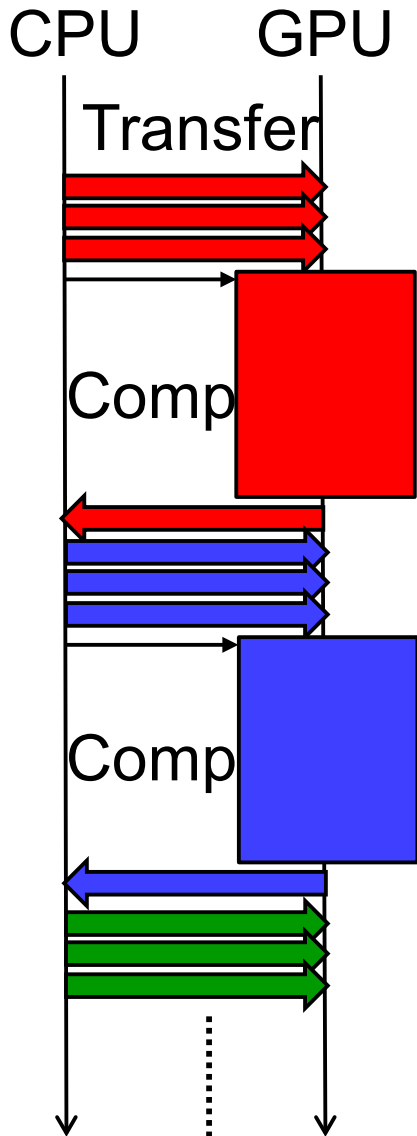
Transfer

Comp

Comp

Example case: We are going to compute multiply for different matrices

- Input data are on host memory
  - $C1 = A1 \times B1$
  - $C2 = A2 \times B2$
  - ….
  - $Cn = An \times Bn$

- In default, GPU cannot compute during transfer
→ cudaStream is useful for hiding transfer costs

This is also useful for speed-up of mm-cuda, by dividing matrices into pieces
→ Samples are at
/gs/hs1/tga-ppcomp/22/mm-str-cuda

# Asynchronous Executions with cudaStream (1)

What are streams?

- GPU's "service counters" that accept tasks from CPU
  - Each stream looks like a queue
- "Tasks" from CPU to GPU include
  - Data transfer (Host → Device)
  - GPU kernel function call
  - Data transfer (Device → Host)

"Tasks" here are a bit different from "omp task"

Ask GPU something (tasks)

stream

CPU

GPU
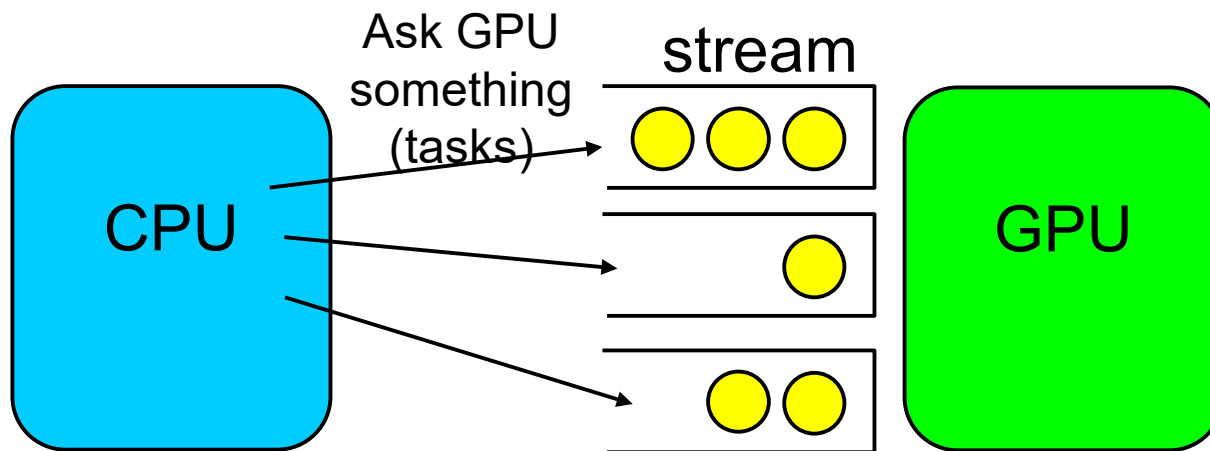
All of sample programs are using the "default stream"

Asynchronous Executions with cudaStream (2)
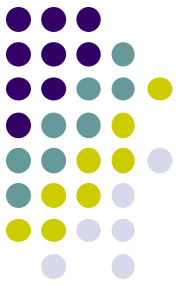
Create a stream

```
cudaStream_t str;
cudaStreamCreate(&str); // Create a stream
```

Data transfer using a specific stream

```
cudaMemcpyAsync(dst, src, size, type, str);
```

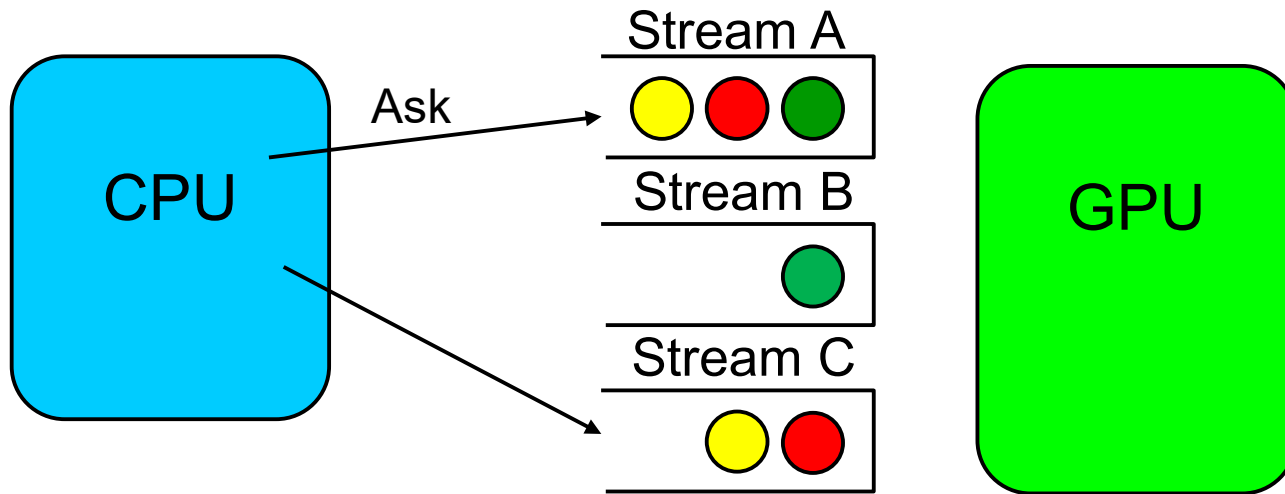Call GPU kernel function using a stream

```
func<<<gs, bs, 0, str>>>( … );
// 3rd parameter is related to for "shared memory"
```

Wait until all tasks on a stream are finished

```
cudaStreamSynchronize(str);
```

How GPU Executes Tasks

Stream A

Ask

CPU

Stream B

Stream C

GPU

- Tasks on the same stream is done in FIFO
- If tasks are in different streams, and have different kinds, they may be done simultaneously
  - Kinds: H→D, kernel, D→H
  - Note: If tasks are in the same kind, no speed up

Speed Up with Overlap of Computation and Transfer

n streams can be used for n independent tasks

- C1 = A1 × B1 (includes H->D, Calc, D->H)
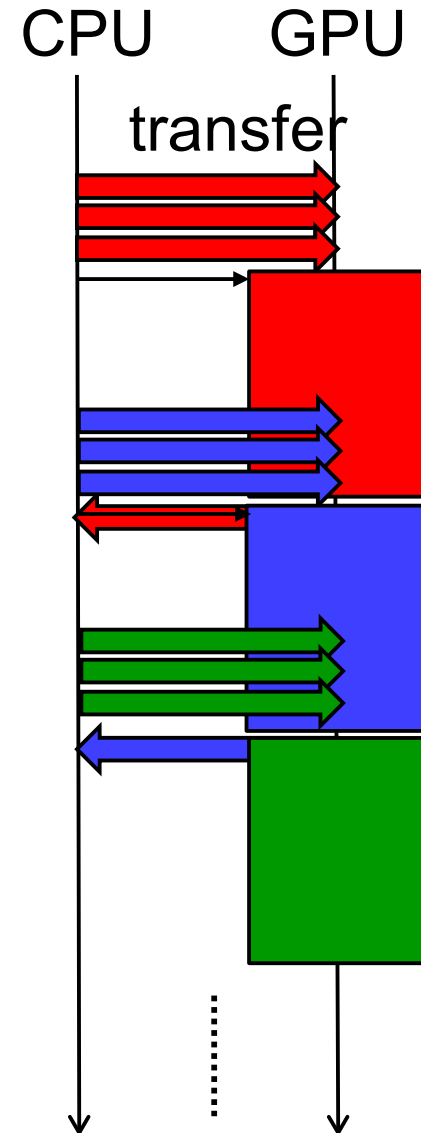- C2 = A2 × B2

  ....

- Cn = An × Bn
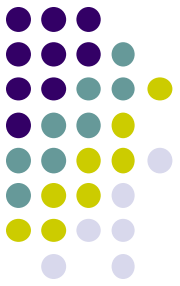
→ We will see speed up since

(Total comp time + Total trans time)

is improved to

max(Total comp time, Total trans time)

This is not a unique solution;
Use 2 or 3 streams repeatedly → we can save memory and stream resources

cudaMallocHost() is used instead of malloc()
This speeds up cudaMemcpyAsync()

CPU   GPU

transfer

More Things to Study

- Using CUDA shared memory
  - fast and small memory than device memory

- Unified memory in recent CUDA
  - cudaMemcpy can be omitted for automatic data transfer

- Using Tensor-core to accelerate deep learning
  - Only on V100 GPUs or later
  - Unfortunately, TSUBAME3 has older P100 ☹

- Using multiple GPUs towards petascale computation
  - MPI+CUDA, MPI+OpenACC

- More and more…

Assignments in GPU Part (Abstract)

Choose <u>one of</u> [G1]—[G3], and submit a report

Due date: May 26 (Thursday)

[G1] Parallelize "diffusion" sample program by OpenACC or CUDA

[G2] Evaluate speed of "mm-acc" or "mm-cuda" in detail

[G3] (Freestyle) Parallelize *any* program by OpenACC or CUDA.

Next Class:

- Part 3: MPI Programming (1)
  - Introduction to distributed memory parallel programming