

# Practical Parallel Computing (実践的並列コンピューティング)

Part3: MPI (1)  
May 19, 2022

Toshio Endo  
School of Computing & GSIC  
[endo@is.titech.ac.jp](mailto:endo@is.titech.ac.jp)

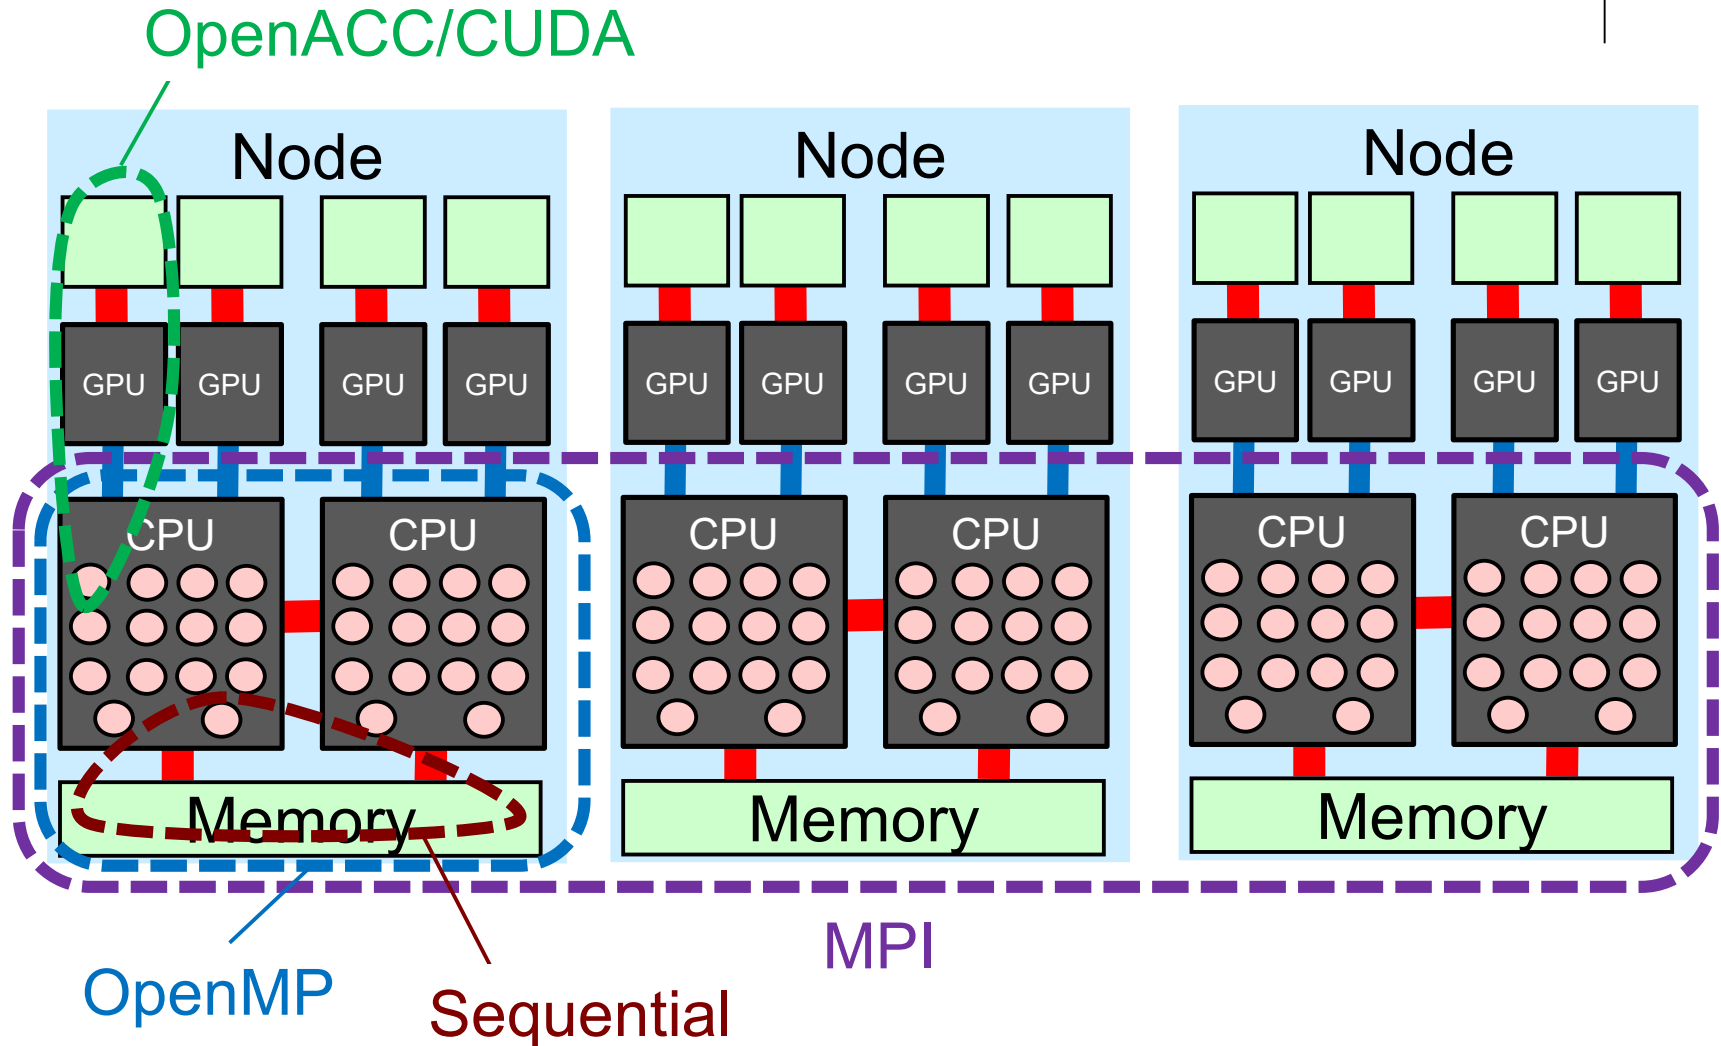


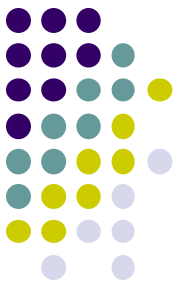


# Overview of This Course

- Part 0: Introduction
  - 2 classes
- Part 1: OpenMP for shared memory programming
  - 4 classes
- Part 2: GPU programming
  - 4 classes ← We are here (1/4)
  - OpenACC (1.5 classes) and CUDA (2.5 classes)
- Part 3: **MPI** for distributed memory programming
  - 4 classes ← We are here (1/4)

# Parallel Programming Methods on TSUBAME

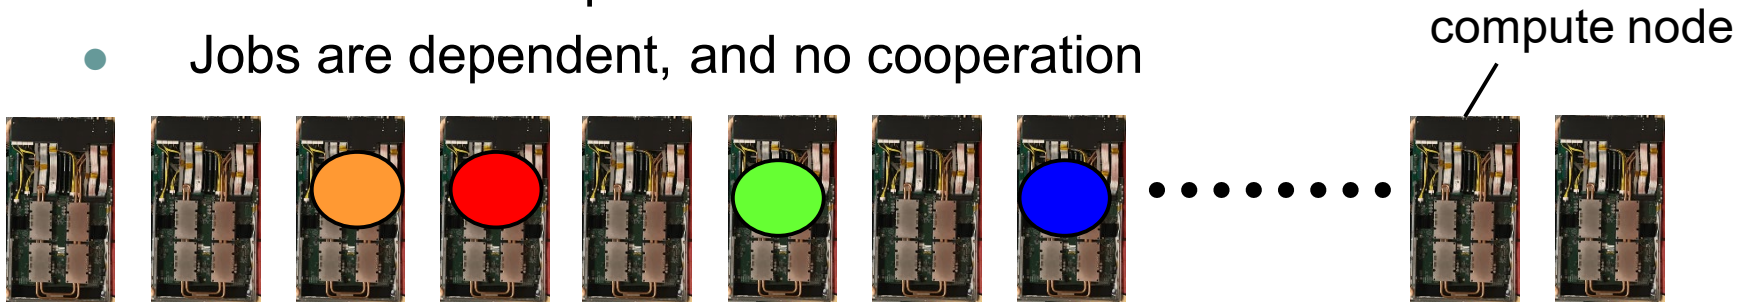




# How Multiple Nodes are Used

## 1. Submit several jobs into job scheduler

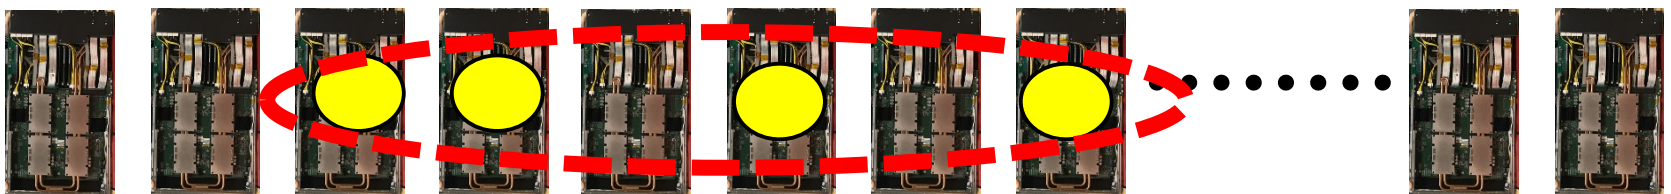
- cf) Program executions with different parameters → Parameter Sweep
- Jobs are dependent, and no cooperation



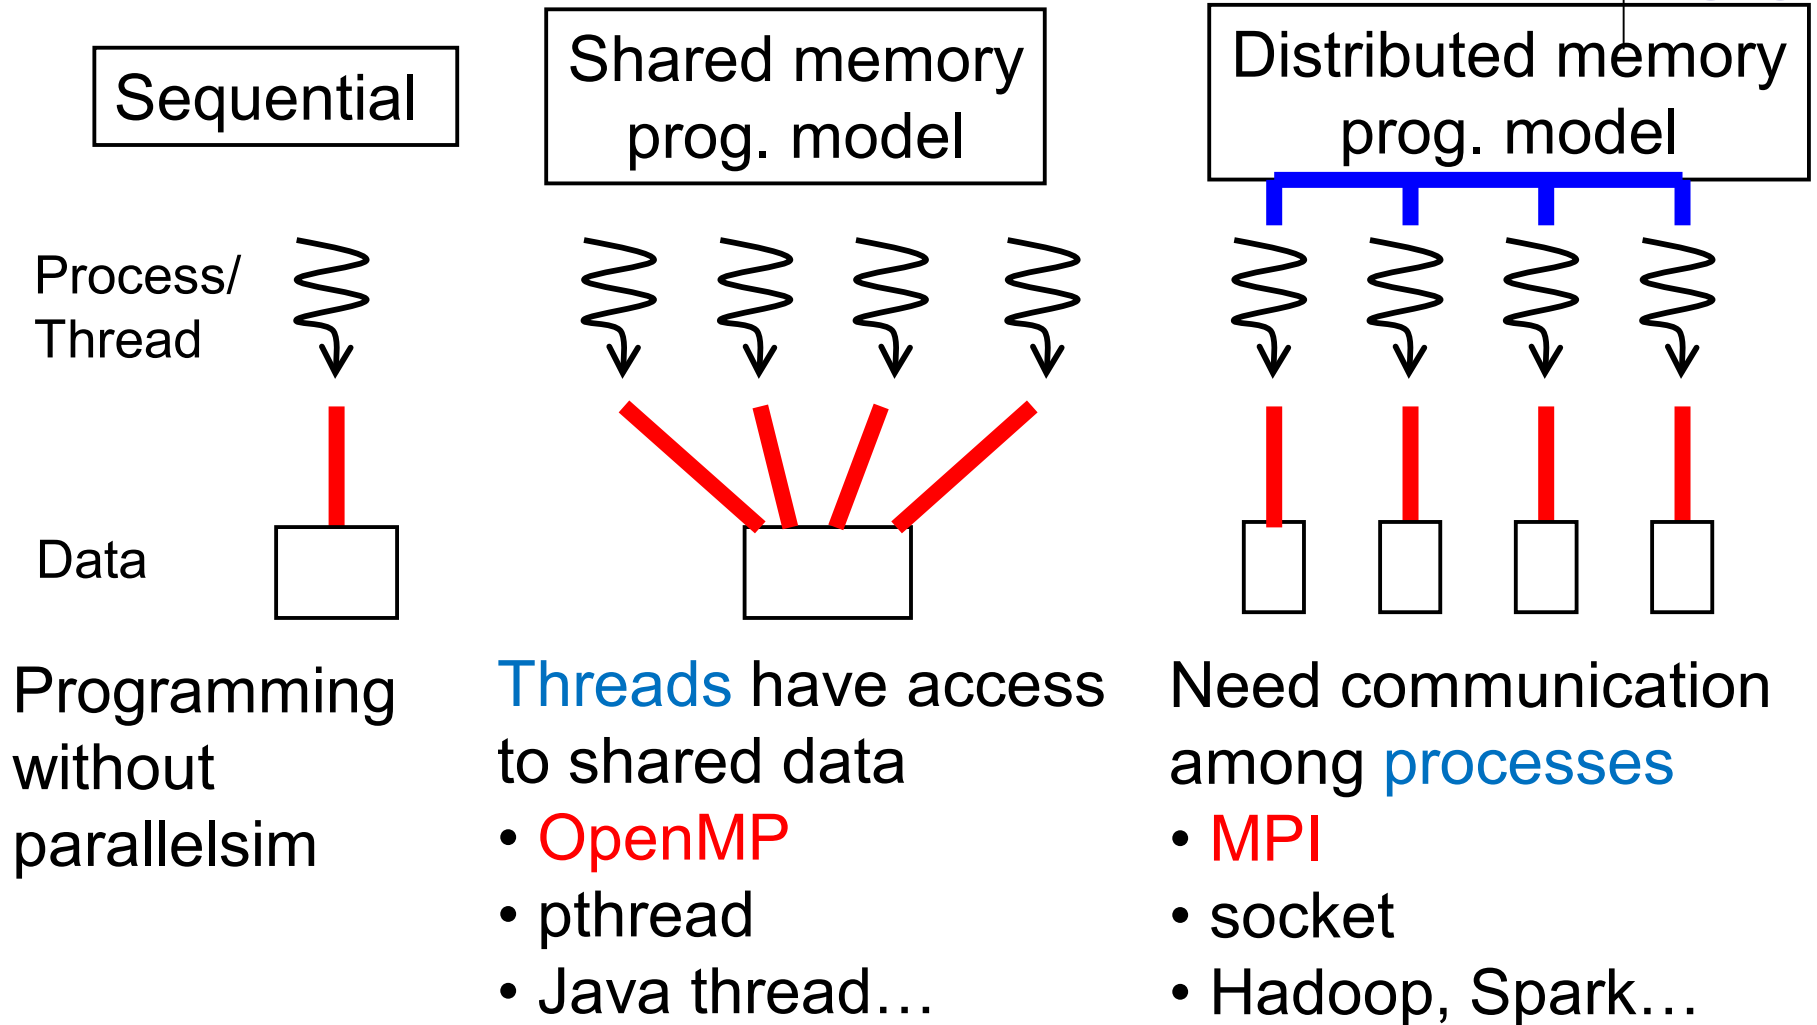
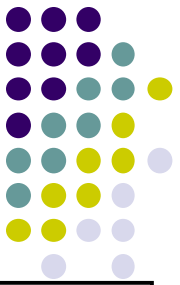
## 2. Use distributed memory programming → A single job can use multiple nodes

- Socket programming, Hadoop, Spark...
- And **MPI**

Part 3



# Classification of Parallel Programming Models



# MPI (message-passing interface)



- Parallel programming interface based on distributed memory model
- Used by C, C++, Fortran programs
  - Programs call MPI library functions, for **message passing** etc.
- There are several MPI libraries
  - OpenMPI (default)      ← OpenMPI ≠ OpenMP ☹️
  - Intel MPI, SGI MPE, MVAPICH, MPICH...

# Differences from OpenMP



In MPI,

- An execution consists of multiple **processes** (not threads)
  - We can use multiple nodes 😊
  - The number of running processes is basically constant
- No variables are shared. Instead **message passing** is used
  - Data distribution has to be programmed
- No smart syntaxes such as “omp for” or “omp task” 😞
  - Task distribution has to be programmed 😞



# First MPI Sample

- [/gs/hs1/tga-ppcomp/22/hello-mpi](#)

*[make sure that you are at a interactive node (r7i7nX) ]*

**module load cuda openmpi** *[Do once after login]*

**cd ~/t3workspace** *[In web-only route]*

**cp -r /gs/hs1/tga-ppcomp/22/hello-mpi .**

**cd hello-mpi**

**make**

*[An executable file “hello” is created]*

**mpixec -n 7 ./hello**

Number of  
processes

Name of program  
(using options are ok)



# Compiling and Executing MPI Programs



Case of OpenMPI library on TSUBAME3.0

- To compile
  - `module load cuda openmpi`, and then use `mpicc`
  - For sample programs, “make” command works
- To execute
  - `mpiexec -n 7 ./hello`

Required for module dependency☹

Number of processes

↑ These methods uses 1 (current) node.

# When We Want to Use More Processes than CPU Cores



- An interactive node has 7 CPU cores (= 14 Hyperthreads)
- If number of processes is larger than CPU cores, we need “--oversubscribe” option of OpenMPI
- `mpiexec -n 7 ./hello`
- `mpiexec -n 14 --oversubscribe ./hello`

Needed when we use  
processes > physical cores  
(7, in an interactive node)

# From Single Node To Multiple Nodes



- So far, we used a single node
  - An interactive node with 7 cores = 14 Hyperthreads
  - For reports, a single node is enough
- MPI can support multiple nodes!

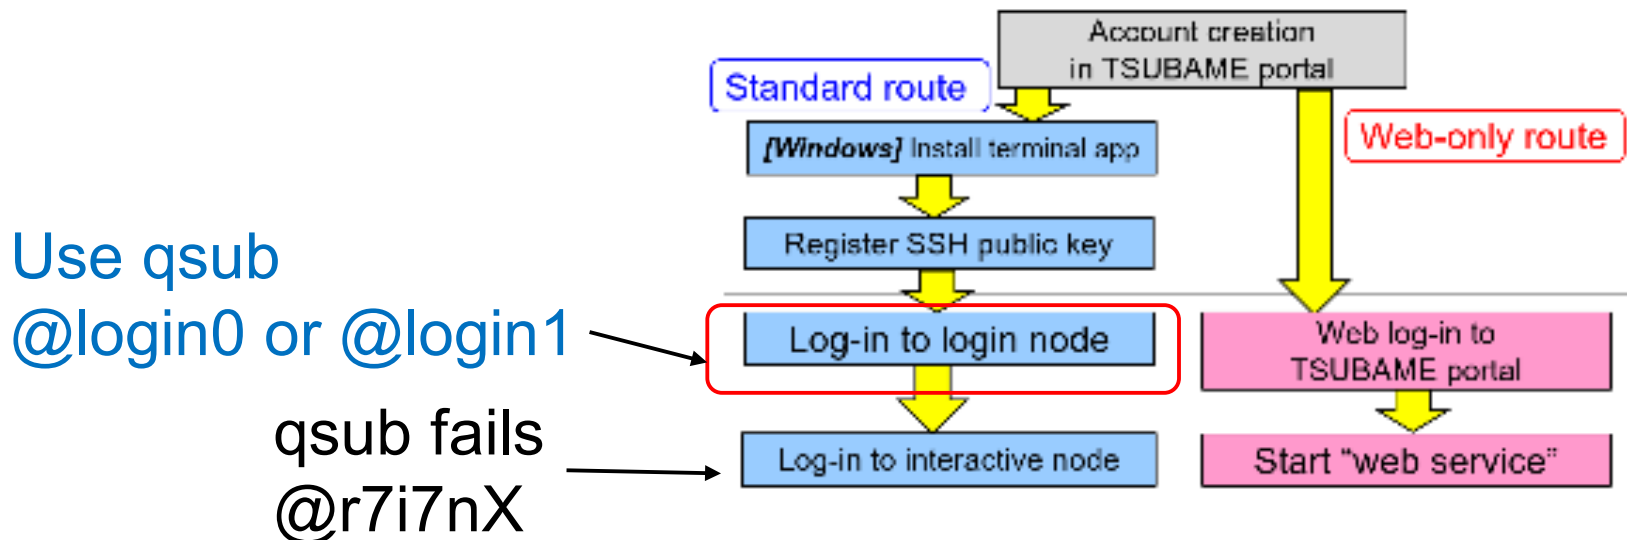
To use multiple nodes on TSUBAME3.0,  
we need “job submission” !

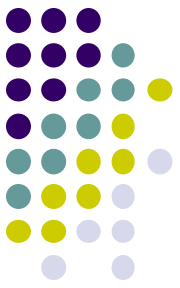


# Notes on Job Submission in “Standard route”



- On an interactive node via “standard route”, qsub/qstat commands may not be found
- Please use qsub/qstat **on a login node**
  - (out of iqrsh)





# Submit an MPI Job (case of OpenMPI)

- We are going to execute a job using  
4 processes  $\times$  2 nodes = 8 processes

(1) Make a script file: `job.sh`

```
#!/bin/sh
#$ -cwd
#$ -l q_core=2
#$ -l h_rt=00:10:00

. /etc/profile.d/modules.sh
module load cuda openmpi

mpixexec -n 8 -npnode 4 ./hello
```

4core node x 2

Module preparation

Number of  
processes

Number of  
processes  
per node

Program name  
(and option)

(2) Submit the job with “`qsub`”

`qsub job.sh`

↑  $\leq 0:10:00$  and  $\leq 2$ nodes

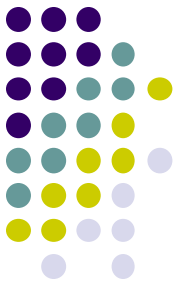
`qsub -g tga-ppcomp job.sh`

↑  $> 0:10:00$  or  $\geq 3$ nodes

Be careful for TSUBAME point

# Notes in This Lecture

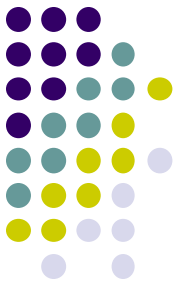
(Also see OpenMP(4) slides)



- Try executions without TSUBAME points first
- まず無料利用を考慮してください
  - $h_{rt} \leq 0:10:00$
- If necessary for reports, you can use up to 36,000 points in total per student
- 本講義のレポートの作成に必要な場合のみ、一人あたり合計で36,000ポイントまで利用を認めます
  - $36,000 \text{ points} \div F_{\text{node}} \times 10 \text{ hours}$
  - You can check point consumption on TSUBAME portal
- The TSUBAME group name is [tga-ppcomp](#)

Users need to follow the rules at [www.t3.gsic.titech.ac.jp](http://www.t3.gsic.titech.ac.jp)

利用時には [www.t3.gsic.titech.ac.jp](http://www.t3.gsic.titech.ac.jp) に示される規則を守る必要があります<sub>4</sub>



# Resource Types on TSUBAME3.0

- Choose one of resource types (number of cores, mainly)
  - It is like “instance types” in cloud systems
  - Please specify “proper” one

Resource type name	Physical CPU cores	Memory (GB)	GPUs
f_node	28	240	4
h_node	14	120	2
q_node	7	60	1
q_core	4	30	0
s_core	1	7.5	0
s_gpu	2	15	1

← Same as an interactive node

`#$ -l [resource_type] = [Number]`

Example:

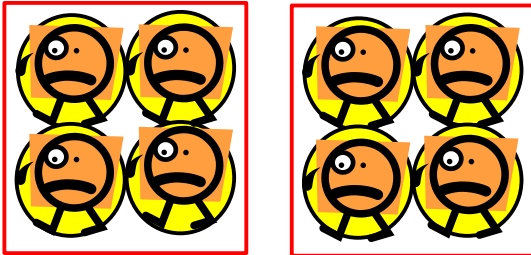
`#$ -l s_core=1` ← The minimum resource allocation

# Nodes, Cores, MPI Processes



```

:
#$ -l q_core=2
:
mpirun -n 8 -npnode 4
...
    
```

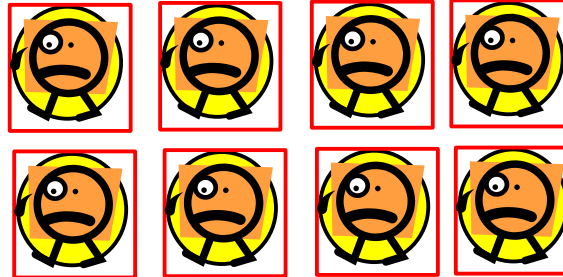


2 (virtual) nodes are prepared  
Each node has 4 cores (q\_core)

4 processes are created per node. Totally 8 are created  
→ 2 nodes are used

```

:
#$ -l s_core=8
:
mpirun -n 8 -npnode 1
...
    
```

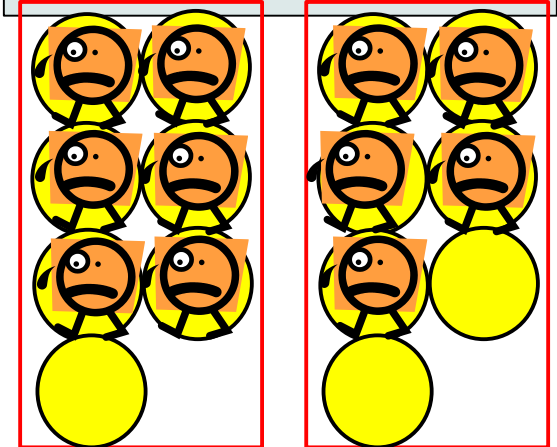


8 (virtual) nodes are prepared  
Each node has 1 cores (s\_core)

1 processes are created per node. Totally 8 are created  
→ 8 nodes are used

```

:
#$ -l q_node=2
:
mpirun -n 11 -npnode 6
...
    
```



2 (virtual) nodes are prepared  
Each node has 7 cores (q\_node)

6 processes are created per node. Totally 11 are created  
→ 2 nodes are used  
(There are idle cores)



# An MPI Program Looks Like



```
#include <stdio.h>
```

```
#include <mpi.h>
```

```
int main(int argc, char *argv[])
```

```
{
```

```
    MPI_Init(&argc, &argv); ← Initialize MPI
```

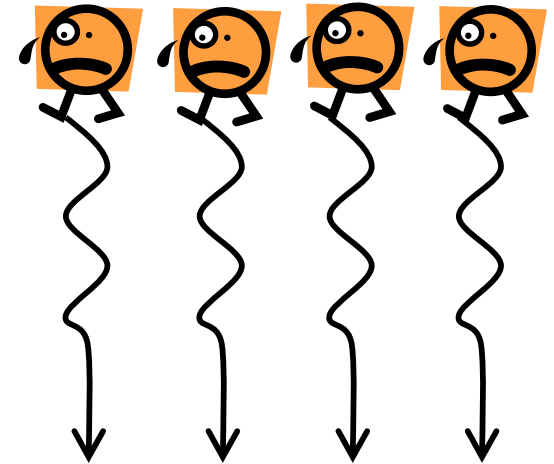
```
    (Computation/communication)
```

```
    MPI_Finalize();
```

```
    ← Finalize MPI
```

```
}
```

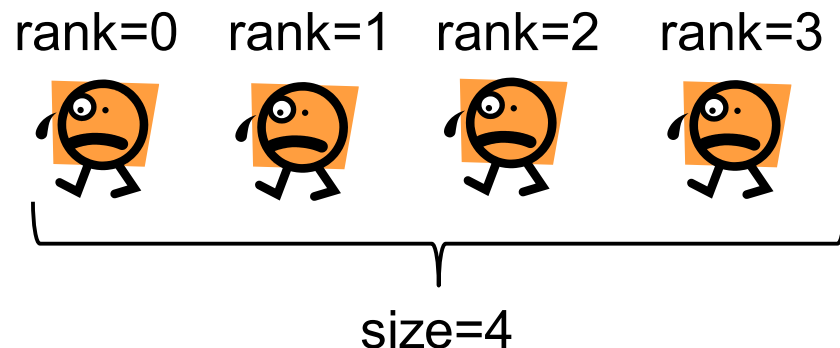
If number of  
processes=4



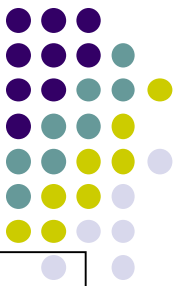


# ID of Each MPI Process

- Each process has its ID (0, 1, 2...), called **rank**
  - `MPI_Comm_rank(MPI_COMM_WORLD, &rank);`  
→ Get its rank
  - `MPI_Comm_size(MPI_COMM_WORLD, &size);`  
→ Get the number of total processes
  - $0 \leq \text{rank} < \text{size}$
  - The rank is used as target of message passing



# “mm” sample: Matrix Multiply



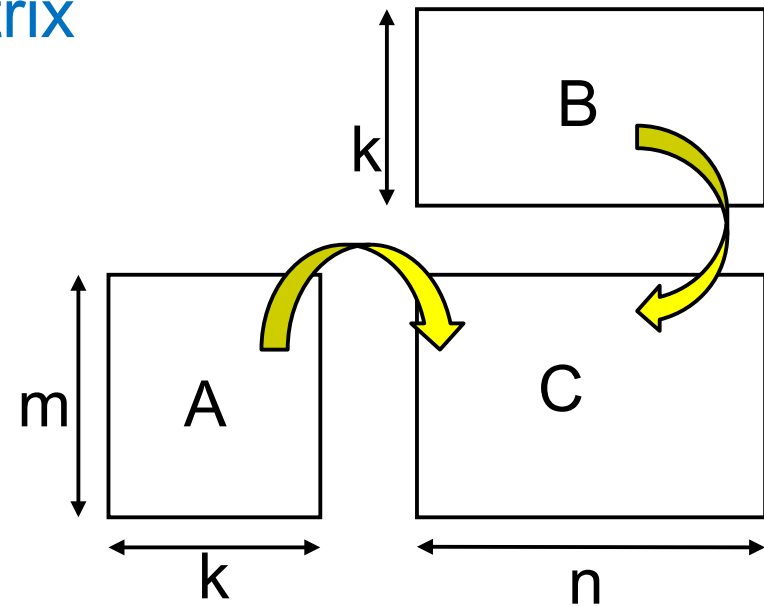
MPI version available at </gs/hs1/tga-ppcomp/22/mm-mpi/>

A: a  $(m \times k)$  matrix, B: a  $(k \times n)$  matrix

C: a  $(m \times n)$  matrix

$$C \leftarrow A \times B$$

- Algorithm with a triple for loop
- Supports variable matrix size.
  - Each matrix is expressed as a 1D array by *column-major* format



Execution:

`mpiexec -n [np] ./mm [m] [n] [k]` (interactive)

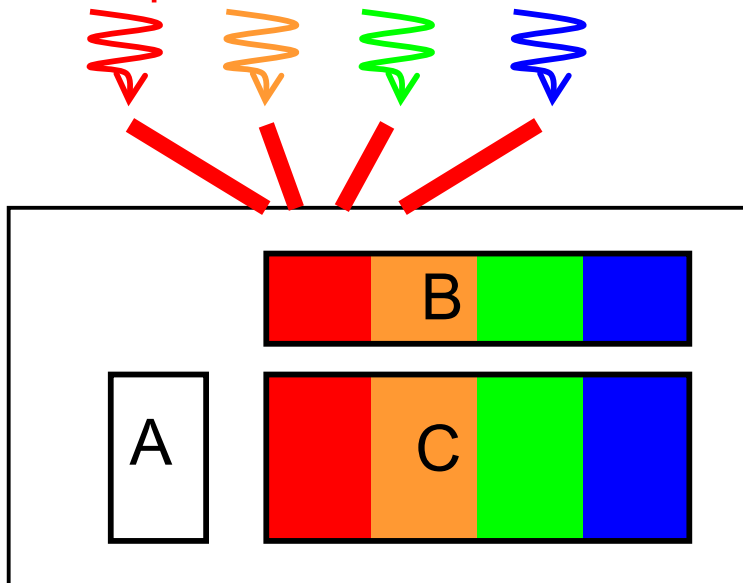
`mpiexec -n [np] -npernode [nn] ./mm [m] [n] [k]` (in job script)

# Why Distributed Programming is More Difficult (case of mm-mpi)



Shared memory with OpenMP:

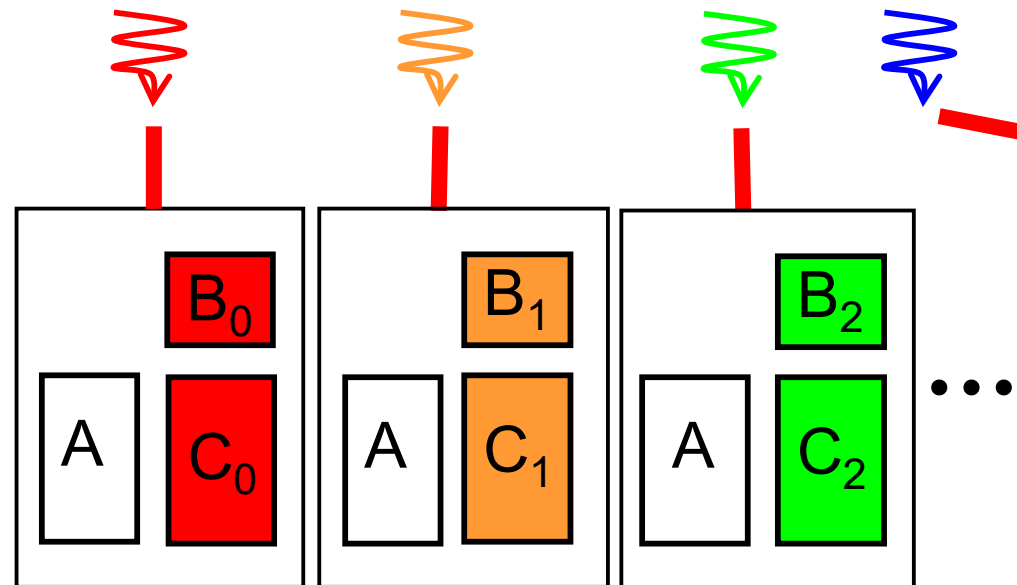
Programmers consider how **computations** are divided



In this case, matrix A is accessed by all threads  
→ Programmers **do not have to know** that

Distributed memory with MPI:

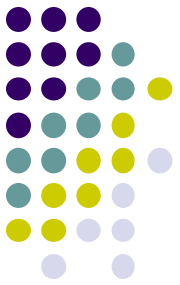
Programmers consider how **data and computations** are divided



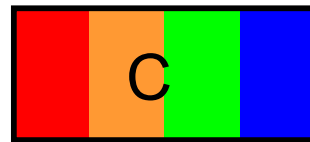
Programmers **have to design** which data is accessed by each process

# Programming Data Distribution

(case of mm-mpi)



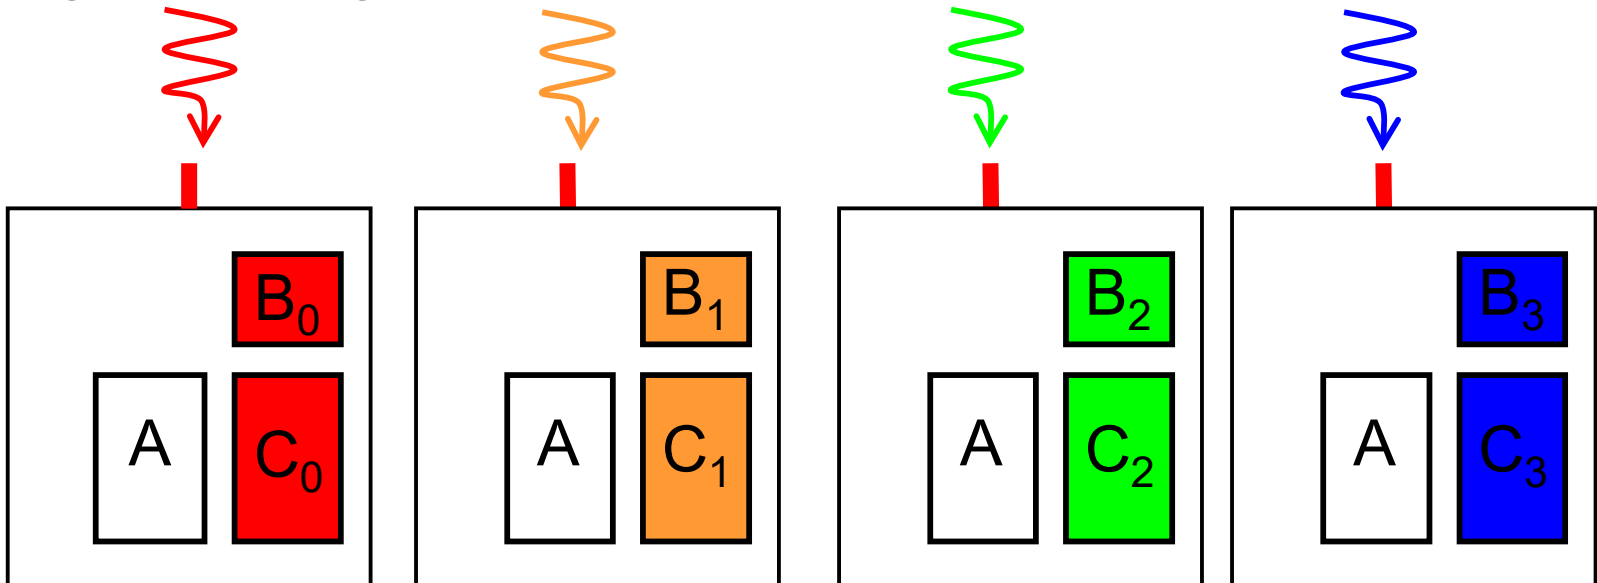
Design distribution method:



I will divide B, C vertically.

I will put replicas of A on every process...

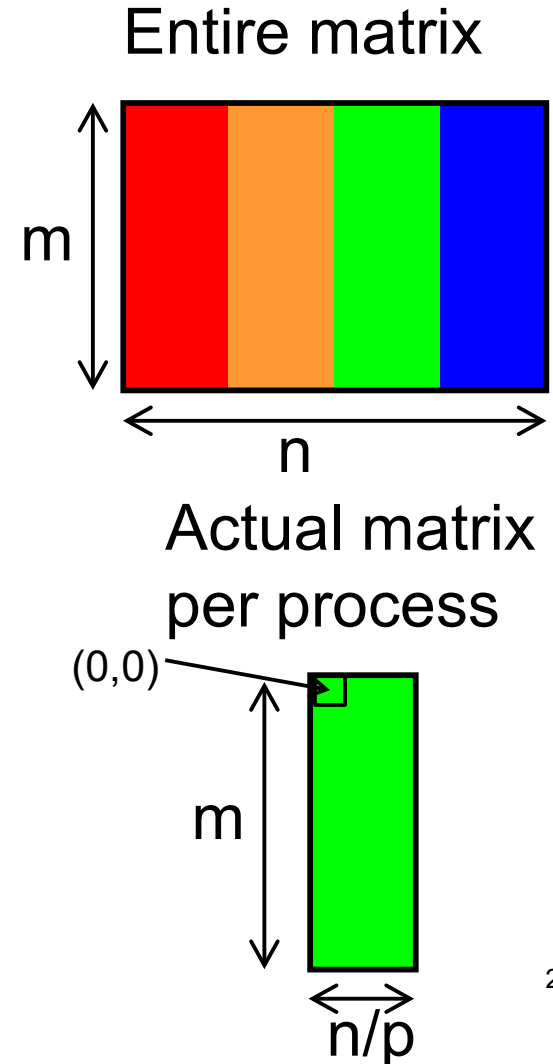
Programming actual location:



# Programming Actual Data Distribution



- We want to distribute a  $m \times n$  matrix among  $p$  processes
  - We assume  $n$  is divisible by  $p$
- Each process has a partial matrix of size  $m \times (n/p)$ 
  - We need to “malloc”  
 $m \times (n/p) \times \text{sizeof}(\text{data-type})$  size
  - We need to be aware of relation between partial matrix and entire matrix
    - $(i, j)$  element in partial matrix owned by Process  $r \Leftrightarrow (i, n/p \times r + j)$  element in entire matrix

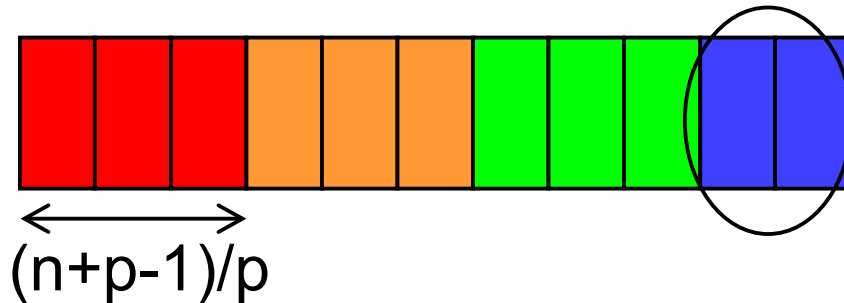




# Considering Indivisible Cases

- What if data size  $n$  is indivisible by  $p$ ?
  - We let  $n=11$ ,  $p=4$ 
    - How many data each process take?
    - $n/p = 2$  is not good (C division uses round down). Instead, we should use round up division
- $(n+p-1)/p = 3$  works well

Note that the “final” process takes less than others



See `divide_length()` function in `mm-mpi/mm.c`  
It calculates the range the process should take  
→ Outputs are **first index  $s$**  and **last index  $e$**



# Notes in Time Measurement

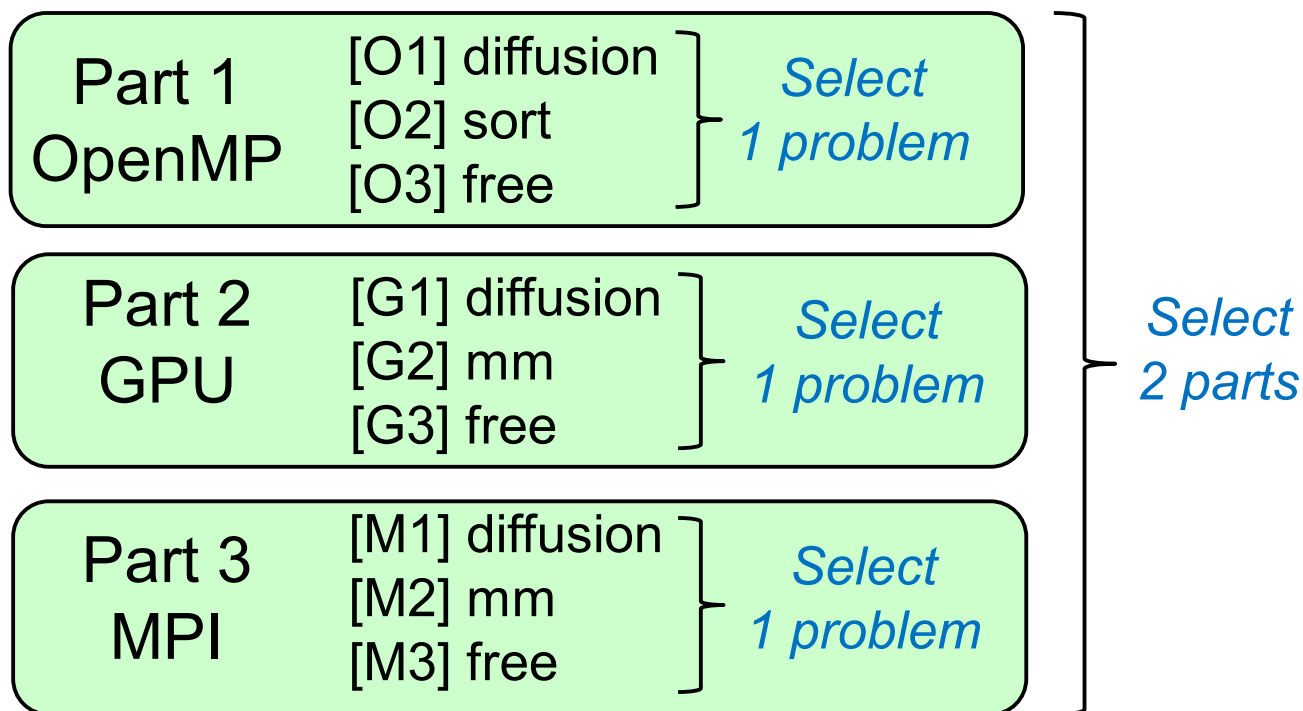
- In mm-mpi, gettimeofday() is used for time measurement
- For accurate measurement, we should call **MPI\_Barrier(MPI\_COMM\_WORLD)** before measurement
  - This synchronizes all processes
  - All processes need to call this



# Assignments in this Course



- There is homework for each part. Submissions of reports for **2 parts** are required





# Assignments in MPI Part (1)

Choose one of [M1]—[M3], and submit a report

Due date: June 9 (Thursday)

[M1] Parallelize “diffusion” sample program by MPI.

- Do not forget to change Makefile and job.sh appropriately
- Use deadlock-free communication
  - see `neicomm_safe()` in `neicomm-mpi` sample

Optional:

- To make array sizes (NX, NY) variable parameters
- To consider the case with NY is indivisible by p
  - see `divide_length()` in `mm_mpi` sample
- To improve performance further. Blocking, 2D division, etc

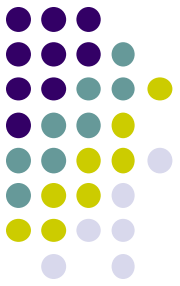


# Assignments in MPI Part(2)

[M2] Improve “mm-mpi” sample in order to reduce memory consumption

Optional:

- To consider indivisible cases
- To try advanced algorithms, such as SUMMA
  - the paper “*SUMMA: Scalable Universal Matrix Multiplication Algorithm*” by Van de Geijn
  - <http://www.netlib.org/lapack/lawnspdf/lawn96.pdf>



# Assignments in MPI Part (3)

[M3] (Freestyle) Parallelize *any* program by MPI.

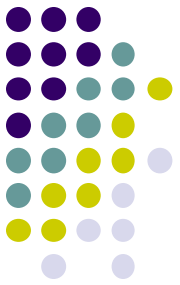
- cf) A problem related to your research
- More challenging one for parallelization is better
  - cf) Partial computations have dependency with each other

# Notes in Report Submission (1)



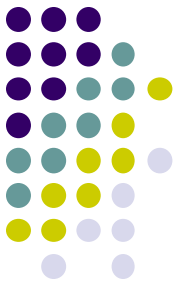
- Submit the followings via **T2SCHOLA**
  - (1) **A report document**
    - PDF, MS-Word or text file
    - 2 pages or more
    - in English or Japanese (日本語もok)
  - (2) **Source code files** of your program
    - Try “zip” to submit multiple files

# Notes in Report Submission (2)



The report document should include:

- Which problem you have chosen
- How you parallelized
  - It is even better if you mention efforts for high performance or new functions
- Performance evaluation on TSUBAME
  - With varying number of processes
    - On an interactive node,
      - $1 \leq P \leq 14$  ( $P$ =number of processes)
      - If  $P > 7$ , use “`mpiexec -n [P] --oversubscribe ...`”
    - To use more CPU cores, you need to do “job submission” (optional)
  - With varying problem sizes
  - Discussion with your findings
  - Other machines than TSUBAME are ok, if available



# Next Class

- MPI (2)
  - Basic message passing
  - How to parallelize diffusion sample with MPI
    - Related to [M1]