

Practical Parallel Computing (実践的並列コンピューティング)

Part3: MPI (4)
May 30, 2021

Toshio Endo
School of Computing & GSIC
endo@is.titech.ac.jp







Improving MPI diffusion by Overlapping of Communication

related to [M1], but optional

Overview of MPI “diffusion”

(See MPI (2) Slides)



```
for (t = 0; t < nt; t++) {
```

```
    if (rank > 0) Send B to rank-1
```

```
    if (rank < size-1) Send D to rank+1
```

```
    if (rank > 0) Recv A from rank-1
```

```
    if (rank < size-1) Recv E from rank+1
```

(1) Communication
in “old” array

```
    Computes points in rows B-D
```

```
    Switch old and new arrays
```

(2) Computation
“old” array \Rightarrow “new” array

```
}
```


Actually this should be fixed to avoid deadlock

Considering Performance of MPI Programs



(Simplified) Execution time of an MPI program =

- Computation time ← including memory access
- + Communication time ← including congestion
- + Others ← load imbalance, I/O...



Computation Time & Communication Time (1)



How are they determined? (very simplified discussion)

1. Aspect of software

Computation time

- Longer if computation costs are larger
 - $O(mnk/p)$ in matmul,
 - $O(NX NY NT/p)$ in diffusion

per process

Communication time

- Longer if communication costs are larger
 - $O(mk)$ in memory reduced matmul
 - $O(NX NT)$ in diffusion

per process

Computation Time & Communication Time (2)




2. Aspect of hardware

Computation time

- Shorter if processor speed is faster
 - 850GFlops per node on TSUBAME3 (CPU only)

Communication time

- Shorter if network speed is faster
 - 400Gbps per node on TSUBAME3



Speed of actual software is slower than the “peak” performance

Parameters for Network Speed



What parameters describes network speed?

- **Bandwidth**: Data amounts that network can transport per unit time → **Larger is better**
 - bps: X bits per second
 - B/s: X Bytes per second
 - On TSUBAME3, 400Gbps = 50GB/s per node
- **Network latency**: Time to transport minimum data (1bit, for example) → **Smaller is better**
 - On TSUBAME3, a few us


✂ Additionally, communication time may suffer from effects of **network topology**: how nodes/switches are connected to each other

Bandwidth and Latency



Is “latency” reciprocal of “bandwidth”?

→ No, because data are transported in “pipe-lined” style



$$T = M / B + L$$

T: Communication time

M: Data size

B: Bandwidth

L: Network latency



※ Be aware of difference between “Byte” and “bit”: 1Byte=8bit

※ In some contexts, T, not L, may be called “latency”

Idea of Overlapping



If “some computations” do not require contents of message, we may start them beforehand



T_P is divided into


- T_{P1} : can be overlapped
- T_{P2} : cannot be overlapped

$$T = \max(T_N, T_{P1}) + T_{P2}$$

Overlapping in Stencil Computation (related to [M1], but not required)



When we consider data dependency in detail, we can find computations that do not need data from other processes



Rows C, D, E do not need data from other processes
→ They can be computed without waiting for finishing communication


On the other hand, rows B, F need received data

For such purposes, non-blocking communications (MPI_Isend, MPI_Irecv...) are helpful again

Implementation without Overlapping (Not Fast!)



```
for (t = 0; t < nt; t++) {  
    Start Send B to rank-1, Start Send F to rank+1  
    (MPI_Isend)  
    Start Recv A from rank-1, Start Recv G from  
    rank+1 (MPI_Irecv)  
    Waits for finishing all communications  
    (MPI_Wait for 4 times)  
    Compute rows B--F  
    Switch old and new arrays  
}
```



$$T = T_N + T_P$$

Implementation with Overlapping



```
for (t = 0; t < nt; t++) {  
    Start Send B to rank-1, Start Send F to rank+1  
    (MPI_Isend)  
    Start Recv A from rank-1, Start Recv G from  
    rank-1 (MPI_Irecv)  
    Compute rows C--E  
    Waits for finishing all communications  
    (MPI_Wait)  
    Compute rows B, F  
    Switch old and new arrays  
}
```



computations are
divided

$$T = \max(T_N, T_{P1}) + T_{P2} < T_N + T_{P1} + T_{P2} = T_N + T_P$$

Another Improvement: Reducing Communication Amounts



Multi-dimensional division may reduce communication



Each process communicate with
upper/lower/right/left processes


- **Comp**: $O(NY \ NX/p)$
 - **Comm**: $O(NX)$
- per 1 process, 1 iteration

- **Comp**: $O(NY \ NX/p)$
 - **Comm**: $O((NY+NX)/p^{1/2})$
- per 1 process, 1 iteration
→ Comm is reduced

Multi-dimensional division and Non-contiguous data (1)



- MD division may need communication of non-contiguous data



In Row-major format, we need send/recv of non-contiguous data for left/right borders

But “fragmented communication” degrades performance! (since Latency > 0)
How do we do?

Multi-dimensional division and Non-contiguous data (2)



Solution (1):

- Before sending, copy non-contiguous data into another contiguous buffer
- After receiving, copy contiguous buffer to non-contiguous area

Solution (2):


- Use MPI_Datatype
 - Skipped in the class; you may use Google



Performance of Collective Communication


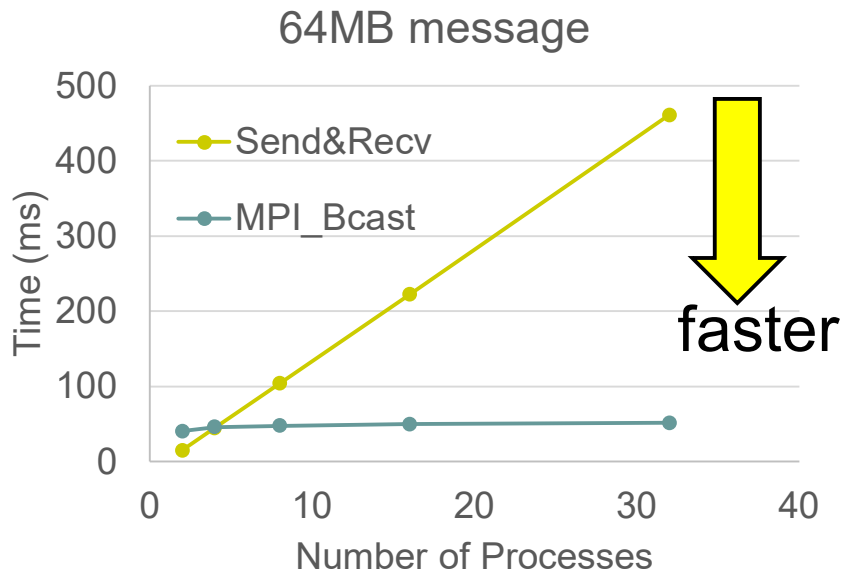
also in MPI (3) slides

It is Better to Use Collective Communications if Appropriate



measured
on TSUBAME2

- Comparing MPI_Bcast and MPI_Send&Recv
1 process per node is invoked (to measure network)
In the latter, rank 0 called MPI_Send for p-1 times to other processes



- In most cases, MPI_Bcast is faster


Why are Collective Communications Fast?






- Since MPI library uses scalable communication algorithms

- Case of **broadcast**:


Flat tree algorithm (slow)



Binomial tree algorithm



Scatter&Allgather
algorithm



(again)

Model of Communication Time





Illustration of peer-to-peer communication of data size M



$$T = M / B + L$$

T : Communication time

M : Data size

B : Bandwidth

L : Network latency

※ Be aware of difference between “Byte” and “bit”: 1Byte=8bit


※ Actually it is more complex for process's place, effects of network topology, congestion, packet size...

Cost Model of Broadcast Algorithms



- Case of “broadcast” of size M data
 - p : number of processes, B : network bandwidth, L : network latency


Flat tree algorithm



$$p(M/B + L)$$

→ *Slow*

Binomial tree algorithm



$$(\log p)(M/B + L)$$

※ difference between $p-1$ and p is ignored


Broadcast by Scatter&Allgather (1)



- Scatter&Allgather algorithm

- (1) The root process divide the message into p parts
- (2) Scatter
- (3) Allgather

R. Thakur and W. Gropp. Improving the performance of collective operations in mpich. EuroPVM/MPI conference, 2003.




Broadcast by Scatter&Allgather (2)



- Scatter&Allgather algorithm

- (1) The root process divide the message into p parts
- (2) Scatter
- (3) Allgather



Comparison of Broadcast Algorithms

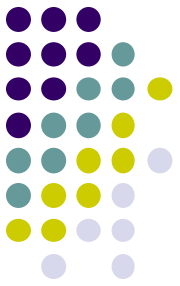


- Consider two extreme cases
 - If M is sufficiently large: $M/B+L \rightarrow M/B$
 - If M is close to zero: $M/B+L \rightarrow L$


	Flat Tree	Binomial Tree	Scatter&Allgather
Cost (General)	$p(M/B+L)$	$(\log p) (M/B+L)$	$2M/B + (p + \log p)L$
Cost with very large M	$p M/B$	$(\log p) M/B$	$2 M/B \rightarrow \text{Fastest}$
Cost with very small M	$p L$	$(\log p) L \rightarrow \text{Fastest}$	$(p + \log p) L$

Many MPI libraries implement multiple algorithms

They switch them automatically according to message size M 😊



Using Multiple GPUs with MPI+CUDA

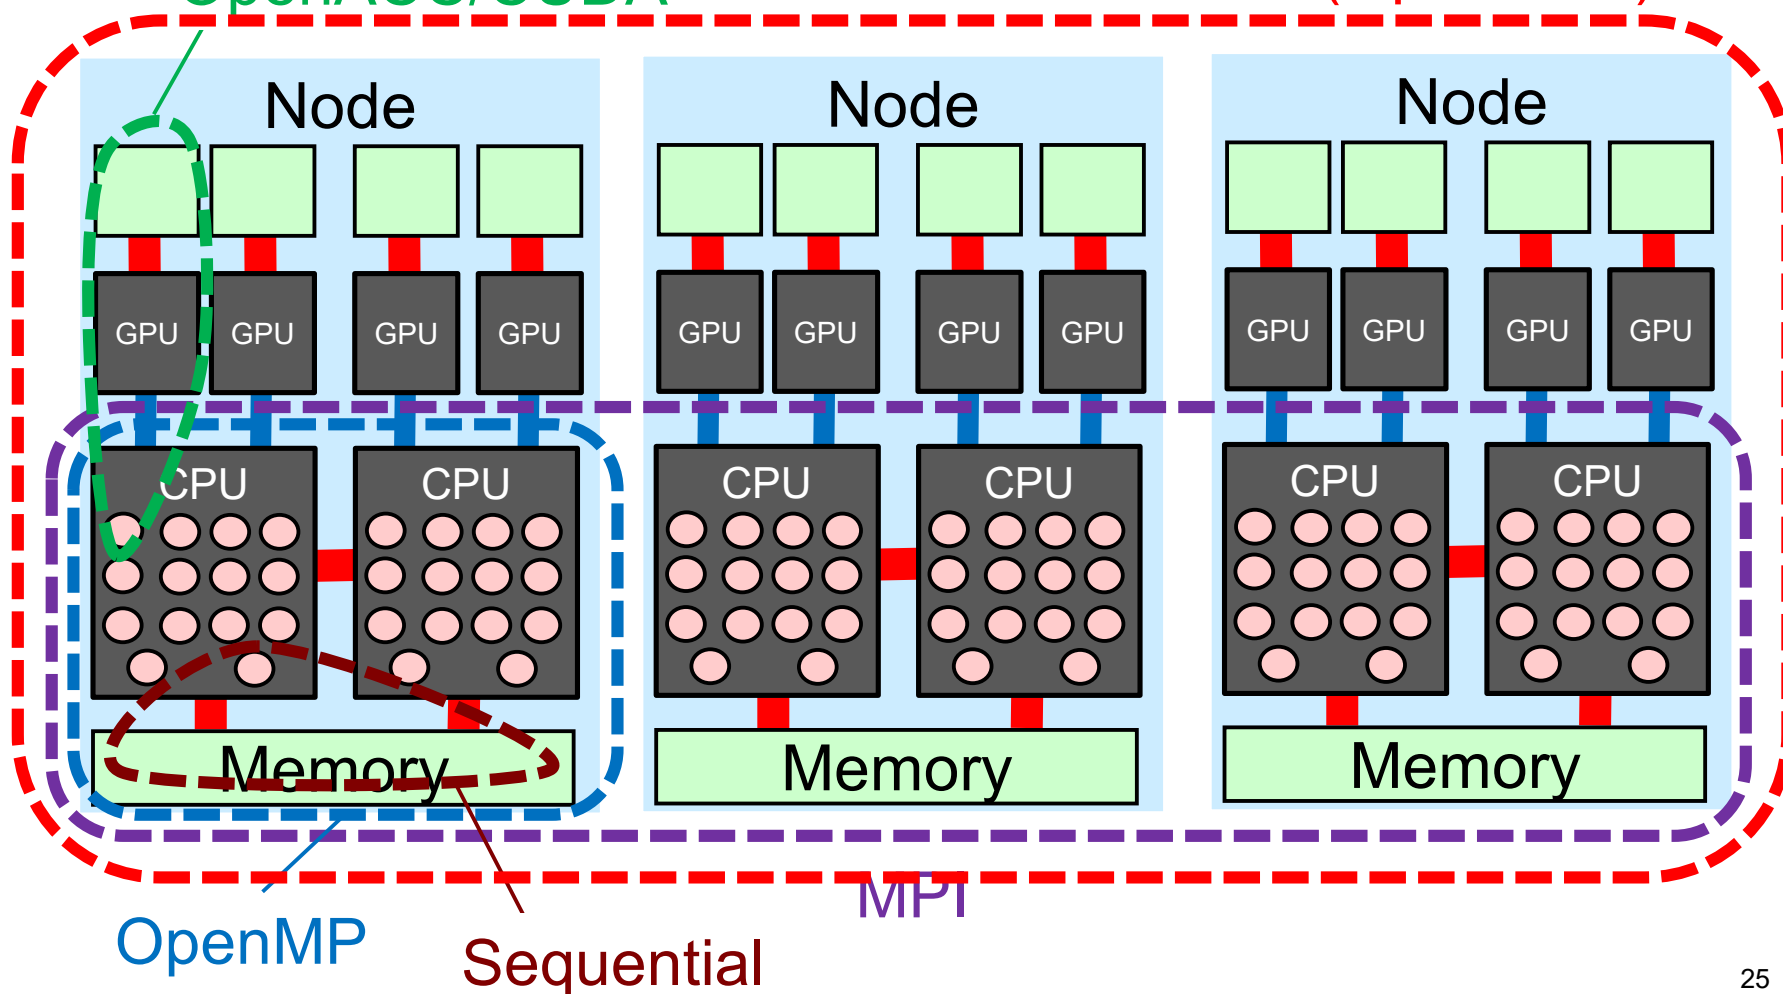


Parallel Programming Methods on TSUBAME



OpenACC/CUDA

MPI+CUDA (OpenACC)






Using Multiple GPUs

- GPUs on a single node
 - Up to 4 GPUs on TSUBAME3.0 f_node
 - OpenMP + CUDA
 - 1 thread uses 1 GPU
 - 1 thread switches multiple GPUs
 - `cudaSetDevice()` is called many times
- GPUs on multiple nodes
 - MPI + CUDA
 - 1 process uses 1 GPU ([mm-mpi-cuda](#) sample)



Using Multiple GPUs with MPI

- Basic idea:
 - (1) Start processes on multiple nodes by MPI
 - (2) Each process uses its local GPU by CUDA



Sample: [/gs/hs1/tga-ppcomp/22/mm-mpi-cuda/](https://github.com/tga-ppcomp/22/mm-mpi-cuda/)

Compiling mm-mpi-cuda Sample




```
module load cuda openmpi [Do once after login]  
cd ~/t3workspace [In web-only route]  
cp -r /gs/hs1/tga-ppcomp/22/mm-mpi-cuda .  
cd mm-mpi-cuda  
make  
[An executable file "mm" is created]
```

In this Makefile,

- nvcc is used as the compiler
- mpic++ is used as the linker, with CUDA libraries

Note: This may not work on other systems or future TSUBAME



Executing mm-mpi-cuda

- Interactive use is only for one node
- To use multiple nodes, **job submission** is required

In standard route,
Use qsub on the
login node

qsub job2q.sh → q_node (1GPU) x 2 are used → 2GPUs in total

qsub job2f.sh → f_node (4GPU) x 2 are used → 8 GPUs in total

job2f.sh

```
#!/bin/sh
#$ -cwd
#$ -l f_node=2
#$ -l h_rt=0:10:00

. /etc/profile.d/modules.sh
module load cuda openmpi

mpiexec -n 8 -npnode 4 -x LD_LIBRARY_PATH ./mm 2048 2048 2048
```

8 processes, 4 processes per node

To avoid scheduler's problem




Issues in Multiple GPUs per node

- `f_node` or `h_node` has multiple GPUs (4 or 2)
 - In default, all processes use “GPU 0” on the node → slow 😞
- Each process should use distinct GPUs
- ➔ In `mm.cu`, `cudaSetDevice(int dev)` is called first
 - specifies the GPU to be used
 - `dev`: GPU number in the node (0, 1, 2...)
 - In this sample, GPU number (`rank % num of devices`) is used



Data Transfer

- mm sample does not use communication
- If we want to do, the basic method is
 - (1) Copy data on GPU memory to CPU (cudaMemcpy)
 - (2) Transfer between processes (MPI_Send/MPI_Recv)
 - (3) Copy data on CPU memory to GPU (cudaMemcpy)



NOTE:
Recent MPI supports
GPU direct,
to direct communication
on GPU memory



We Have Learned

- Part 1: Shared memory parallel programming with OpenMP
- Part 2: GPU programming with OpenACC and CUDA
- Part 3: Distributed memory parallel programming with MPI

Many common strategies towards faster software:

- To understand source of bottleneck
- Reducing computation and communication
- Overlapping computation and communication
- To understand property of architecture

Assignments in MPI Part (Abstract)



Choose one of [M1]—[M3], and submit a report
Due date: **June 9 (Thursday)**

[M1] Parallelize “diffusion” sample program by MPI.

[M2] Improve mm-mpi sample in order to reduce memory consumption.

[M3] (Freestyle) Parallelize *any* program by MPI.

For more detail, please see MPI (1) slides on May 19



- Thank you for participating in practical parallel computing

Let's enjoy high performance computing!