

Practical Parallel Computing (実践的並列コンピューティング)

Part 2: GPU

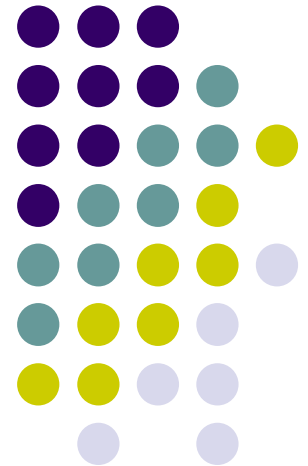
No 3: Threads in CUDA

May 11, 2023

Toshio Endo

School of Computing & GSIC

endo@is.titech.ac.jp





Overview of This Course

- Part 0: Introduction
 - 2 classes
- Part 1: OpenMP for shared memory programming
 - 4 classes
- Part 2: **GPU** programming
 - 4 classes **← We are here (3/4)**
 - OpenACC (1.5 classes) and **CUDA (2.5 classes)**
- Part 3: **MPI** for distributed memory programming
 - 4 classes

Comparing OpenMP/OpenACC/CUDA




	OpenMP	OpenACC	CUDA
Processors	CPU	CPU+GPU	
File extension	.c, .cc		.cu
To start parallel (GPU) region	#pragma omp parallel	#pragma acc kernels	func<<<..., ...>>>()
To specify # of threads	export OMP_NUM_THREADS=...	(num_gangs, vector_length etc)	
Desirable # of threads	# of CPU cores or less	# of GPU cores or “more”	
To get thread ID	omp_thread_num()	-	blockIdx, threadIdx
Parallel for loop	#pragma omp for	#pragma acc loop	-
Task parallel	#pragma omp task	-	-
To allocate device memory	-	#pragma acc data	cudaMalloc()
To copy to/from device memory	-	#pragma acc data #pragma acc update	cudaMemcpy()
Functions on GPU	-	#pragma acc routine	__global__, __device__

※ “# of XXX” = “The number of XXX”

Threads in CUDA



Specify 2 numbers (at least) for number of threads, when calling a GPU kernel function



cf) func <<< 4, 3 >>> (); → 12 threads


Number of thread blocks
= gridDim

Number of threads per block
= blockDim

The reason is related to GPU hardware
Thread block \Leftrightarrow SMX, Thread \Leftrightarrow CUDA core



Thread Block ID, Thread ID



For every thread, `gridDim.x = 4`, `blockDim.x = 3`

Note: In order to see the entire sequential ID, we should compute
`blockIdx.x * blockDim.x + threadIdx.x`

The Case of add-cuda Sample



- </gs/hs1/tga-ppcomp/23/add-cuda/>
- We want to do `for (i = 0; i < 100; i++) { DA[i] += DB[i]; }`


[CPU side]

```
add<<<20, 5>>>(...);
```

*20x5=100 threads
will execute add function*

[GPU side]

```
__global__ void add(int *DA, int *DB)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    DA[i] += DB[i];
    return;
}
```



How is Number of Threads Designed? (1)



On CUDA, a different strategy is required from on OpenMP

- On OpenMP, number of threads (OMP_NUM_THREADS) should be \leq CPU cores (or hyper threads)
 - The number is basically determined by hardware
 - ≤ 14 on q_node node, ≤ 56 on f_node
- On CUDA, it is better to use number of thread \geq GPU cores
 - ≥ 3584 on TSUBAME3's P100 GPU
 - You can use >1,000,000 threads!

How is Number of Threads Designed? (2)



We have to decide 2 numbers
<<<block number, block size>>>

A better way would be

- (1) We decide **total** number of threads P
- (2) We tune each block size BS
 - Good candidates are 16, 32, 64, ... 1024
- (3) Then block number is P/BS
 - We consider indivisible cases later



“mm” sample: Matrix Multiply (related to [G2])



CUDA versions are at


- </gs/hs1/tga-ppcomp/23/mm-v1-cuda/>
- </gs/hs1/tga-ppcomp/23/mm-cuda/>

A: a $(m \times k)$ matrix, B: a $(k \times n)$ matrix

C: a $(m \times n)$ matrix

$$C \leftarrow A \times B$$

- Supports variable matrix size
- Execution: `./mm [m] [n] [k]`



On CUDA, We need to design

- (1) How we parallelize computation
- (2) How we put data on host memory & device memory




How We Parallelize Computation

In mm, we can compute different C elements in parallel

- On the other hand, it is harder to parallelize dot-product loop


OpenMP

- Parallelize column-loop
(or row-loop)



CUDA (mm-v1-cuda)

- We can create many threads
- 1 thread computes 1 row
 - We use m threads



✂ This is not the unique way



Parallelism in mm-v1-cuda

- It is ok to make >1000 , >10000 threads on CUDA
- We use m threads for m rows computation

```
matmul_kernel<<<m/BS, BS>>>(. . . . .);
```

gridDim

blockDim (BS=16 in
this sample)

1 element for 1 row \rightarrow No need of “i” loop in this sample

Note: $\langle\langle\langle m, 1 \rangle\rangle\rangle$ also works, but speed is not good
 $\langle\langle\langle 1, m \rangle\rangle\rangle$ causes an error if $m > 1024$ (CUDA's rule)


If Number of Threads is Indivisible by BlockDim



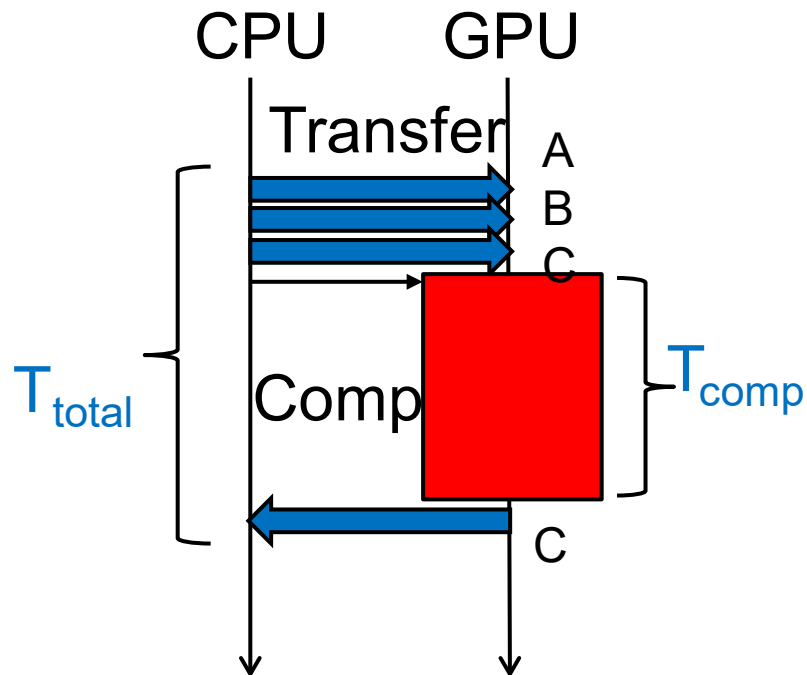
- m : the number of threads
- BS : blockDim
- If m may be indivisible by BS , we should use $\lll(m+BS-1)/BS, BS\ggg$
- But # of threads may be larger m . “Extra” threads ($id \geq m$) should not work. See mm-v1-cuda/mm.cu

Example: $m=7$, $BS=5 \rightarrow \lll 2, 5 \ggg$

10 threads start working, but 3 threads should do nothing



Data Transfer in mm-v1-cuda



(1) A, B, C are copied from CPU to GPU

- `cudaMemcpy(DA, A, ...) ...`

(2) Computation is done on GPU

(3) C is copied from GPU to CPU

- `cudaMemcpy(C, DC, ...)`



Notes in Time Measurement

- `clock()`, `gettimeofday()` must be called from CPU
- For accurate measurement, we should call `cudaDeviceSynchronize()` before measurement
 - Actually GPU kernel function call and `cudaMemcpy(HostToDevice)` are non-blocking



Speed of mm-v1-cuda

- Measured with a P100 GPU on TSUBAME3
 - `./mm 1000 1000 1000` and so on
- The program outputs 2 speeds
 - Speed **with** data transfer costs → shown on the table
 - Speed **without** data transfer costs

m=n=k	mm-acc (Gflops)	mm-v1-cuda (Gflops)
1000	158	15
2000	193	29
4000	214	50
6000	198	70
8000	198	85

Why slow?

Discussion on Speed (related to [G2])



mm-v1-cuda is slower than mm-acc!

- In mm-acc, i-loop and j-loop has “loop independent”
→ $m \times n$ elements are computed in parallel
- In mm-v1-cuda, we use m threads in total
 - We should use more threads on a GPU!
 - At least, $\geq 3,584$ = number of CUDA cores
 - We see $m=8000$ threads are still insufficient, and slow

Improvement: How to Use More Threads




In mm, computation of **each C element** is independent with each other

mm-v1-cuda

- 1 thread computes 1 row


→ We use m threads



mm-cuda

- 1 thread computes 1 element

→ We use $m \times n$ threads !!




Creating Threads with 2D/3D IDs



- Now we want to make $m*n$ (may be $>1,000,000$) threads
 - `<<<(m*n)/BS, BS>>>` is ok, but coding is bothersome
- On CUDA, gridDim and blockDim may have “dim3” type, 3D vector structure with x, y, z fields

cf) func `<<< dim3(4,2,1), dim3(3,2,1) >>> ();` → 48 threads




✂ This example is the case of 2D (Z dimensions are 1)

Thread IDs in multi-dimensional cases



In the case of func `<<< dim3(4,2,1), dim3(3,2,1) >>> ();`




- For every thread,
gridDim.x=4, gridDim.y=2, gridDim.z=1
blockDim.x=3, blockDim.y=2, blockDim.z=1
- For the thread with blue mark,
blockIdx.x=1, blockIdx.y=1, blockIdx.z=0
threadIdx.x=2, threadIdx.y=0, threadIdx.z=0



Threads in mm-cuda Sample

- The total number of threads are $m \times n$
- How do we determine gridDim, blockDim?
 - `<<<m, n>>>` does not work for constraints explained later ☹
- Here, we determine blockDim as $x=16, y=16 \rightarrow 256$ threads per block
 - Then gridDim is computed from M, N
- x is mapped to row index, y is mapped to column index (⌘)



⌘ A different mapping is possible, but inefficient (explained next time)



Code in mm-cuda

gridDim

blockDim

```
matmul_kernel<<<dim3(m / BS, n / BS, 1), dim3(BS, BS, 1)>>>  
(DA, DB, DC, m, n, k);
```

BS=16 in this sample

Actually, we use rounding up

In matmul_kernel function,

:

$j = \text{blockIdx.y} * \text{blockDim.y} + \text{threadIdx.y};$

$i = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x};$

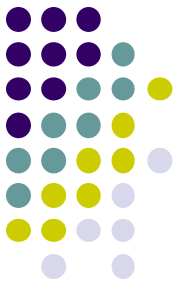
: This thread computes C_{ij} ← Only k loop




About Programming Efforts

- So far, mm-cuda with $(m * n)$ threads has been explained
 - How fast is it? Please measure it (Related to [G2])
- On the other hand, codes are more complex than mm-acc
Programmer have to
 - Call cudaMalloc, cudaMemcpy
 - Determine the number of blocks, threads
 - Determine the structure of for-loop
 - We have to change code largely when parallelization method is changed
 - In OpenACC, adding “#pragma acc loop independent” works
 - :

CUDA Rules on Number of Threads



`func<<<A, B>>> (...);` (*A, B are integers*) is same as
`func<<<dim3(A,1,1), dim3(B,1,1)>>> (...);`




Size of blockDim has severe limitation ☹

Cf) `<<<m, n>>>` causes an error if $n > 1024$ ☹



Rules for Memory/Variables

- Variables declared in GPU kernel functions are “**thread private**”



- Device memory is **shared** by all CUDA threads
 - Be careful to avoid race condition problem (multiple threads write same address)
 - Reading same address is ok
- Do not forget host memory and device memory are separated



Two Types of GPU Kernel Functions

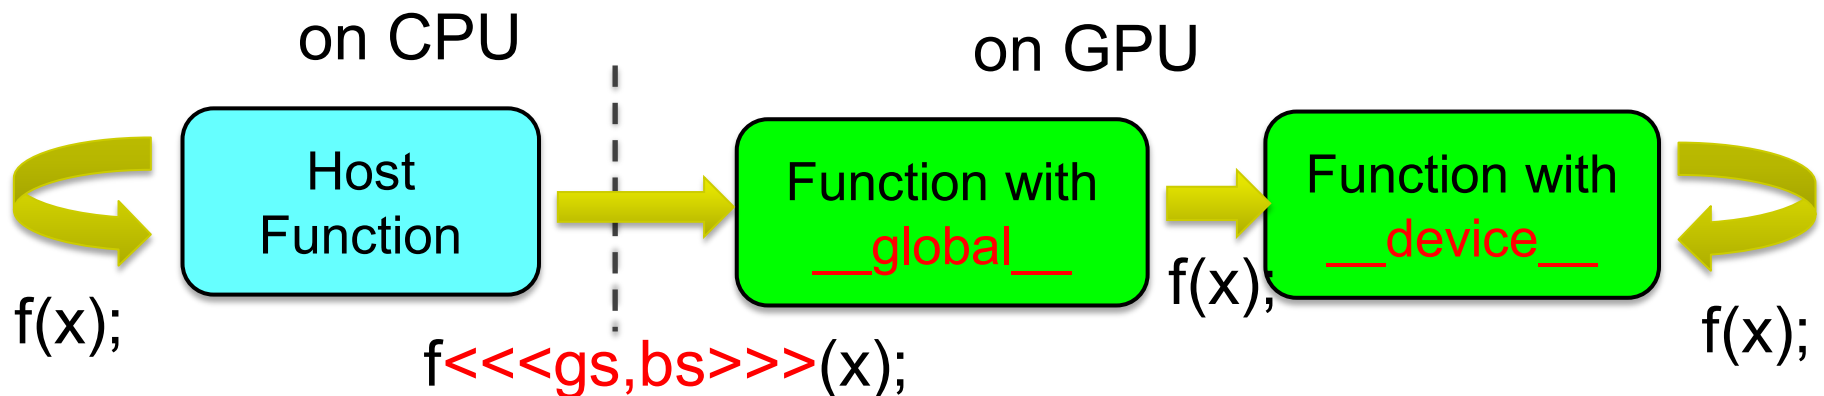
1) Functions with `__global__` keyword

- “Gateway” from CPU
- Return value type must be “void”

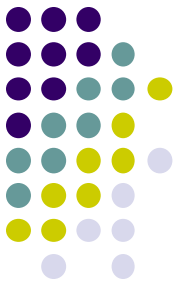
2) Function with `__device__` keyword

- Callable only from GPU
- Can have return values
- Recursive call is OK


→ In OpenACC, `#pragma acc routine`



“diffusion” Sample Program related to [G1]



An example of diffusion phenomena:




The ink spreads gradually, and finally the density becomes uniform (Figure by Prof. T. Aoki)

Available at </gs/hs1/tga-ppcomp/23/diffusion/>
You can use </gs/hs1/tga-ppcomp/23/diffusion-cuda/>

- Execution: `./diffusion [nt]`
 - nt: Number of time steps



Discussion on diffusion sample



Both arrays have to be on GPU device memory when computations are done

Consideration of Parallelizing Diffusion with CUDA related to [G1]



- x, y loops can be parallelized
- t loop cannot be parallelized

[Data transfer from CPU to GPU]

```
for (t = 0; t < nt; t++) {
```

```
    for (y = 1; y < NY-1; y++) {  
        for (x = 1; x < NX-1; x++) {  
            :  
        }  
    }
```

```
}
```

[Data transfer from GPU to CPU]


GPU computation must be a distinct function (GPU kernel function)

It's better to transfer data *out of* t-loop



Considering CUDA Threads

- How do we design threads on CUDA?
- There are several choices in [G1]
 - 1thread = 1row
 - We use NY threads in total → only x-loop in kernel function
 - 1thread = 1column
 - We use NX threads in total → only y-loop in kernel function
 - 1thread = 1element
 - We use NX x NY threads in total → No loop in kernel function!
 - This looks fast since the number of threads is very large





Mapping between Threads and Data




mm-cuda:

Matrices has
column-major format



diffusion:

2D array has
row-major format



```
j = blockIdx.y * blockDim.y +  
threadIdx.y;  
i = blockIdx.x * blockDim.x +  
threadIdx.x;  
: This thread computes Cij
```

```
y = blockIdx.y * blockDim.y +  
threadIdx.y;  
x = blockIdx.x * blockDim.x +  
threadIdx.x;  
: This thread computes [y][x]
```


[Q] What if the dimensions are exchanged?



Considering gridDim/blockDim (1)

```
func <<< dim3 ( ?, ?, ? ), dim3 ( ?, ?, ? ) >>> (...);
```

gridDim blockDim



(1) We decide total number of threads

→ (NX, NY, 1) threads

- See notes on the next page

(2) We tune each block size (blockDim)

→ Good candidates are (4, 4, 1), (8, 8, 1), (16, 16, 1), (32, 32, 1)

- The number must be ≤ 1024
- How about non-square blocks?

(3) Then block number (gridDim) is determined

We should consider indivisible cases

Considering gridDim/blockDim (2)



- In diffusion, Points $[1, NX-1) \times [1, NY-1)$, excluded boundary, should be computed

There are choices:



(A) Create $NX \times NY$ threads

- Thread (x,y) computes (x,y)
- Threads with below IDs do nothing
 - $x == 0$ or $y == 0$ or $x \geq NX-1$ or $y \geq NY-1$

(B) Create $(NX-2) \times (NY-2)$ threads

- Thread (x,y) computes $(x+1,y+1)$
- Threads with below IDs do nothing
 - $x \geq NX-2$ or $y \geq NY-2$

Either is ok 😊



Discussion on Data Transfer of Diffusion



Both codes will work, but how about speeds?

[Data transfer from CPU to GPU]

```
for (t = 0; t < nt; t++) {  
    :
```

```
    for (y = 1; y < NY-1; y++) {  
        for (x = 1; x < NX-1; x++) {  
            :  
        }  
    }  
}
```

```
}
```

[Data transfer from GPU to CPU]

Computation: $O(NX \ NY \ nt)$
Transfer: $O(NX \ NY)$

```
for (t = 0; t < nt; t++) {  
    :
```

[Data transfer from CPU to GPU]

```
for (y = 1; y < NY-1; y++) {  
    for (x = 1; x < NX-1; x++) {  
        :  
    }  
}
```

```
}
```

[Data transfer from GPU to CPU]

Computation: $O(NX \ NY \ nt)$
Transfer: $O(\underline{NX \ NY} \ nt)$

What Can be Done in GPU Functions?



- Basic computations (+, -, *, /, %, &&, ||...) are OK
- if, for, while, return are OK
- Device memory access is OK
- Host memory access is NG
- Calling host functions is NG
- Calling most of functions in libc or other libraries for CPUs are NG
 - Several mathematical functions, sin(), sqrt()... are OK
 - printf() is OK
 - Calling malloc()/free() on GPU is OK, if the size must be small
 - Usually, use cudaMalloc() on CPU

Assignments in GPU Part (Abstract)



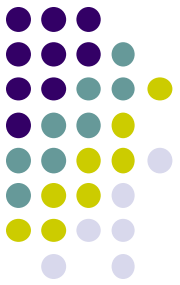
Choose one of [G1]—[G3], and submit a report

Due date: May 25 (Thursday)

[G1] Parallelize “diffusion” sample program by OpenACC or CUDA

[G2] Evaluate speed of “mm-acc” or “mm-cuda” in detail

[G3] (Freestyle) Parallelize *any* program by OpenACC or CUDA.



Next Class:

- GPU Programming (4)
 - Discussion on speed and GPU hardware
- Also please note due date of OpenMP assignment is May 11 (Today!)