

Practical Parallel Computing (実践的並列コンピューティング)

Part 2: GPU

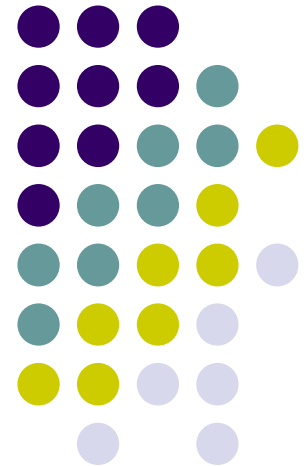
No 4: Effects of GPU Architecture

May 15, 2023

Toshio Endo

School of Computing & GSIC

endo@is.titech.ac.jp





Overview of This Course

- Part 0: Introduction
 - 2 classes
- Part 1: OpenMP for shared memory programming
 - 4 classes
- Part 2: **GPU** programming
 - 4 classes **← We are here (4/4)**
 - OpenACC (1.5 classes) and **CUDA (2.5 classes)**
- Part 3: **MPI** for distributed memory programming
 - 4 classes

Comparing OpenMP/OpenACC/CUDA



	OpenMP	OpenACC	CUDA
Processors	CPU	CPU+GPU	
File extension	.c, .cc		.cu
To start parallel (GPU) region	#pragma omp parallel	#pragma acc kernels	func<<<..., ...>>>()
To specify # of threads	export OMP_NUM_THREADS=...	(num_gangs, vector_length etc)	
Desirable # of threads	# of CPU cores or less	# of GPU cores or “more”	
To get thread ID	omp_thread_num()	-	blockIdx, threadIdx
Parallel for loop	#pragma omp for	#pragma acc loop	-
Task parallel	#pragma omp task	-	-
To allocate device memory	-	#pragma acc data	cudaMalloc()
To copy to/from device memory	-	#pragma acc data #pragma acc update	cudaMemcpy()
Function on GPU	-	#pragma acc routine	__global__, __device__

※ “# of XXX” = “The number of XXX”

Speed of GPU Programs and GPU Architecture



Advanced topic

Case 1: How should block-size be determined?

When creating 1,000,000 threads,

- `<<<1, 1000000>>>` causes an error
 - blockDim must be ≤ 1024
- `<<<1000000, 1>>>` can work, but slow
- `<<<1000, 1000>>>` is faster \rightarrow Why?




Case 2: How should each thread access memory?

- In mm-cuda, $(x = \text{row}, y = \text{col})$ and $(x = \text{col}, y = \text{row})$ shows different speed


Knowledge of GPU architecture helps understanding of speeds

Why Do We Have to Specify both `gridDim` and `blockDim`?



- and why did NVIDIA decide so?

→ Hierarchical structure of GPU processor is considered



Structure of P100 GPU
(16nm, 15Billion transistors)

1 GPU = 56 **SMXs**

1 **SMX** = 64 CUDA cores
(16 cores x 4 groups)

→ 1GPU=3,584 CUDA cores



Mapping between Threads and Cores





- 1 thread blocks (or more) run on 1 SMX
 - At least 56 blocks are needed to use all SMXs on P100
 - `gridDim (gx*gy*gz)` should be ≥ 56
- 1 thread (or more) run on a CUDA core
 - At least $56 \times 64 = 3584$ threads in total are needed to use all CUDA cores on P100
 - `Total threads (gx*gy*gz * bx*by*bz)` should be ≥ 3584
- 32 consecutive threads (in a block) are batched (called a warp) and scheduled
 - At least 32 threads per block are needed for performance
 - `blockDim (bx*by*bz)` should be ≥ 32




Warp: Internal Execution Unit

thread < **warp** < thread block < grid




- Threads in a thread block are internally divided into “**warp**”, a group of contiguous 32 threads
- 32 threads in a warp always are executed synchronously
 - They execute the same instruction simultaneously
 - Only 1 program counter for 32 threads → GPU hardware is simplified
 - Actually 32 threads are executed on 16 CUDA cores





Observations due to Warps

- If number of threads per block (blockDim) is not $32 \times n$, it is inefficient
 - Even if blockDim=1, the system creates a warp for it
- Characteristics in memory addresses accessed by threads in a warp affect the performance
 - Coalesced accesses are fast




⌘ In multi-dimensional cases (blockDim.y>1 or blockDim.z>1), “neighborhood” is defined by x-dimension




Coalesced Memory Access

- When threads in a warp access “neighbor” address on memory (**coalesced access**), it is more efficient

Coalesced access
→ **Faster**



Non-coalesced access
→ **Slower**






Accesses in mm-cuda Sample

- **mm-cuda**: ($x = \text{row}, y = \text{col}$) \rightarrow coalesced and fast
- **mm-nc-cuda**: ($x = \text{col}, y = \text{row}$) \rightarrow non-coalesced and slow
 - </gs/hs1/tga-ppcomp/23/mm-nc-cuda>

We should see “what data are accessed by threads in a warp simultaneously”



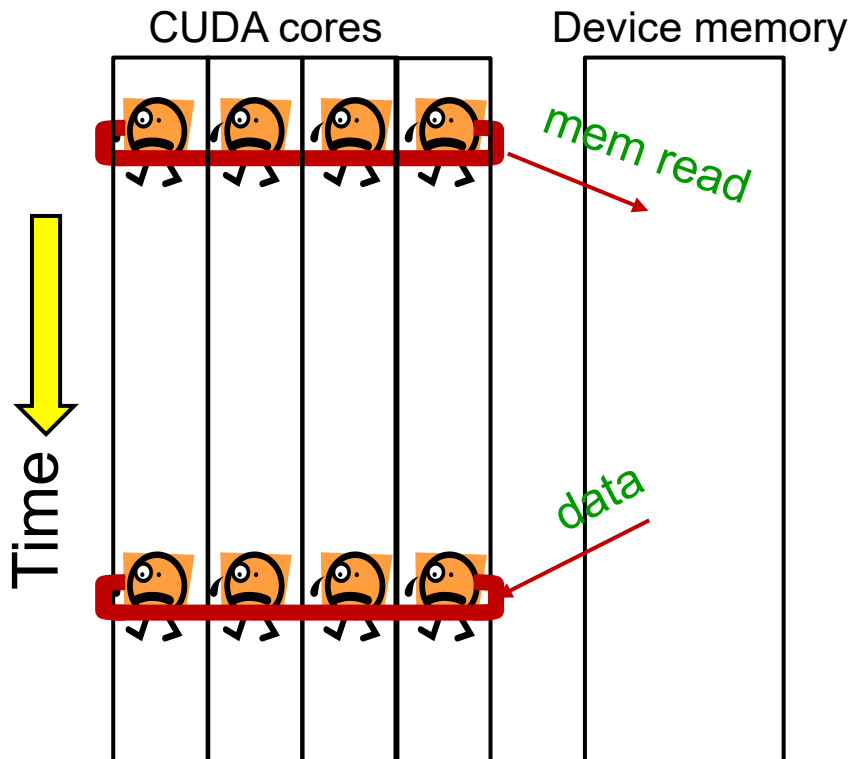
matrices in column-major format

Why $\#threads \gg \#cores$ Works Well on GPUs?




- GPU supports very fast (~ 1 clock) context switches
→ With many threads, memory access latency can be hidden

$\#threads == \#cores$



$\#threads > \#cores$




Considering Branches in Parallel Programs



Consider this code. How long is execution time?

```
if (thread-id % 2 == 0) {  
    : // (A) 30msec  
} else {  
    : // (B) 20msec  
}
```

On CPU (OpenMP)




On GPU, threads in a warp must execute the same instruction. What happens?





Branches on GPU (1)



Some threads are made sleep
Both “then” and “else” are executed!

→ Answer to previous question is **50ms** !

⌘ Similar cases happen in for, while...



Branches on GPU (2)


- As exceptional cases, if threads in a warp “agree” in branch condition, either “then” part or “else” part is executed → Efficient!
 - If there is difference of opinion (previous page), it is called a **divergent branch**
- Agreement among buddies (threads in a warp) is important for speed

Considering Data Transfer Costs of mm Sample



- In mm sample, the speed is degraded by **data transfer costs** ☹️
- This can be improved by combination of:
 1. Split computation
 2. Using CUDA streams


→ The faster sample is at
</gs/hs1/tga-ppcomp/23/mm-str-cuda/>





Split mm Computation (1)


- Computation of “ $C \leftarrow A \times B$ ” is split by splitting B and C vertically
 - $C_1 \leftarrow A \times B_1, C_2 \leftarrow A \times B_2, \dots, C_n \leftarrow A \times B_n$
- The n computations are independent each other



A is reused for all computations



Split mm Computation (2)



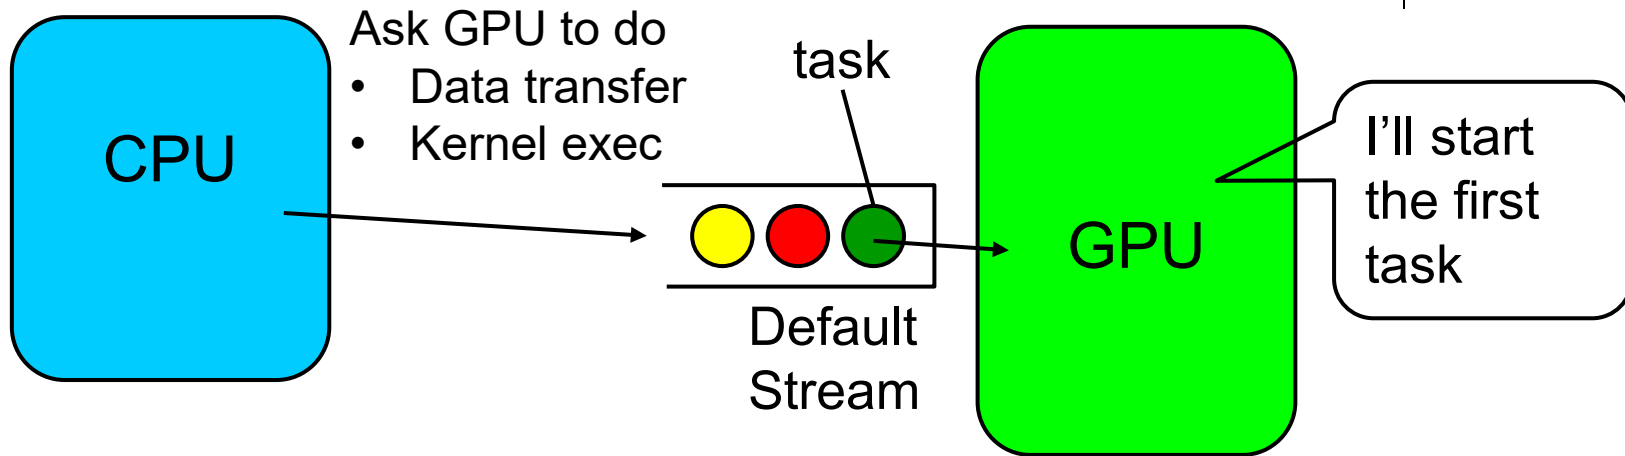
Algorithm:

- (1) Copy A from CPU to GPU
- (2) For each partition i (sequentially)
 - (1) Copy B_i and C_i to GPU
 - (2) Compute $C_i \leftarrow A \times B_i$
 - (3) Copy back C_i to GPU

This does NOT improve speed yet, since neither total computation costs nor total transfer costs change

→ **cudaStream** is useful for hiding transfer costs

How GPU Executes Tasks (Without multiple streams)




- A GPU is idle until asked to do something by CPU
- CPU asks the GPU to do one of followings (called **tasks** here)
 - Data transfer (Host → Device) or
 - GPU Kernel function execution or
 - Data transfer (Device → Host)
- Then the task is put on a FIFO queue, called **default stream**
- GPU takes a task from the stream and executes it in FIFO

Asynchronous Executions with `cudaStream` (1)



What are `streams`?

- GPU's “service counters” that accept tasks from CPU
 - In addition to default stream user program can create streams,
 - Each stream looks like a FIFO queue



All of CUDA sample programs, except `mm-str-cuda`, are using the single “default stream”

Asynchronous Executions with `cudaStream` (2)



Create a stream

```
cudaStream_t str;  
cudaStreamCreate(&str); // Create a stream
```

Data transfer using a specific stream

```
cudaMemcpyAsync(dst, src, size, type, str);
```

Call GPU kernel function using a stream


```
func<<<gs, bs, 0, str>>>( ... );  
// 3rd parameter is related to for “shared memory”
```

Wait until all tasks on a stream are finished

```
cudaStreamSynchronize(str);
```

✂ The default stream is expressed as `(cudaStream_t)0`

How GPU Executes Tasks with Multiple Streams



- Rule: Tasks on the same stream are done in FIFO
 - The GPU considers that “tasks on one stream have dependencies, so I’ll do them in the order”
- If tasks are in different streams, and have different kinds, they may be done simultaneously
 - Kinds: Host→Device, kernel, Device→Host
 - Note: If tasks are in the same kind, no speed up

mm-str-cuda sample: Overlapping Computation and Transfer




n streams can be used for n independent “task sets”

- $C1 \leftarrow A \times B1$ (includes H->D, Calc, D->H)
- $C2 \leftarrow A \times B2$
-
- $Cn \leftarrow A \times Bn$

→ We will see speed up since
(Total comp time + Total trans time)
is improved to roughly
 $\max(\text{Total comp time}, \text{Total trans time})$

This is not a unique solution;
It is ok to use 2 or 3 streams repeatedly → we can
save GPU memory and stream resources

`cudaMallocHost()` is used instead of `malloc()`
This speeds up `cudaMemcpyAsync()`





More Things to Study

- Using CUDA shared memory
 - fast and small memory than device memory
- Unified memory in recent CUDA
 - `cudaMemcpy` can be omitted for automatic data transfer
 - Google with “`cudaMallocManaged`”
- Using Tensor-core to accelerate deep learning
 - Only on V100 GPUs or later
 - Unfortunately, TSUBAME3 has older P100 ☹
- Using multiple GPUs towards petascale computation
 - MPI+CUDA, MPI+OpenACC
- More and more...

Assignments in GPU Part (Abstract)



Choose one of [G1]—[G3], and submit a report

Due date: May 25 (Thursday)

[G1] Parallelize “diffusion” sample program by OpenACC or CUDA

[G2] Evaluate speed of “mm-acc” or “mm-cuda” in detail

[G3] (Freestyle) Parallelize *any* program by OpenACC or CUDA.



Next Class:

- Part 3: MPI Programming (1)
 - Introduction to distributed memory parallel programming
- Planned schedule
 - May 18: Part 3 (1)
 - May 22: Part 3 (2)
 - May 25: **cancelled (休講)** & Due for Part2 assignment
 - May 29: Part 3 (3)
 - June 1: Part 3 (4)