

Practical Parallel Computing (実践的並列コンピューティング)

Part 3: MPI

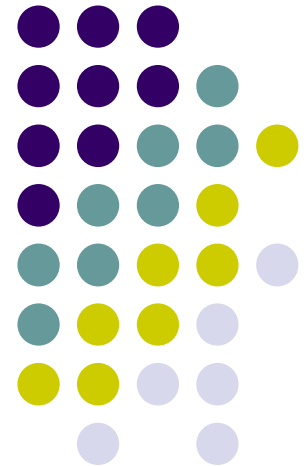
No 1: MPI Introduction

May 18, 2023

Toshio Endo

School of Computing & GSIC

endo@is.titech.ac.jp

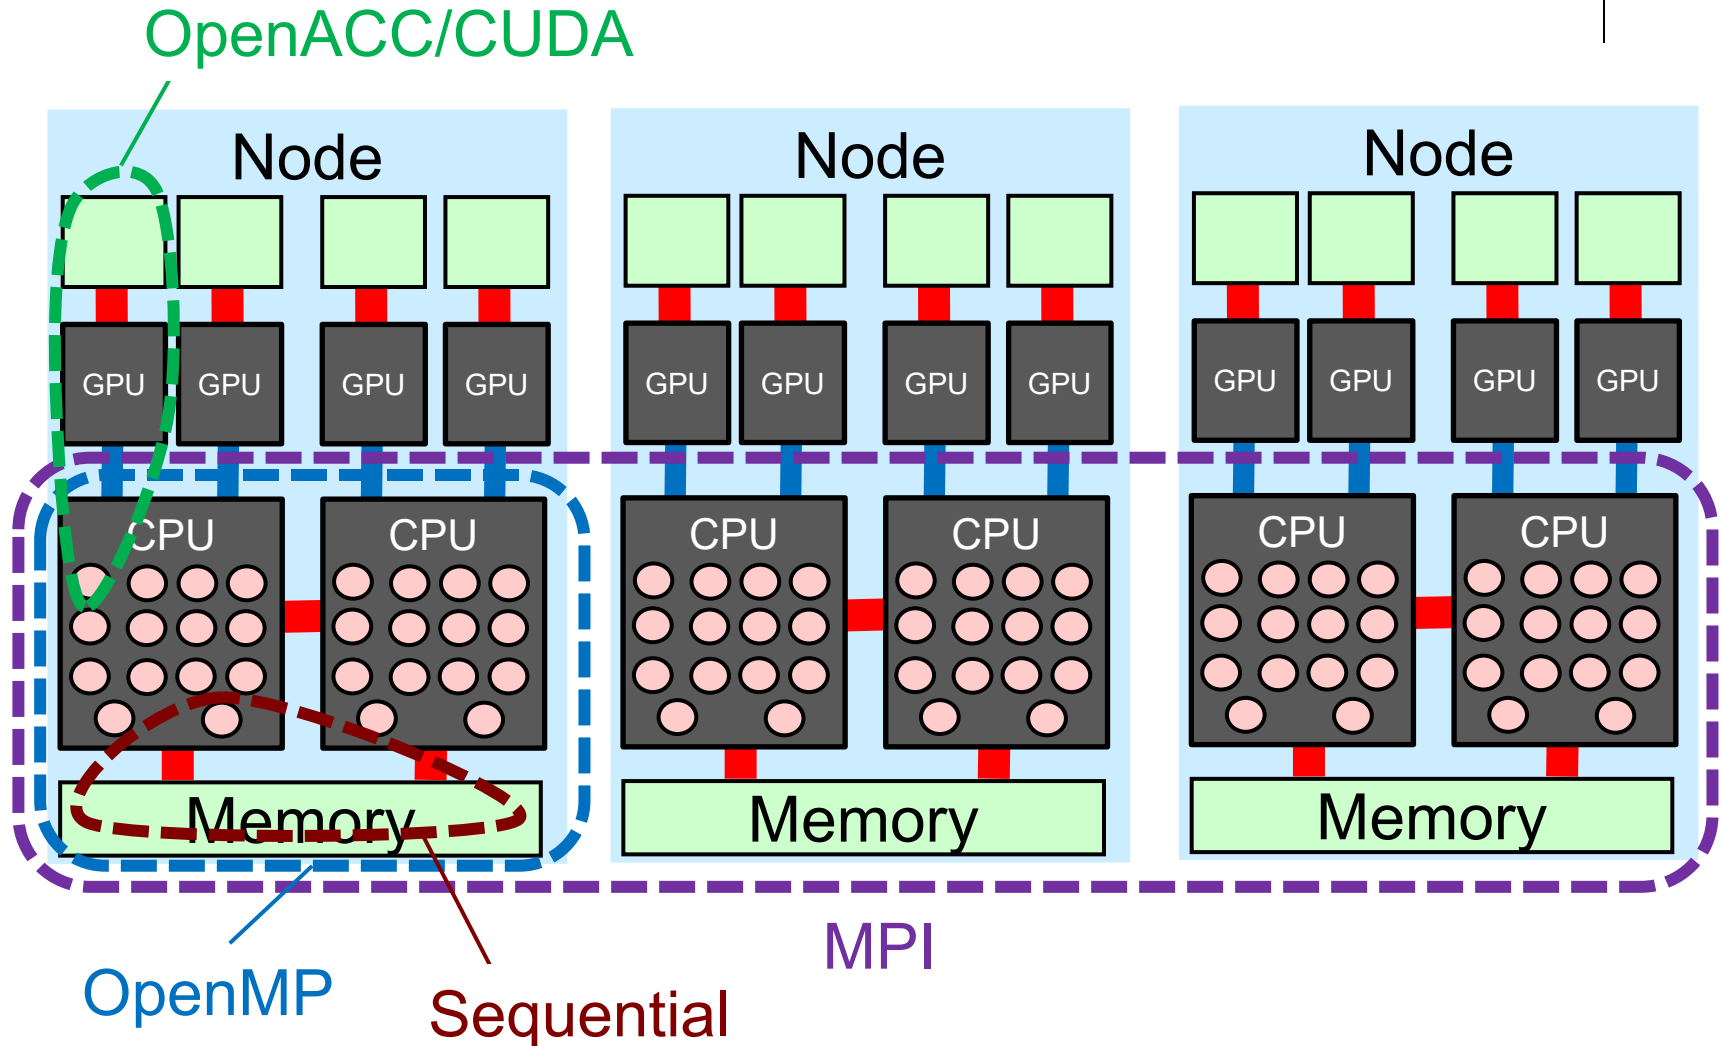
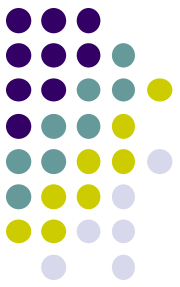


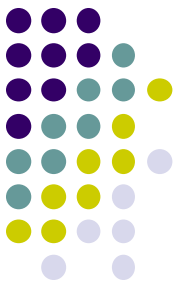


Overview of This Course

- Part 0: Introduction
 - 2 classes
- Part 1: OpenMP for shared memory programming
 - 4 classes
- Part 2: GPU programming
 - 4 classes ← We are here (1/4)
 - OpenACC (1.5 classes) and CUDA (2.5 classes)
- Part 3: **MPI** for distributed memory programming
 - 4 classes ← We are here (1/4)

Parallel Programming Methods on TSUBAME

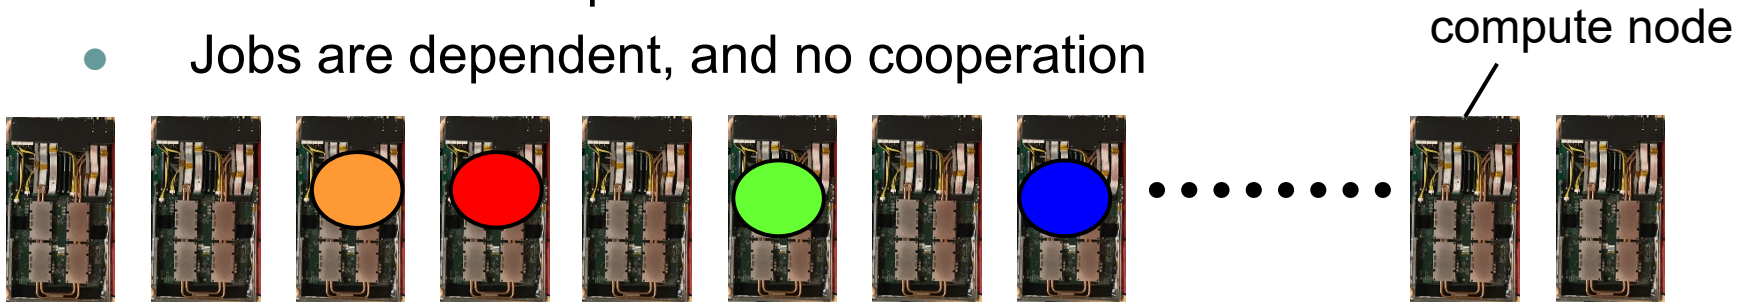




How Multiple Nodes are Used

1. Submit several jobs into job scheduler

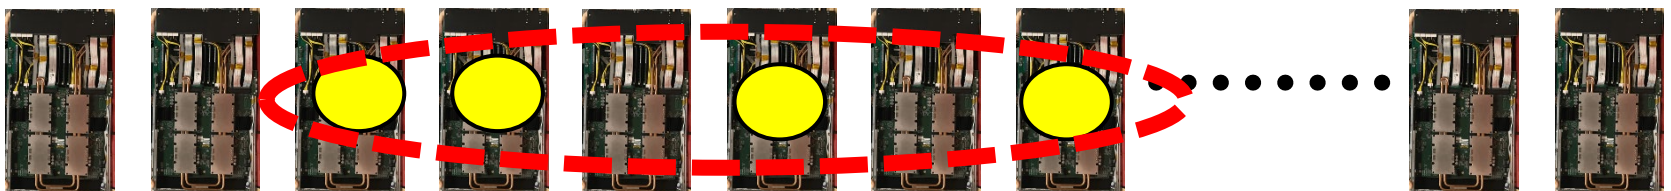
- cf) Program executions with different parameters → Parameter Sweep
- Jobs are dependent, and no cooperation



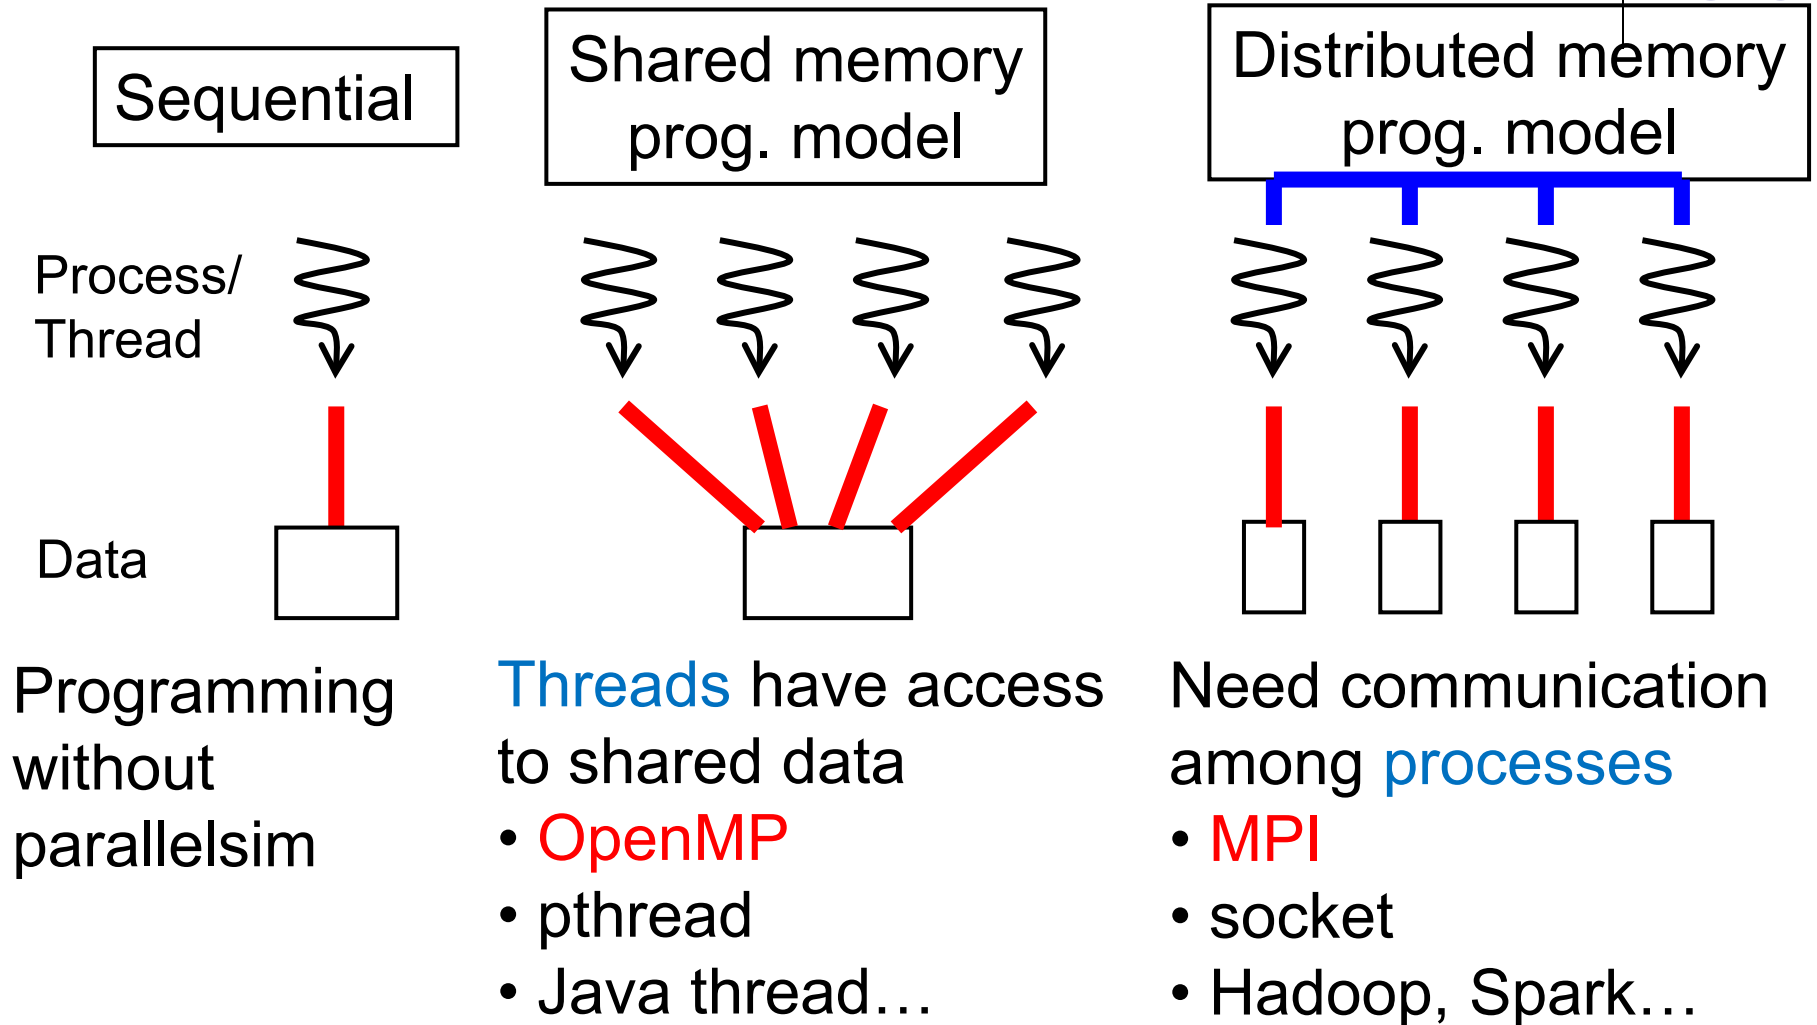
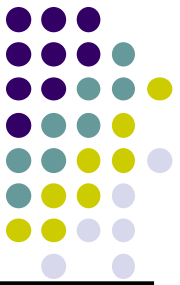
2. Use distributed memory programming → A single job can use multiple nodes

- Socket programming, Hadoop, Spark...
- And **MPI**

Part 3



Classification of Parallel Programming Models



MPI (message-passing interface)



- Parallel programming interface based on distributed memory model
- Used by C, C++, Fortran programs
 - Programs call MPI library functions, for **message passing** etc.
- There are several MPI libraries
 - OpenMPI (default) ← OpenMPI ≠ OpenMP ☹️
 - Intel MPI, SGI MPE, MVAPICH, MPICH...

Differences from OpenMP



In MPI,

- An execution consists of multiple **processes** (not threads)
 - We can use multiple nodes 😊
 - The number of running processes is basically constant
- No variables are shared. Instead **message passing** is used
 - Data distribution has to be programmed
- No smart syntaxes such as “omp for” or “omp task” 😞
 - Task distribution has to be programmed 😞



First MPI Sample

- `/gs/hs1/tga-ppcomp/23/hello-mpi`

[make sure that you are at a interactive node (r7i7nX)]

`module load cuda openmpi` *[Do once after login]*

`cd ~/t3workspace` *[In web-only route]*

`cp -r /gs/hs1/tga-ppcomp/23/hello-mpi .`

`cd hello-mpi`

`make`

[An executable file “hello” is created]

`mpixec -n 7 ./hello`

Number of
processes

Name of program
(using options are ok)

Compiling and Executing MPI Programs



Case of OpenMPI library on TSUBAME3.0

- To compile
 - `module load cuda openmpi`, and then use `mpicc`
 - For sample programs, “make” command works
- To execute
 - `mpiexec -n 7 ./hello`

Required for module dependency ☹️

Number of processes

↑ These methods uses 1 (current) node.

When We Want to Use More Processes than CPU Cores



- An interactive node has 7 CPU cores (= 14 Hyperthreads)
- If number of processes is larger than CPU cores, we need “--oversubscribe” option of OpenMPI
 - `mpiexec -n 7 ./hello`
 - `mpiexec -n 14 --oversubscribe ./hello`

Needed when we use
processes > physical cores
(7, in an interactive node)

From Single Node To Multiple Nodes



- So far, we used a single node
 - An interactive node with 7 cores = 14 Hyperthreads
 - For reports, a single node is enough
- MPI can support multiple nodes!

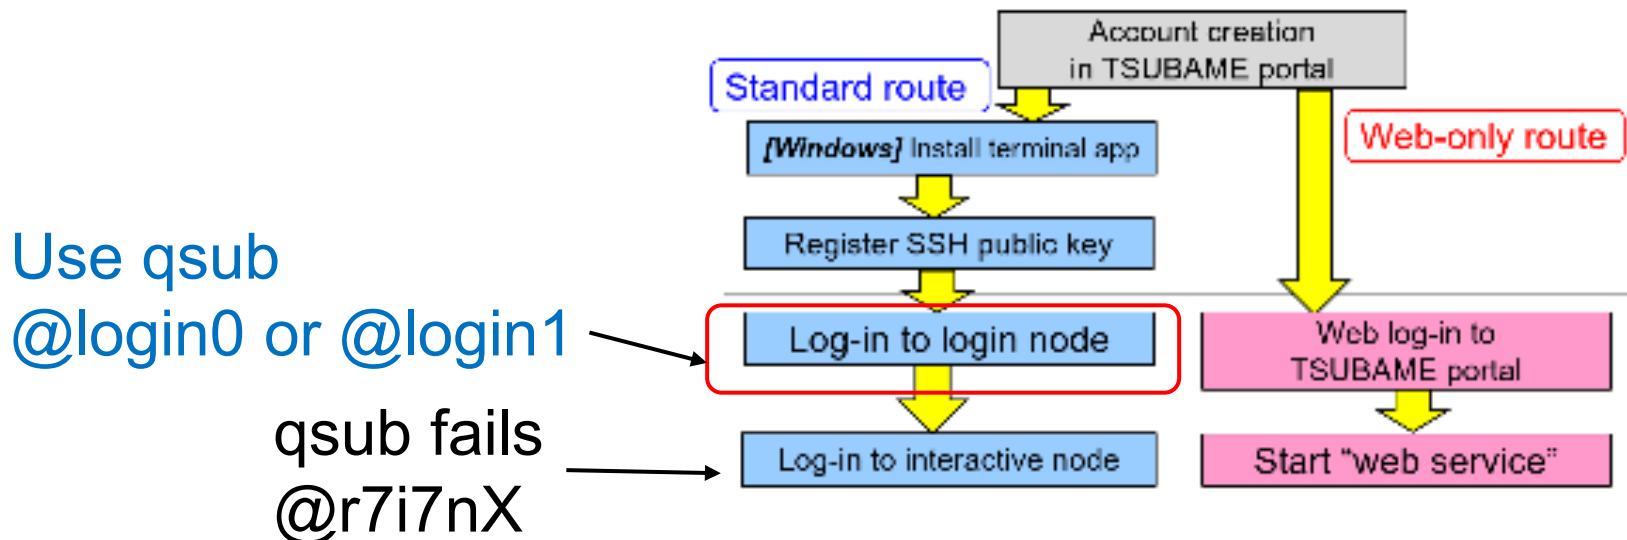
To use multiple nodes on TSUBAME3.0,
we need “job submission” !

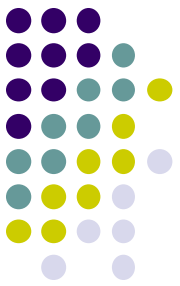


Notes on Job Submission in “Standard route”



- On an interactive node via “standard route”, qsub/qstat commands may not be found
- Please use qsub/qstat **on a login node**
 - (out of iqrsh)





Submit an MPI Job (case of OpenMPI)

- We are going to execute a job using
4 processes \times 2 nodes = 8 processes

(1) Make a script file: `job.sh`

```
#!/bin/sh
#$ -cwd
#$ -l q_core=2
#$ -l h_rt=00:10:00

. /etc/profile.d/modules.sh
module load cuda openmpi

mpirun -n 8 -npnode 4 ./hello
```

4core node x 2

Module preparation

Number of
processes

Number of
processes
per node

Program name
(and option)

(2) Submit the job with “`qsub`”

`qsub job.sh`

$\uparrow \leq 0:10:00$ and ≤ 2 nodes

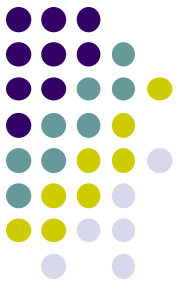
`qsub -g tga-ppcomp job.sh`

$\uparrow > 0:10:00$ or ≥ 3 nodes

Be careful for TSUBAME point

Notes in This Lecture

(also shown in 1-2 slides)



- First, please consider usage of interactive node (web usage/iqrsh)
- まずはインタラクティブノードの利用を検討してください (web usage/iqrsh)
- If necessary for reports, you can use up to 18,000 points in total per student. For more, please ask Endo
- 本講義のレポートの作成の目的で、一人あたり合計で18,000ポイントまで利用を認めます。より必要な場合は遠藤へ相談を
 - 18,000 points \div F_node x 5 hours
 - You can check point consumption on TSUBAME portal
- The TSUBAME group name is [tga-ppcomp](#)

Users need to follow the rules at www.t3.gsic.titech.ac.jp

利用時には www.t3.gsic.titech.ac.jp に示される規則を守る必要があります



Resource Types on TSUBAME3.0

- Choose one of resource types (number of cores, mainly)
 - It is like “instance types” in cloud systems
 - Please specify “proper” one

Resource type name	Physical CPU cores	Memory (GB)	GPUs
f_node	28	240	4
h_node	14	120	2
q_node	7	60	1
q_core	4	30	0
s_core	1	7.5	0
s_gpu	2	15	1

← Same as an interactive node

`#$ -l [resource_type] = [Number]`

Example:

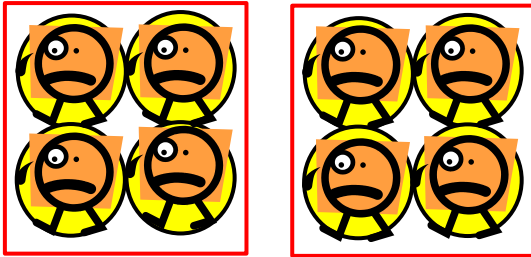
`#$ -l s_core=1` ← The minimum resource allocation

Nodes, Cores, MPI Processes



```

:
#$ -l q_core=2
:
mpirun -n 8 -npnode 4
...
    
```

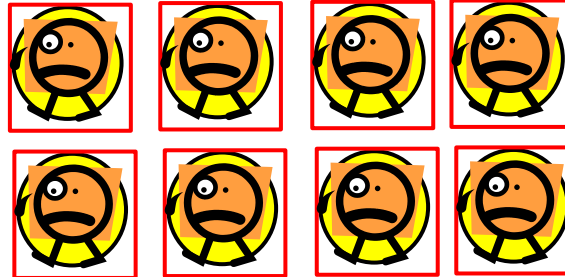


2 (virtual) nodes are prepared
Each node has 4 cores (q_core)

4 processes are created per node. Totally 8 are created
→ 2 nodes are used

```

:
#$ -l s_core=8
:
mpirun -n 8 -npnode 1
...
    
```

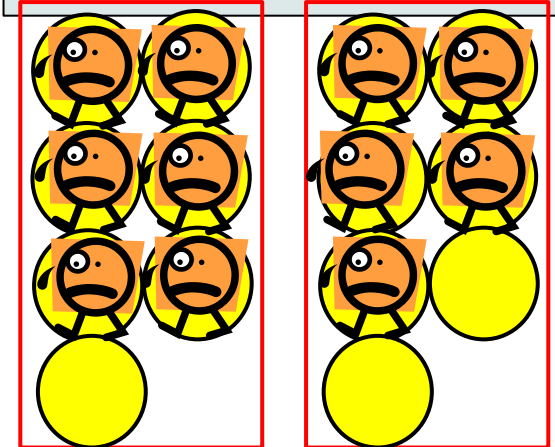


8 (virtual) nodes are prepared
Each node has 1 cores (s_core)

1 processes are created per node. Totally 8 are created
→ 8 nodes are used

```

:
#$ -l q_node=2
:
mpirun -n 11 -npnode 6
...
    
```



2 (virtual) nodes are prepared
Each node has 7 cores (q_node)

6 processes are created per node. Totally 11 are created
→ 2 nodes are used
(There are idle cores)



An MPI Program Looks Like

```
#include <stdio.h>
```

```
#include <mpi.h>
```

```
int main(int argc, char *argv[])
```

```
{
```

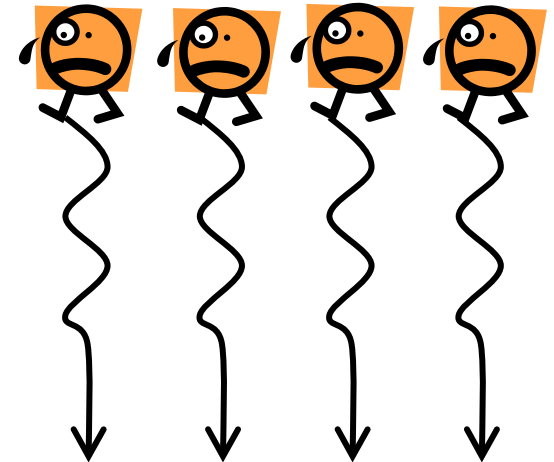
```
    MPI_Init(&argc, &argv); ← Initialize MPI
```

```
    (Computation/communication)
```

```
    MPI_Finalize(); ← Finalize MPI
```

```
}
```

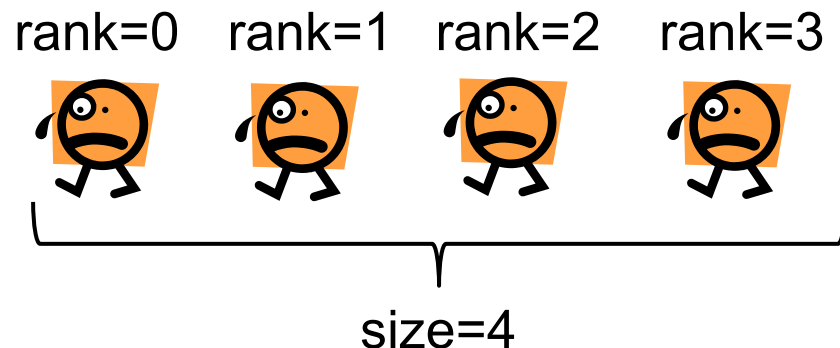
If number of
processes=4



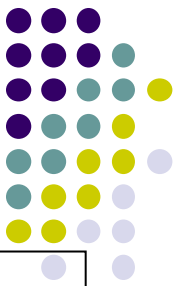


ID of Each MPI Process

- Each process has its ID (0, 1, 2...), called **rank**
 - `MPI_Comm_rank(MPI_COMM_WORLD, &rank);`
→ Get its rank
 - `MPI_Comm_size(MPI_COMM_WORLD, &size);`
→ Get the number of total processes
 - $0 \leq \text{rank} < \text{size}$
 - The rank is used as target of message passing



“mm” sample: Matrix Multiply



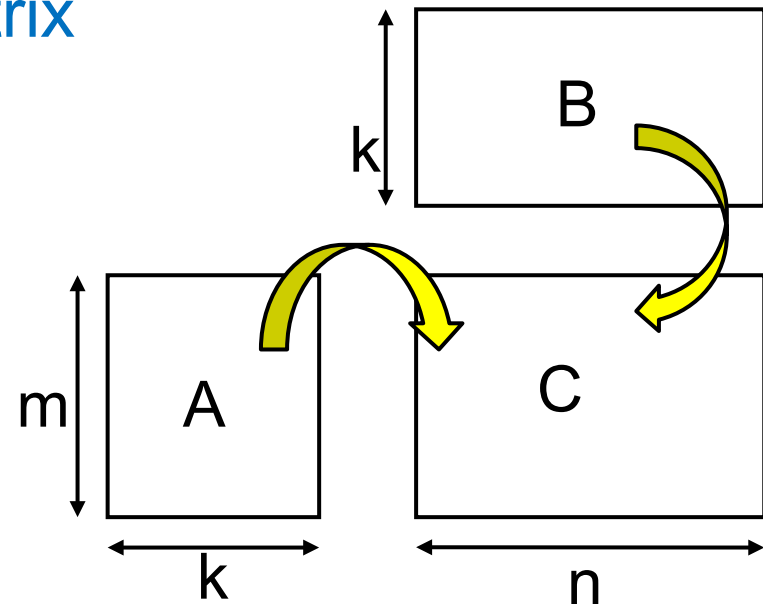
MPI version available at </gs/hs1/tga-ppcomp/23/mm-mpi/>

A: a $(m \times k)$ matrix, B: a $(k \times n)$ matrix

C: a $(m \times n)$ matrix

$$C \leftarrow A \times B$$

- Algorithm with a triple for loop
- Supports variable matrix size.
 - Each matrix is expressed as a 1D array by *column-major* format

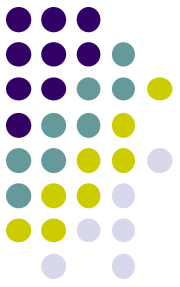


Execution:

`mpiexec -n [np] ./mm [m] [n] [k]` (interactive)

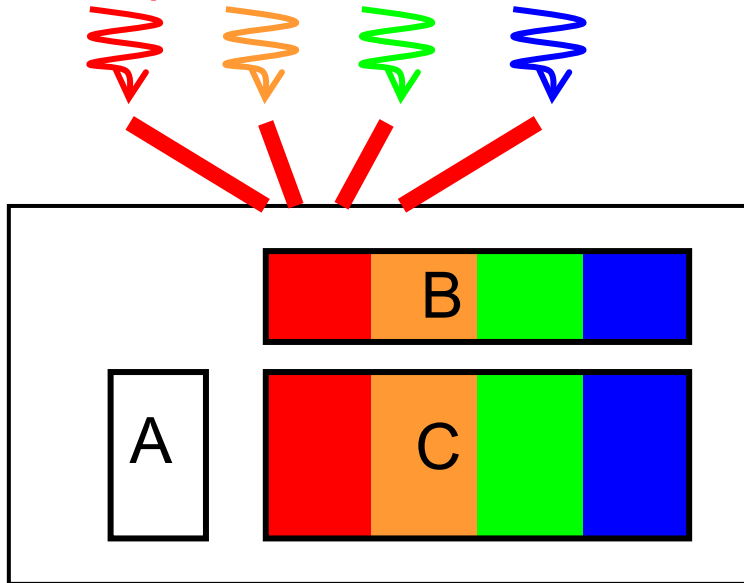
`mpiexec -n [np] -npernode [nn] ./mm [m] [n] [k]` (in job script)

Why Distributed Programming is More Difficult (case of mm-mpi)



Shared memory with OpenMP:

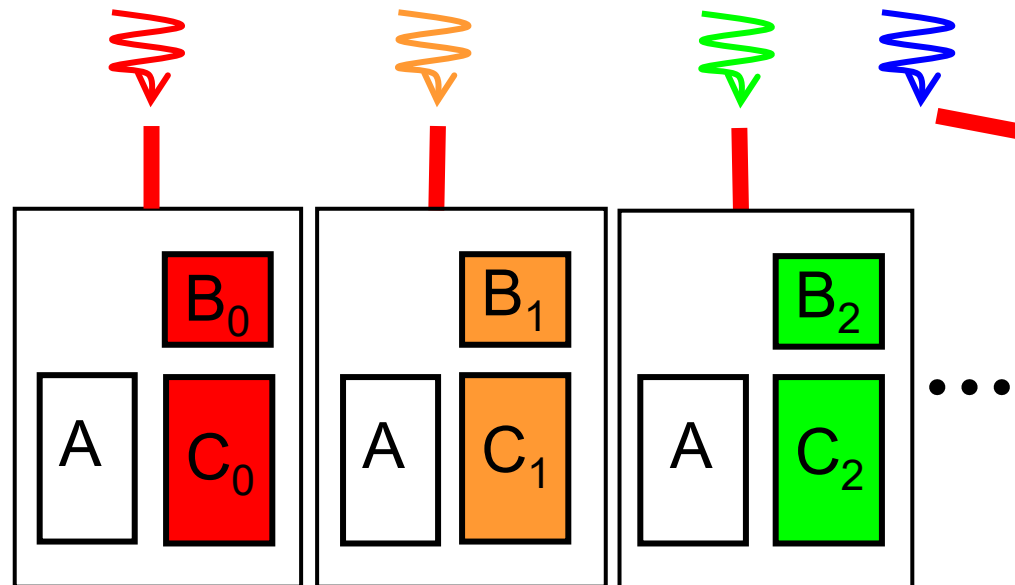
Programmers consider how **computations** are divided



In this case, matrix A is accessed by all threads
→ Programmers **do not have to know** that

Distributed memory with MPI:

Programmers consider how **data and computations** are divided

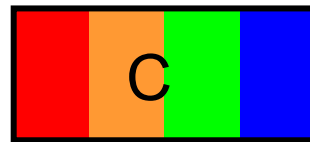


Programmers **have to design** which data is accessed by each process

Programming Data Distribution

(case of mm-mpi)

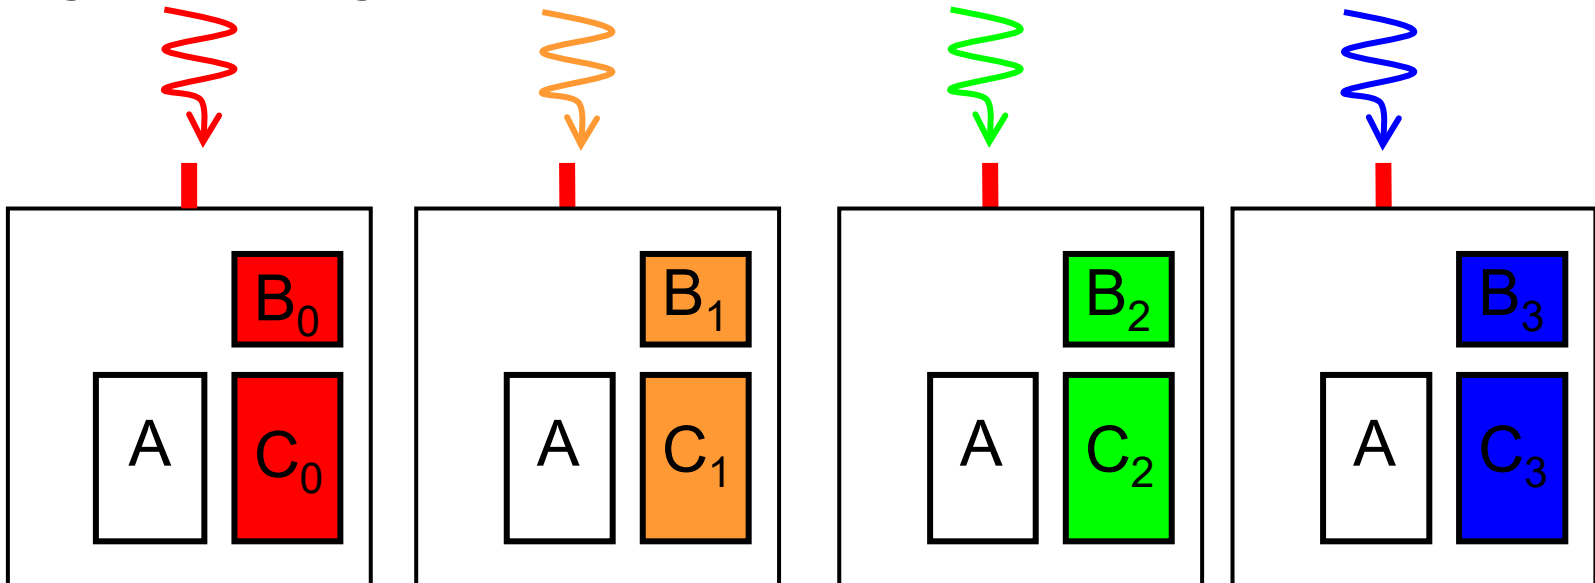
Design distribution method:



I will divide B, C vertically.

I will put replicas of A on every process...

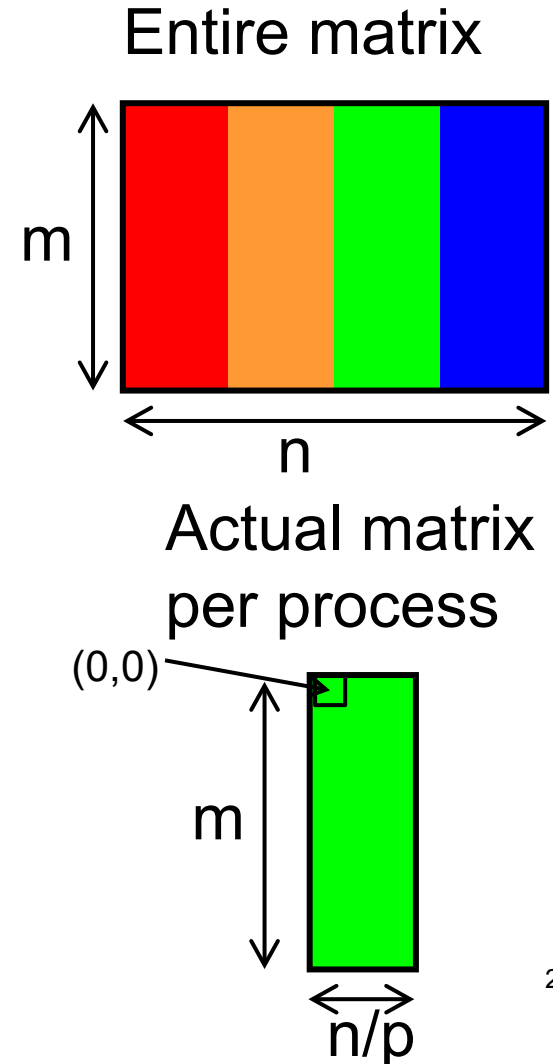
Programming actual location:



Programming Actual Data Distribution



- We want to distribute a $m \times n$ matrix among p processes
 - We assume n is divisible by p
- Each process has a partial matrix of size $m \times (n/p)$
 - We need to “malloc”
 $m \times (n/p) \times \text{sizeof}(\text{data-type})$ size
 - We need to be aware of relation between partial matrix and entire matrix
 - (i, j) element in partial matrix owned by Process $r \Leftrightarrow (i, n/p \times r + j)$ element in entire matrix

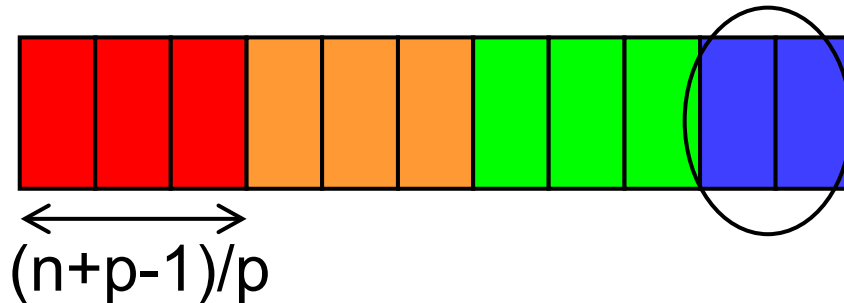




Considering Indivisible Cases

- What if data size n is indivisible by p ?
 - We let $n=11$, $p=4$
 - How many data each process take?
 - $n/p = 2$ is not good (C division uses round down). Instead, we should use round up division
- $(n+p-1)/p = 3$ works well

Note that the “final” process takes less than others



See `divide_length()` function in `mm-mpi/mm.c`
It calculates the range the process should take
→ Outputs are **first index s** and **last index e**



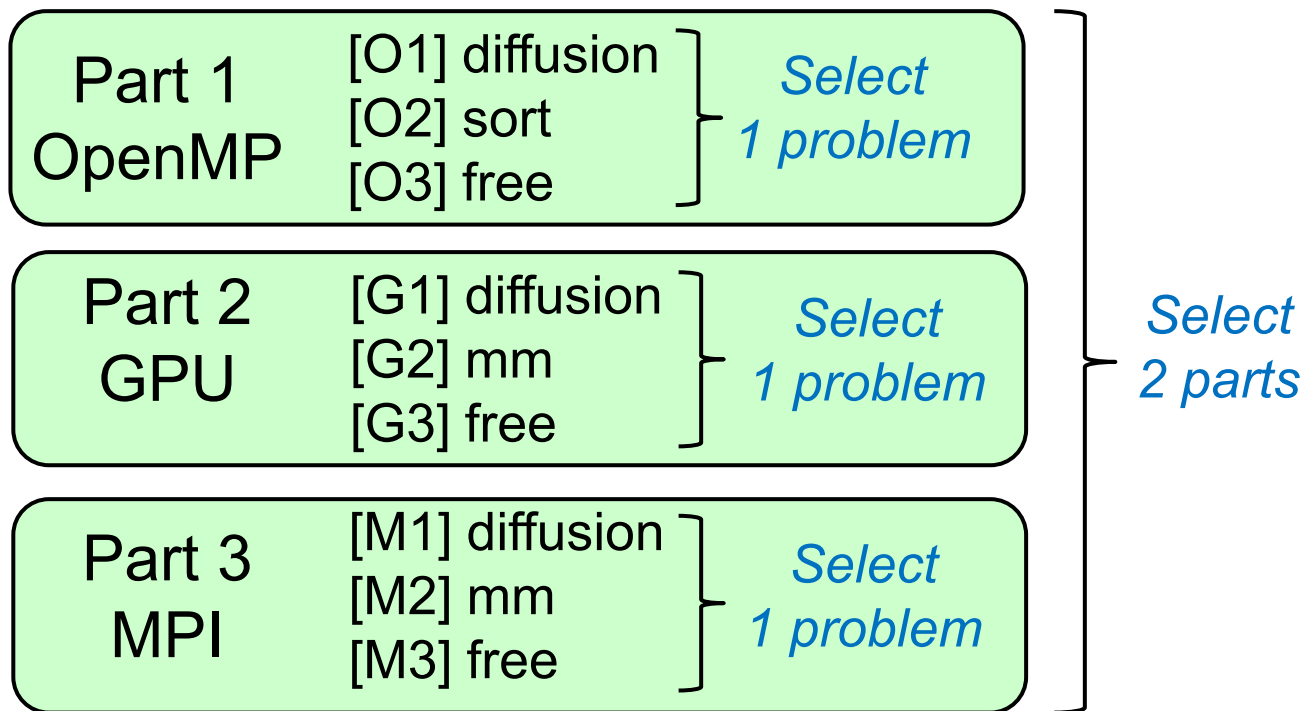
Notes in Time Measurement

- In mm-mpi, gettimeofday() is used for time measurement
- For accurate measurement, we should call **MPI_Barrier(MPI_COMM_WORLD)** before measurement
 - This synchronizes all processes
 - All processes need to call this

Assignments in this Course



- There is homework for each part. Submissions of reports for **2 parts** are required





Assignments in MPI Part (1)

Choose one of [M1]—[M3], and submit a report

Due date: June 12 (Monday)

[M1] Parallelize “diffusion” sample program by MPI.

- Do not forget to change Makefile and job.sh appropriately
- Use deadlock-free communication
 - see `neicomm_safe()` in `neicomm-mpi` sample

Optional:

- To make array sizes (NX, NY) variable parameters
- To consider the case with NY is indivisible by p
 - see `divide_length()` in `mm_mpi` sample
- To improve performance further. Blocking, 2D division, etc



Assignments in MPI Part(2)

[M2] Improve “mm-mpi” sample in order to reduce memory consumption

Optional:

- To consider indivisible cases
- To try advanced algorithms, such as SUMMA
 - the paper “*SUMMA: Scalable Universal Matrix Multiplication Algorithm*” by Van de Geijn
 - <http://www.netlib.org/lapack/lawnspdf/lawn96.pdf>



Assignments in MPI Part (3)

[M3] (Freestyle) Parallelize *any* program by MPI.

- cf) A problem related to your research
- More challenging one for parallelization is better
 - cf) Partial computations have dependency with each other

Notes in Report Submission (1)



- Submit the followings via **T2SCHOLA**
 - (1) **A report document**
 - PDF, MS-Word or text file
 - 2 pages or more
 - in English or Japanese (日本語もok)
 - (2) **Source code files** of your program
 - Try “zip” to submit multiple files

Notes in Report Submission (2)



The report document should include:

- Which problem you have chosen
- How you parallelized
 - It is even better if you mention efforts for high performance or new functions
- Performance evaluation on TSUBAME
 - With varying number of processes
 - On an interactive node,
 - $1 \leq P \leq 7$ (P=number of processes)
 - If you try $P > 7$, use “`mpiexec -n [P] --oversubscribe ...`”
 - To use more CPU cores, you need to do “job submission” (optional)
 - With varying problem sizes
 - Discussion with your findings
 - Other machines than TSUBAME are ok, if available



Next Class

- MPI (2)
 - Basic message passing
 - How to parallelize diffusion sample with MPI
 - Related to [M1]
- Planned schedule
 - May 22: Part 3 (2)
 - May 25: **cancelled (休講)** & Due for Part2 assignment
 - May 29: Part 3 (3)
 - June 1: Part 3 (4)