

Practical Parallel Computing (実践的並列コンピューティング)

Part 3: MPI

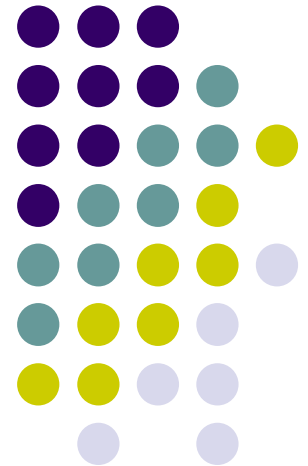
No 3: Communication Costs

May 29, 2023

Toshio Endo

School of Computing & GSIC

endo@is.titech.ac.jp






Overview of This Course

- Part 0: Introduction
 - 2 classes
- Part 1: OpenMP for shared memory programming
 - 4 classes
- Part 2: GPU programming
 - 4 classes ← We are here (1/4)
 - OpenACC (1.5 classes) and CUDA (2.5 classes)
- Part 3: **MPI** for distributed memory programming
 - 4 classes ← We are here (3/4)

“mm” sample: Matrix Multiply




MPI version available at [/gs/hs1/tga-ppcomp/23/mm-mpi/](https://github.com/tga-ppcomp/23/mm-mpi/)

A: a $(m \times k)$ matrix, B: a $(k \times n)$ matrix

C: a $(m \times n)$ matrix

$$C \leftarrow A \times B$$

- Algorithm with a triple for loop
- Supports variable matrix size.
 - Each matrix is expressed as a 1D array by *column-major* format





Execution: `mpiexec -n [#proc] ./mm [m] [n] [k]`

Programming Data Distribution

(for mm-mpi sample)





Design distribution method:



I will divide B, C vertically.

I will put replicas of A on every process...

Programming actual location:



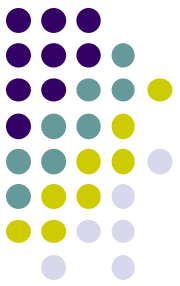
This is not a unique way. Let us discuss other ways

Discussion on Considering Data Distribution




- Choice of data distribution have impact on
 - Communication cost
 - Memory consumption cost
 - In mm-mpi, every process has a copy of matrix $A \rightarrow$ memory consumption is larger ☹️
- Smaller cost is better


Other Data Distribution Methods?




- $C_{i,j}$ requires i -th row of A and j -th column of B



C is divided in col-wise
 \Rightarrow Similarly B
 A is replicated



C is divided in row-wise
 \Rightarrow Similarly A
 B is replicated



C is divided in 2D
 \Rightarrow A:row-wise + replica
 B:col-wise + replica

Total Comm.	0	0	0
Total Mem.	$O(mkp+nk+mn)$	$O(mk+nkp+mn)$	$O(mkp^{1/2}+nkp^{1/2}+mn)$

p : the number of processes

Note: If initial matrix is owned by one process, we need communication before computation


Among them, the third version has lowest memory consumption

Reducing Memory Consumption Further




- Even in the third version, memory consumption is $O(mkp^{1/2} + nkp^{1/2} + mn) > O(mk + nk + mn)$ (theoretical minimum)
- If $p=10000$, we consume 100x larger memory ☹️
→ we cannot solve larger problems on supercomputers
- To reduce memory consumption, we want to **eliminate replica!**
→ But this increases communication costs


Trade-off: a balance achieved between two desirable but incompatible features



Data Distribution of Memory Reduced “mm” (related to [M2])



- Not only B and C, but A is divided among all processes
(In this example, column-wise)




Memory consumption
is smallest

- But computing each C element requires data on other processes → We need **communication** !

How We Proceed Computation with Others' Data



- The following algorithm is not good for memory consumption

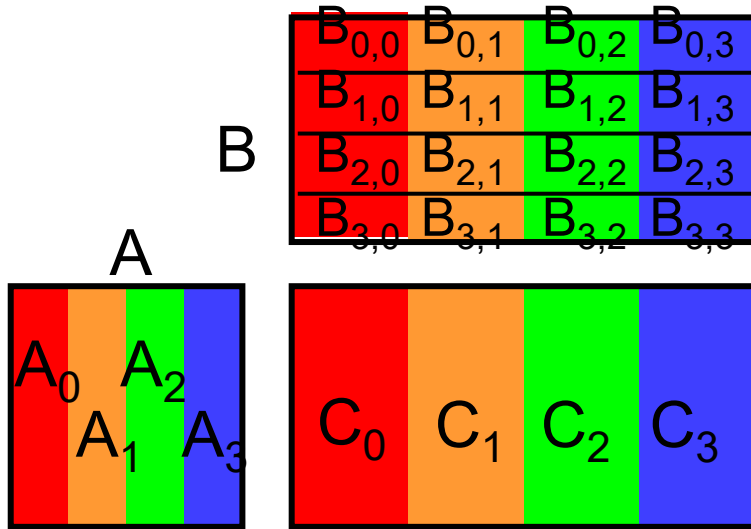


1. Collect entire A from other processes by communication
 2. Compute $C_L = A \times B_L$
- Each process has (entire) A, B_L , C_L → Same as mm-mpi ☹️

We should avoid computation of $C_L = A \times B_L$ at once



Algorithm of Memory Reduced “mm”



If we have A only partially,
we can only do $C_L = A \times B_L$
partially

Algorithm

Step 0:

P_0 sends A_0 to all other processes
Every process P_r computes

$$C_r += A_0 \times B_{0,r}$$

Step 1:

P_1 sends A_1 to all other processes
Every process P_r computes

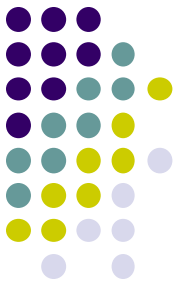
$$C_r += A_1 \times B_{1,r}$$

:

Repeat until Step (p-1)


Total Comm: $O(mkp)$ **Total Mem:** $O(mk+nk+mn)$

Actual Data Distribution



Every process has partial A, B, C

- A_L on process $r \Leftrightarrow A_r$
- B_L on process $r \Leftrightarrow B_r$
- C_L on process $r \Leftrightarrow C_r$



- Additionally, every process should prepare a receive buffer $\rightarrow A'$ in the figure
 - A' (instead of A_L) is used for arguments of `MPI_Recv()`
 - On receivers, A' is used for computation

[Q] What if a process uses A_L for `MPI_Recv()`'s target ?

Programming

Memory Reduced mm



On every **process** r :

```
for (s = 0; s < p; s++) { // s: step no, p: number of processes
```

```
    if ( $r == s$ ) {
```

```
        for (dest = 0; dest < p; dest++)
```

```
            if (dest != r) MPI_Send( $A_L$ , ..., dest, ...);
```

```
    } else
```

```
        MPI_Recv( $A'$ , ..., s, ...);
```

P_s sends its A_L to all other processes

Receives data (P_s 's A_L) and stores it to A'


```
    if ( $r == s$ )
```

```
        Compute  $C_L += A_L \times B_{L,s}$ 
```

```
    else
```

```
        Compute  $C_L += A' \times B_{L,s}$ 
```

```
}
```




Improvements of Memory Reduced Version



Followings are options (NOT mandatory) in assignments [M2]

1. To use SUMMA: scalable universal matrix multiplication algorithm

- See <http://www.netlib.org/lapack/lawnspdf/lawn96.pdf>
- Replica is eliminated, and matrices are divided in 2D



2. To use **collective communications** (explained hereafter)

Peer-to-peer Communications vs Collective Communications



- Communications we have learned are called **peer-to-peer communications**
- A process sends a message. A process receives it



※ `MPI_Irecv`, `MPI_Isend` are also peer-to-peer communications

	Blocking	Non-Blocking
Peer-to-Peer	<code>MPI_Send</code> , <code>MPI_Recv...</code>	<code>MPI_Isend</code> , <code>MPI_Irecv...</code>
Collective	<code>MPI_Bcast</code> , <code>MPI_Reduce...</code>	(<code>MPI_Ibcast</code> , <code>MPI_Ireduce...</code>)

Collective Communications (Group Communications)



- **Collective communications** involves many processes
 - MPI provides several collective communication patterns
 - Bcast, Reduce, Gather, Scatter, Barrier...
 - All processes must call the same communication function



→ Something happens for all of them

One of Collective Communications: Broadcast by MPI_Bcast




cf) rank 0 has “int a” (called **root process**). We want to send it to all other processes

```
MPI_Bcast(&a, 1, MPI_INT, 0, MPI_COMM_WORLD);
```

- All processes (in the communicator) must call MPI_Bcast(), including rank 0


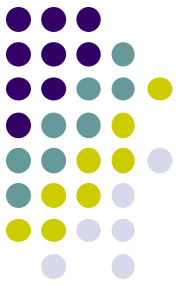
→ All other process will receive the value on memory region **a**



✂ What is the role of 1st argument?

it is “input” on the root process, and “output” on other processes

MPI_Bcast Can Be Used in Memory Reduced MM



- In Step i , rank i becomes the root
 - It sends A_L to all other processes
- This is “broadcast” pattern. **We can use MPI_Bcast!**

Note: Root wants to send A_L . Others want to receive data into A'
→ Different pointers

Solution 1:

```
if (I am rank  $i$ ) copies  $A_L$  to  $A'$   
MPI_Bcast( $A'$ , ... );
```

Solution 2:

```
if (I am rank  $i$ ) {MPI_Bcast( $A_L$ , ...); }  
else {MPI_Bcast( $A'$ , ...); }
```

Reduction by MPI_Reduce




cf) Every process has “int a”. We want the sum of them

```
MPI_Reduce(&a, &b, 1, MPI_INT, MPI_SUM, 0,  
           MPI_COMM_WORLD);
```

operation *root process*

- Every process must call MPI_Reduce()
→ The sum is put on **b** on **root process** (rank 0 now)
- Operation is one of MPI_SUM, MPI_PROD(product), MPI_MAX, MPI_MIN, MPI LAND (logical and), etc.




MPI Version of “pi” Sample

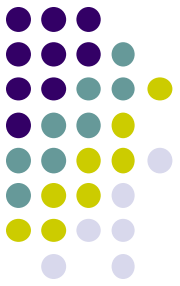


</gs/hs1/tga-ppcomp/23/pi-mpi/>

- Execution: `mpiexec -n [#procs] ./pi [n]`
 - n: Number of division
 - Cf) `./pi 100000000`
- We divide n tasks among processes and calculate total yellow area
 1. Each process calculates local sum
 2. Rank 0 obtains the final sum by **MPI_Reduce**



$$dx = 1/n$$
$$y = \sqrt{1-x^2}$$



Using pi-mpi sample

- [/gs/hs1/tga-ppcomp/23/pi-mpi](#)

[make sure that you are at a interactive node (r7i7nX)]

`module load cuda openmpi` *[Do once after login]*

`cd ~/t3workspace` *[In web-only route]*

`cp -r /gs/hs1/tga-ppcomp/23/pi-mpi .`

`cd pi-mpi`

`make`

[An executable file “pi” is created]


`mpexec -n 7 ./pi 1000000000`

Number of division

Note: Differences with “omp for reduction” in OpenMP



- Syntaxes are completely different
- Computations are also different
 - `#pragma omp for reduction(...)` in OpenMP
 - Do “`sum += a[i]`” in parallel for loop with `reduction(+:sum)`
- `MPI_Reduce(...)` in MPI
 - If each input is an array, output is also an array
 - Operations are done for each index






MPI_Allreduce

- Allreduce = Reduction + Bcast

```
MPI_Allreduce(&a, &b, 1, MPI_INT, MPI_SUM,  
             MPI_COMM_WORLD);
```

- The sum is put on **b** on all processes



Important communication pattern for distributed deep learning → Google “allreduce deep learning”



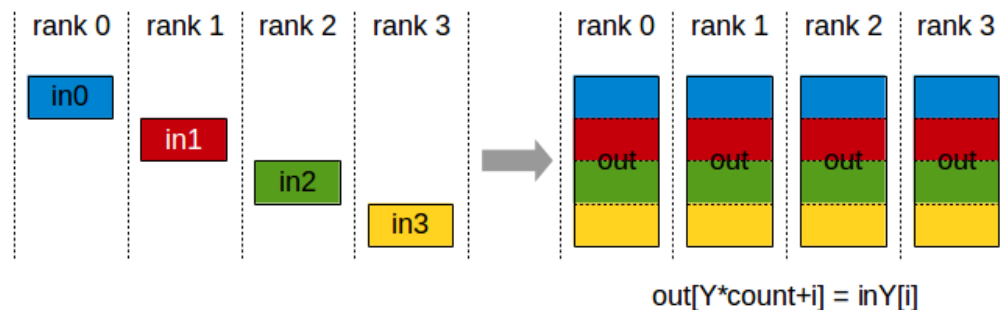
MPI_Barrier

- **Barrier synchronization:** processes are stopped until all processes reach the point
`MPI_Barrier(MPI_COMM_WORLD);`
 - Used in sample programs, to measure execution time more precisely

Other Collective Communications



- **MPI_Scatter**
 - An array on a process is “scattered” to all processes
 - cf) Process 0 has an array of length 10,000. There are 10 processes. The array is divided to parts of length 1,000 and scattered
- **MPI_Gather**
 - Data on all processes are “gathered” to the root process.
 - Contrary to MPI_Scatter
- **MPI_Allgather**
 - Similar to MPI_Gather. Gathered data are put on all processes




*From NCCL manual at
docs.nvidia.com*



Performance of Collective Communication





“Do I Really Need to Learn New Functions?”



- You can still use MPI_Send/MPI_Recv multiple times, but **collective functions are often faster**

On TSUBAME3
1 proc / node

In the graph, rank 0 called MPI_Send for p-1 times to other processes



- MPI_Bcast are faster, especially when p is larger !
- The reason is MPI uses “scalable” communication algorithms:
 - Will be explained in next class

cf) <http://www.mcs.anl.gov/~thakur/papers/mpi-coll.pdf>



FYI: Measurement Method

- Measurement in the previous page was done with </gs/hs1/tga-ppcomp/23/mm-mpi/>
- NOTE: job*.sh in this directory need to consume TSUBAME points
 - job*.sh use > 2nodes


Why are Collective Communications Fast?






- Since MPI library uses scalable communication algorithms

- Case of **broadcast**:


Flat tree algorithm (slow)

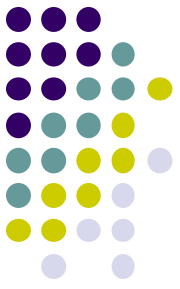


Binomial tree algorithm




Scatter&Allgather
algorithm





Model of Communication Time

Illustration of peer-to-peer communication of data size M



$$T = M / B + L$$

T : Communication time

M : Data size

B : Bandwidth

L : Network latency

※ Be aware of difference between “Byte” and “bit”: 1Byte=8bit


※ Actually it is more complex for process's place, effects of network topology, congestion, packet size...

Cost Model of Broadcast Algorithms



- Case of “broadcast” of size M data
 - p : number of processes, B : network bandwidth, L : network latency


Flat tree algorithm



$$p(M/B + L)$$

→ *Slow*


Binomial tree algorithm



$$(\log p)(M/B + L)$$


※ difference between $p-1$ and p is ignored

Broadcast by Scatter&Allgather Algorithm (1)

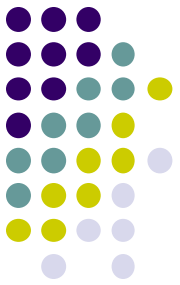


- (1) The root process divide the message into p parts
- (2) Scatter
- (3) Allgather


R. Thakur and W. Gropp. Improving the performance of collective operations in mpich. EuroPVM/MPI conference, 2003.




Scatter&Allgather Algorithm (2)




- (1) The root process divide the message into p parts
- (2) Scatter: i -th part goes to process i
- (3) Allgather



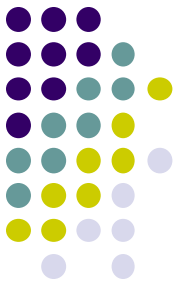
Scatter&Allgather Algorithm (3)




- (1) The root process divide the message into p parts
- (2) Scatter
- (3) Allgather in $\log p$ steps
 - If $p=8$, we use 3 steps



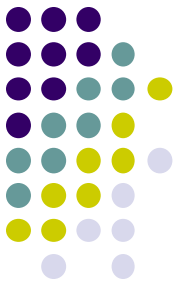
Scatter&Allgather Algorithm (3)




- (1) The root process divide the message into p parts
- (2) Scatter
- (3) Allgather in $\log p$ steps
 - If $p=8$, we use 3 steps



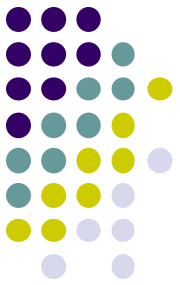
Scatter&Allgather Algorithm (3)




- (1) The root process divide the message into p parts
- (2) Scatter
- (3) Allgather in $\log p$ steps
 - If $p=8$, we use 3 steps



Scatter&Allgather Algorithm (3)



- (1) The root process divide the message into p parts
- (2) Scatter
- (3) Allgather in $\log p$ steps
 - If $p=8$, we use 3 steps




Cost of Scatter&Allgather Algorithm



- Scatter&Allgather algorithm

- (1) The root process divide the message into p parts
- (2) Scatter
- (3) Allgather



Comparison of Broadcast Algorithms



- Consider two extreme cases
 - If M is sufficiently large: $M/B+L \rightarrow M/B$
 - If M is close to zero: $M/B+L \rightarrow L$

	Flat Tree	Binomial Tree	Scatter&Allgather
General Cost	$p(M/B+L)$	$(\log p) (M/B+L)$	$2M/B + (p + \log p)L$
Cost with very large M (L is ignored)	$p M/B$	$(\log p) M/B$	$2 M/B$ \rightarrow Fastest
Cost with very small M (M is ignored)	$p L$	$(\log p) L$ \rightarrow Fastest	$(p + \log p) L$

Many MPI libraries implement multiple algorithms

They switch them automatically according to message size M 😊

Assignments in MPI Part (Abstract)



Choose one of [M1]—[M3], and submit a report
Due date: **June 12 (Monday)**

[M1] Parallelize “diffusion” sample program by MPI.

[M2] Improve mm-mpi sample in order to reduce memory consumption.

[M3] (**Freestyle**) Parallelize *any* program by MPI.

For more detail, please see 3-1 slides on May 19



Next Class

- MPI (4)
 - Other topics about MPI and parallel computing
 - Improvement of MPI version of diffusion further
- Planned schedule
 - June 1: Part 3 (4) & TSUBAME tour
 - Please come to the lecture room as usual