

Practical Parallel Computing (実践的並列コンピューティング)


Part 0: Introduction No 2: Parallel Architecture & Sample Programs

Apr 11, 2024

Toshio Endo

School of Computing & GSIC

endo@is.titech.ac.jp





Overview of This Course


- Part 0: Introduction
 - 2 classes ← We are here (2/2)
- Part 1: **OpenMP** for shared memory programming
 - 4 classes
- Part 2: **GPU** programming
 - OpenACC and CUDA
 - 4 classes
- Part 3: **MPI** for distributed memory programming
 - 3 or 4 classes

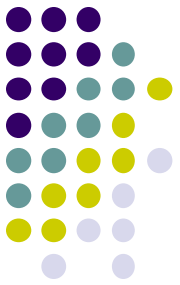
Different Parallel Programming Methods



- Why do we learn several programming methods?
 - OpenMP, OpenACC/CUDA, MPI in this lecture

Reason: Programming methods depend on **structure of computer hardware** (or **computer architecture**) we will use






Software Runs on Hardware

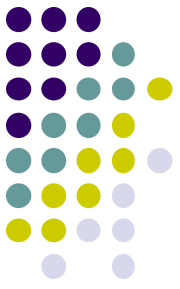
- Software = Algorithm + Data
- Hardware (architecture) \doteq Processor + Memory

Note: This is so simplified discussion



Software Example


```
int a[3] = {10, 20, 30};  
int i;  
  
for (i = 0; i < 3; i++) {  
    a[i] = a[i] * 2;  
}
```




What is Parallel Architecture?

- Parallel architecture has MULTIPLE components
- Two basic types:

Shared memory
parallel architecture



Distributed memory
parallel architecture




- Different programming methods are used for different architecture

Modern SCs use Both!



Modern SCs are combination of “shared” and “distributed”
“shared memory” in a node


“distributed memory” among nodes, connected by network



✂ Moreover, each processor (core) may have *SIMD parallelism*, such as SSE, AVX...

A processor (core) can do several computations at once


SIMD is out of scope of this class






(Confusing) Terminology

- In old days, definition of “**processor**” was simple




- Since around 2005, “**multicore processor**” became popular

A processor package

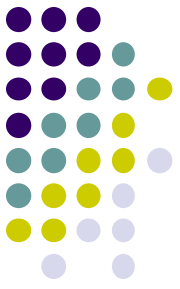


=




A processor core

※ *Hyperthreading* makes discussion more complex:
1 physical core = 2 logical cores
In this slide, “core” basically means physical core

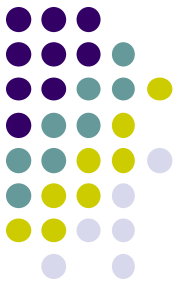


A TSUBAME4 Node (1)

- 2 processor packages (CPU) × 96 cores
→ A TSUBAME4 node has **192 cores**



- GPUs are (still) omitted in this figure




A TSUBAME3 Node (2)

- A node has 2 CPUs + 4 GPUs
 - Each GPU (H100) has 132SMs = 16,896 cores



=



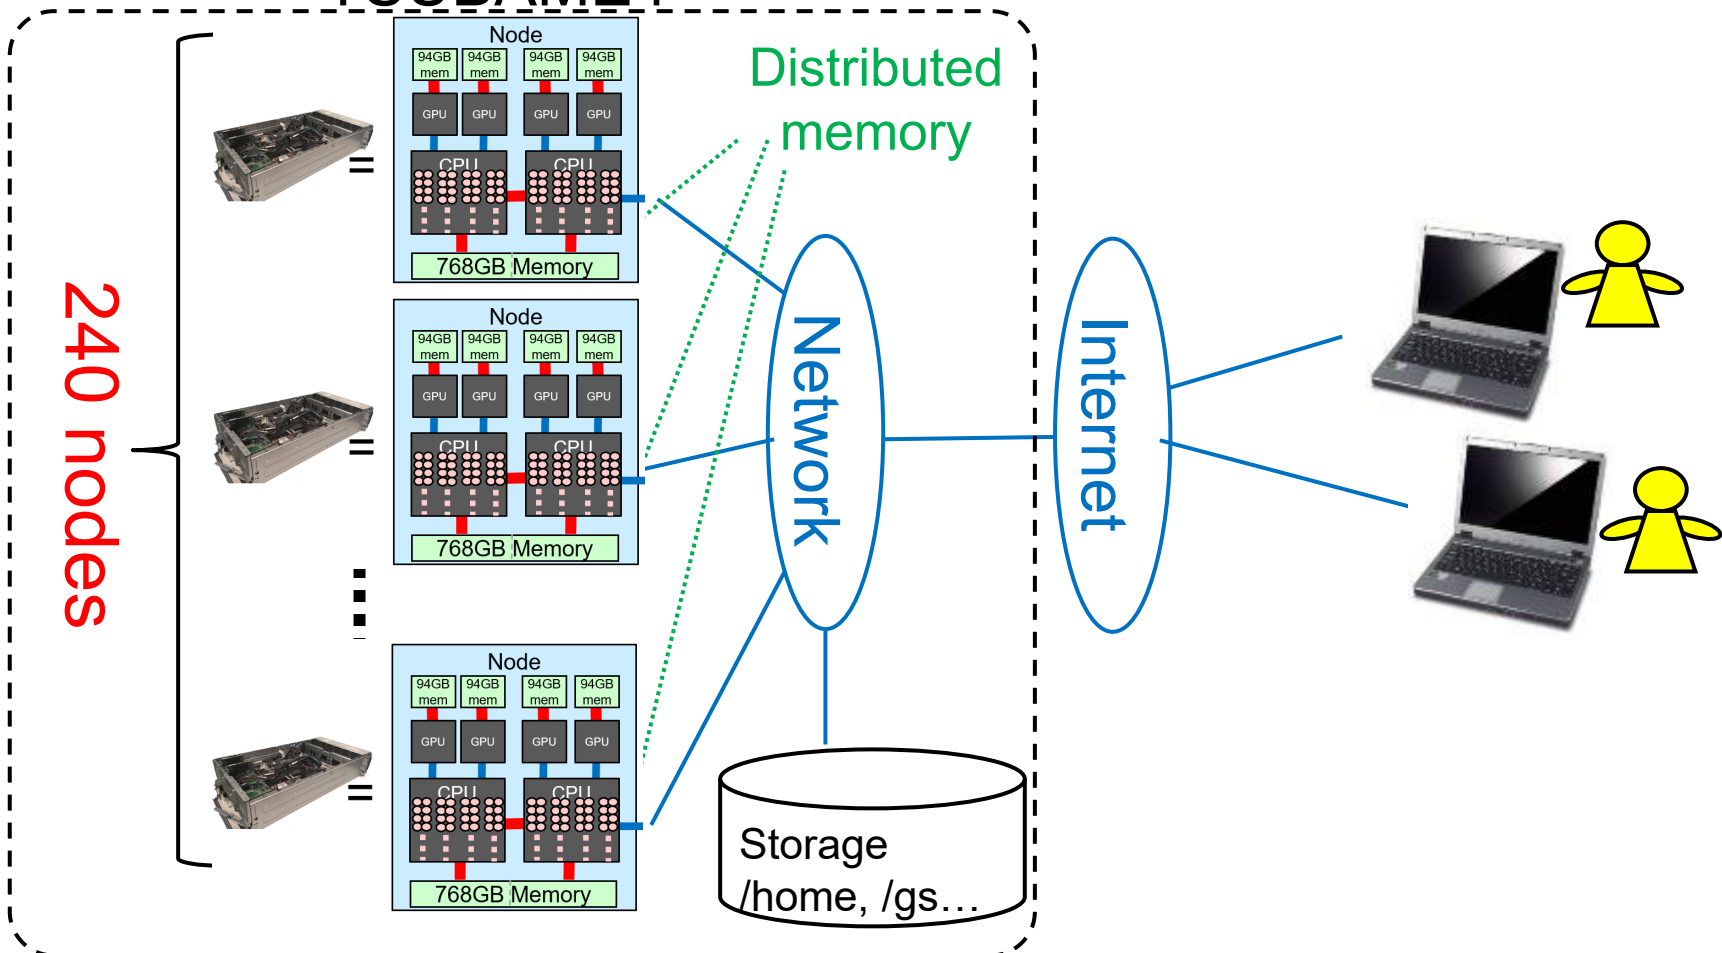
Network:
connected to
other nodes

TSUBAME4 System

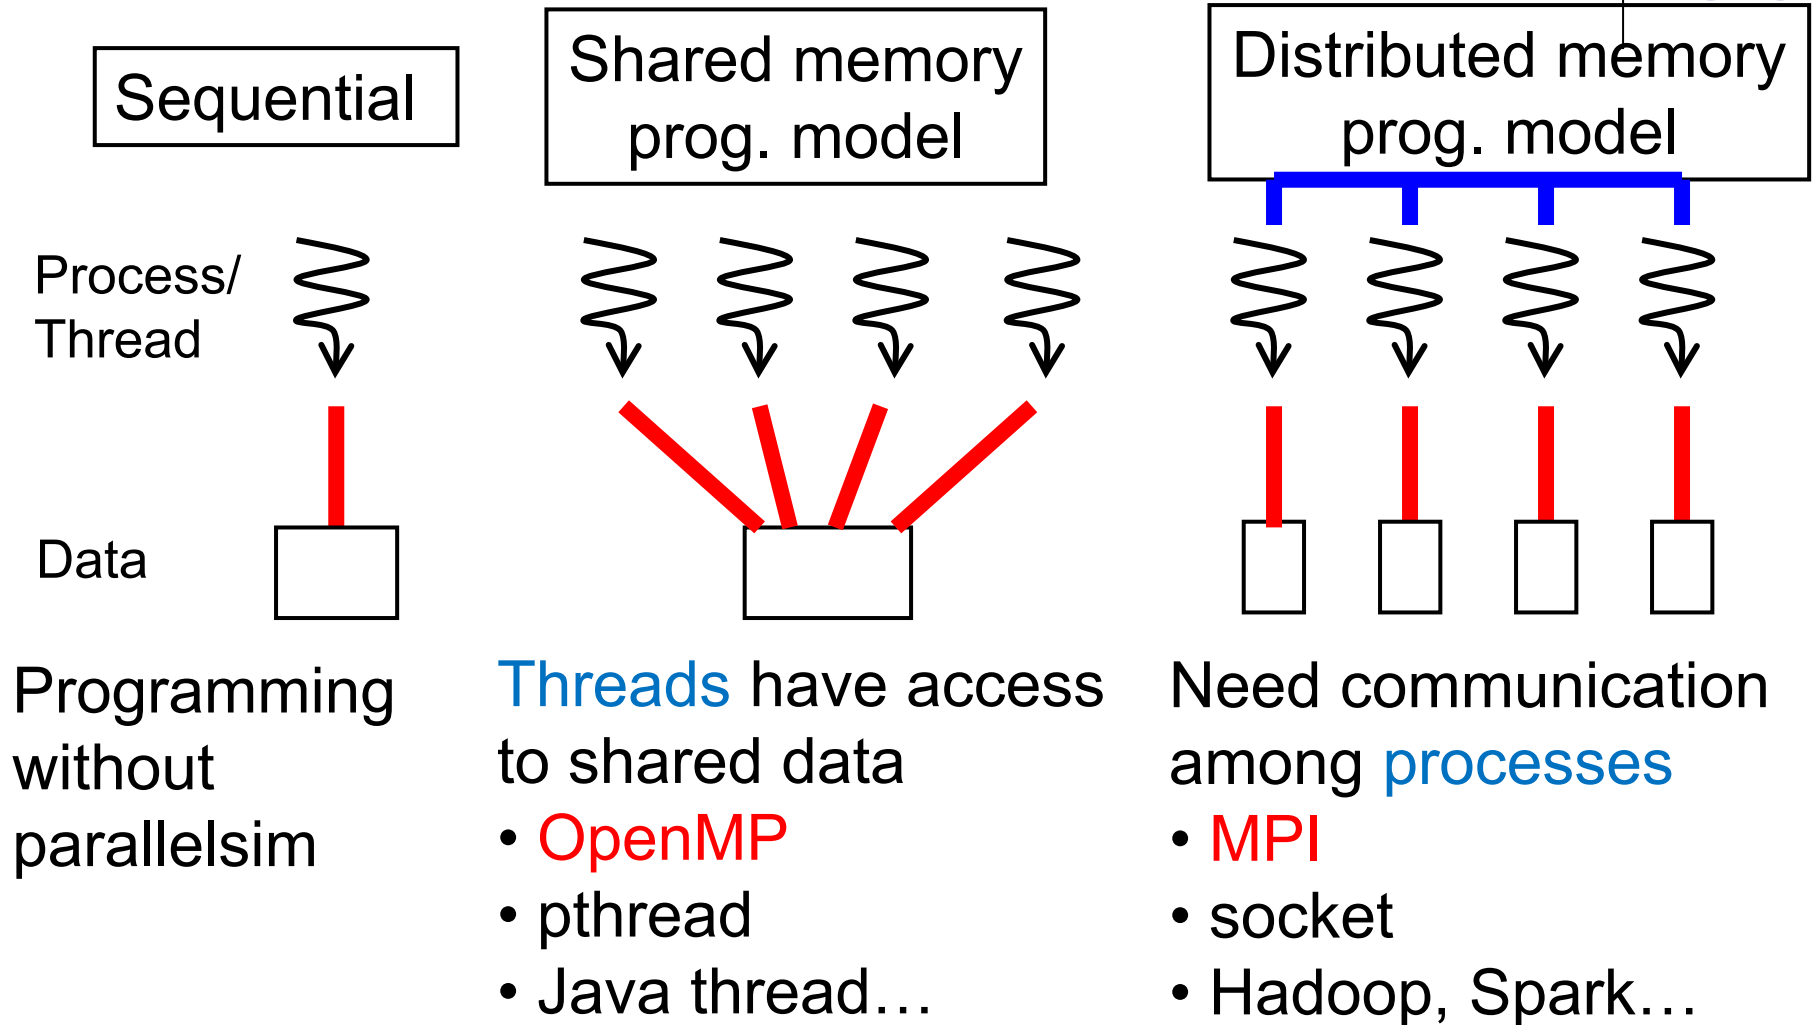



- 240 nodes (and storage) are connected by fast network

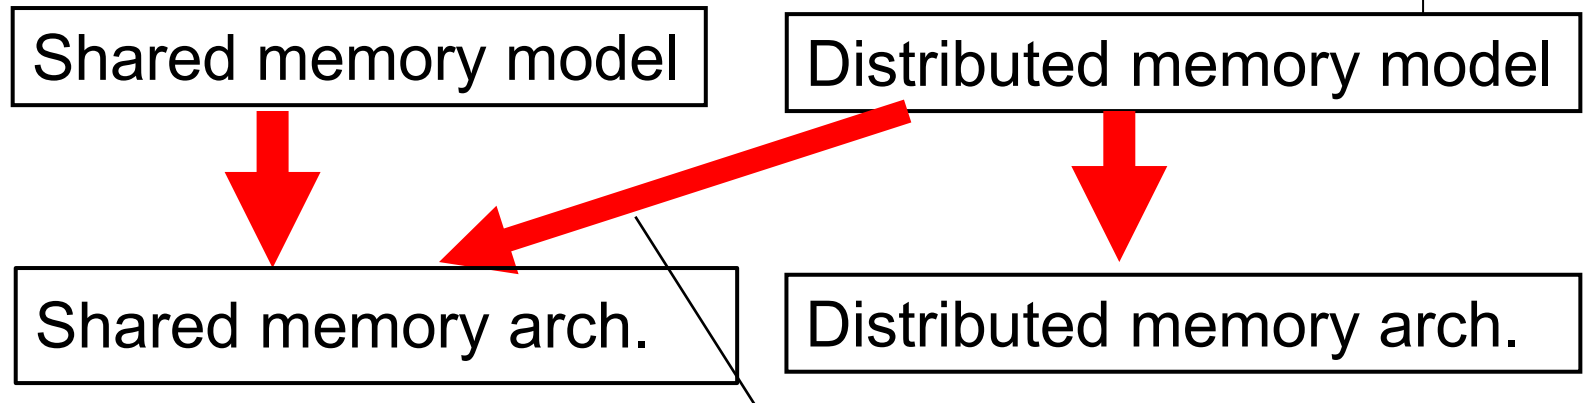
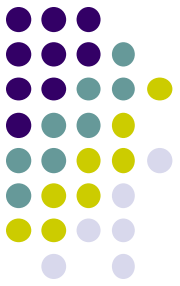
TSUBAME4



Classification of Parallel Programming Models




Programming Models on Architecture



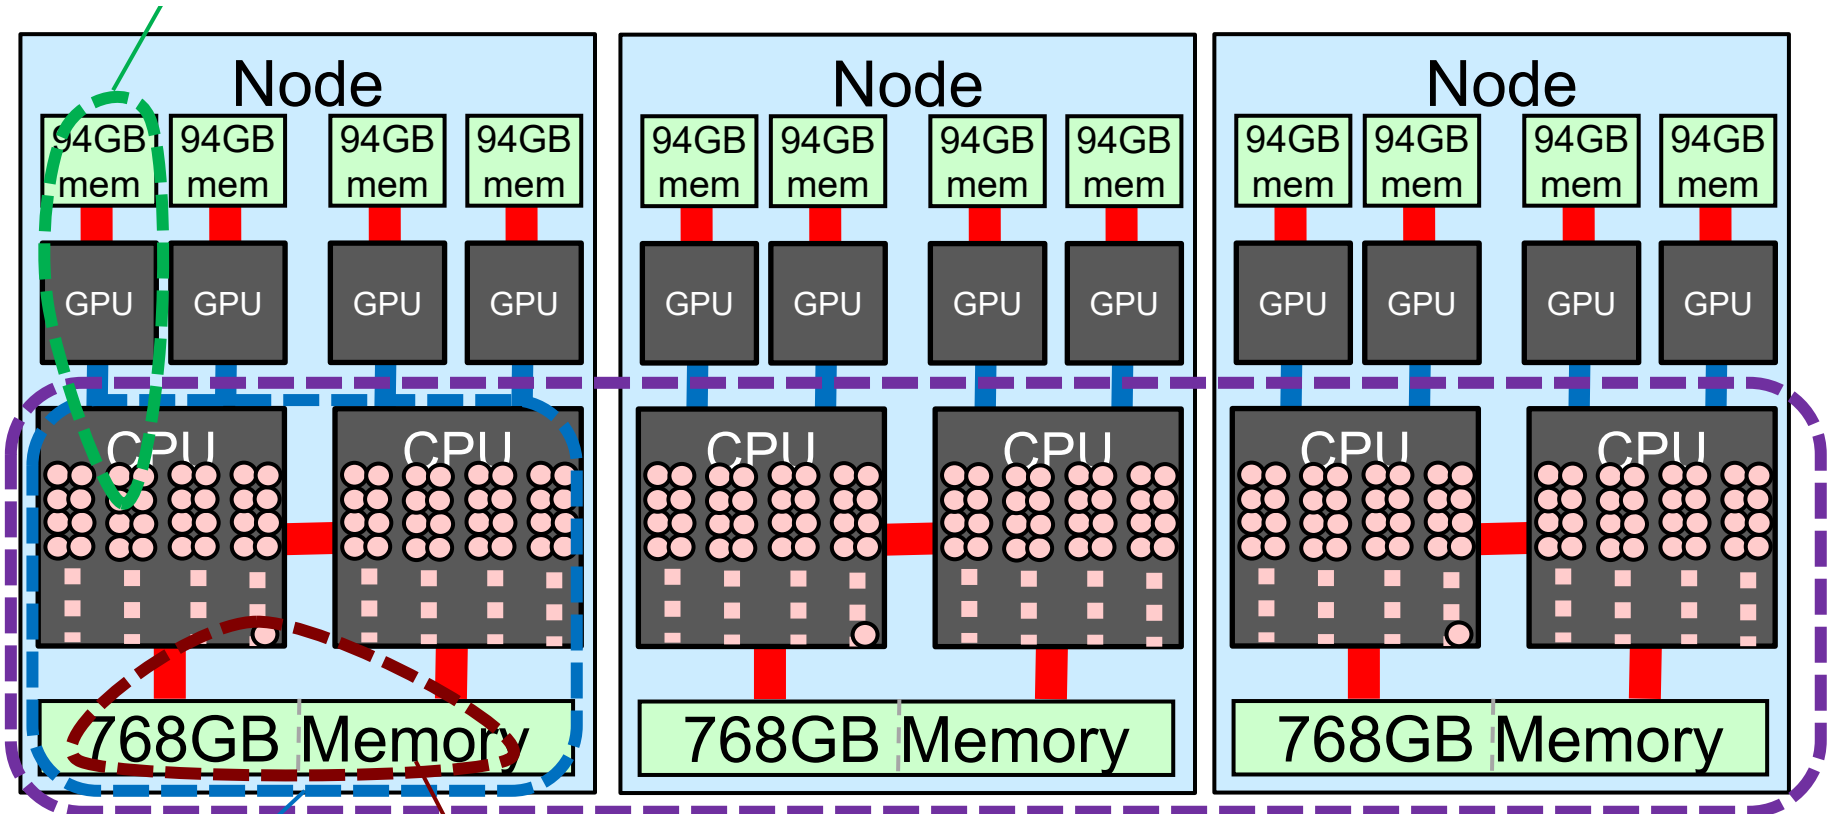
It's OK to make multiple
processes on a node

- Shared memory model (Part 1) can use only cores in a single node (up to 192 cores on TSUBAME4)
- Distributed memory model (Part 3) supports large scale parallelism ($192 \times 240 = 46,080$ cores on TSUBAME4)

Parallel Programming Methods on TSUBAME



OpenACC/CUDA



MPI

OpenMP


Sequential



Standard route

Web-only route

TSUBAME Interactive Node



A node is partitioned into 8, each of which has

- 1/8 node = 24 CPU cores + 96GB memory
+ 0.5 GPU (7680cores+46GB mem)

A user can use one partition

A partition may be shared by several users → you may suffer from slow down

Sample Programs in this Lecture



- Samples are at </gs/bs/tga-ppcomp/24/> directory
 - You have to be a member of [tga-ppcomp](#) group
 - If “[ls /gs/bs/tga-ppcomp/24](#)” works well, you are a member
 - There are sub-directories per sample
- Sequential (non-parallel) sample programs are
 - [mm](#): matrix multiplication
 - [pi](#): approximation of pi (π)
 - [diffusion](#): simple simulation of diffusion phenomena
 - [fib](#): Fibonacci number
 - [sort](#): quick-sort sample

Make Copies of Sample In Case of mm



- Samples in /gs/... are “*read-only*”, so make copies of samples into somewhere in your home directory
- Compile programs inside your home directory

Example:

Create a directory ~/ppc24 (name is arbitrary, do once)

```
mkdir ~/ppc24
```

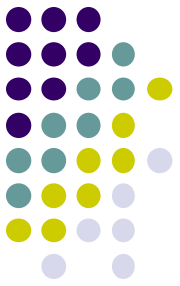
Copy the “**mm**” sample as ~/ppc24/mm

```
cd ~/ppc24
```

```
cp -r /gs/bs/tga-ppcomp/24/mm .
```

don't forget
space & dot

Executing Sample In Case of mm



```
cd ~/ppc24  
[make sure that you have copied mm directory]  
cd mm  
ls  
[you will see 3 files of mm.c, Makefile, job.sh]  
make    [this creates an executable file "mm"]  
./mm 2000 2000 2000  
[this is the execution of mm sample]
```

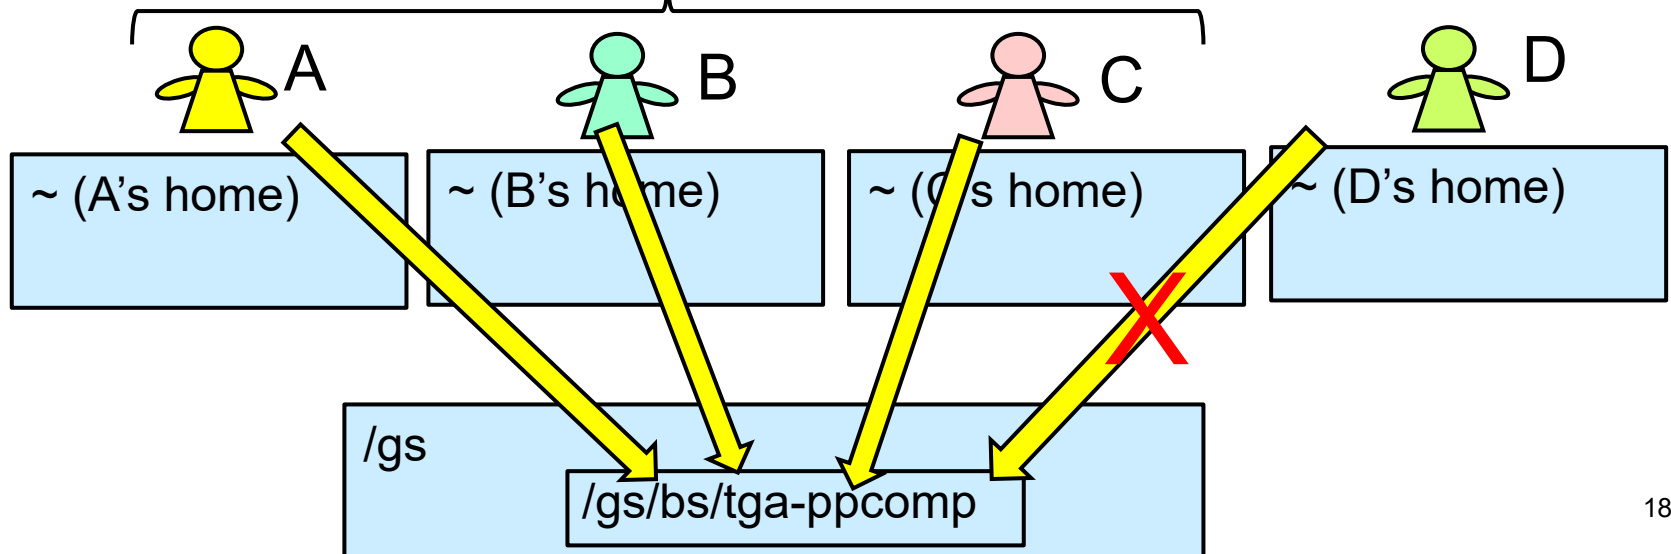
- grey texts are comments; do not type

Disk Storage of TSUBAME



- ~/... (home directory) is assigned to each user
 - Actually, ~ is an alias to /home/??/uX00000 (username)
- /gs/... are shared by multiple users
 - /gs/bs/tga-ppcomp/... is shared by tga-ppcomp members
- Above storage can be accessed from all computing nodes and log-in nodes in TSUBAME

tga-ppcomp group members




Web-only route

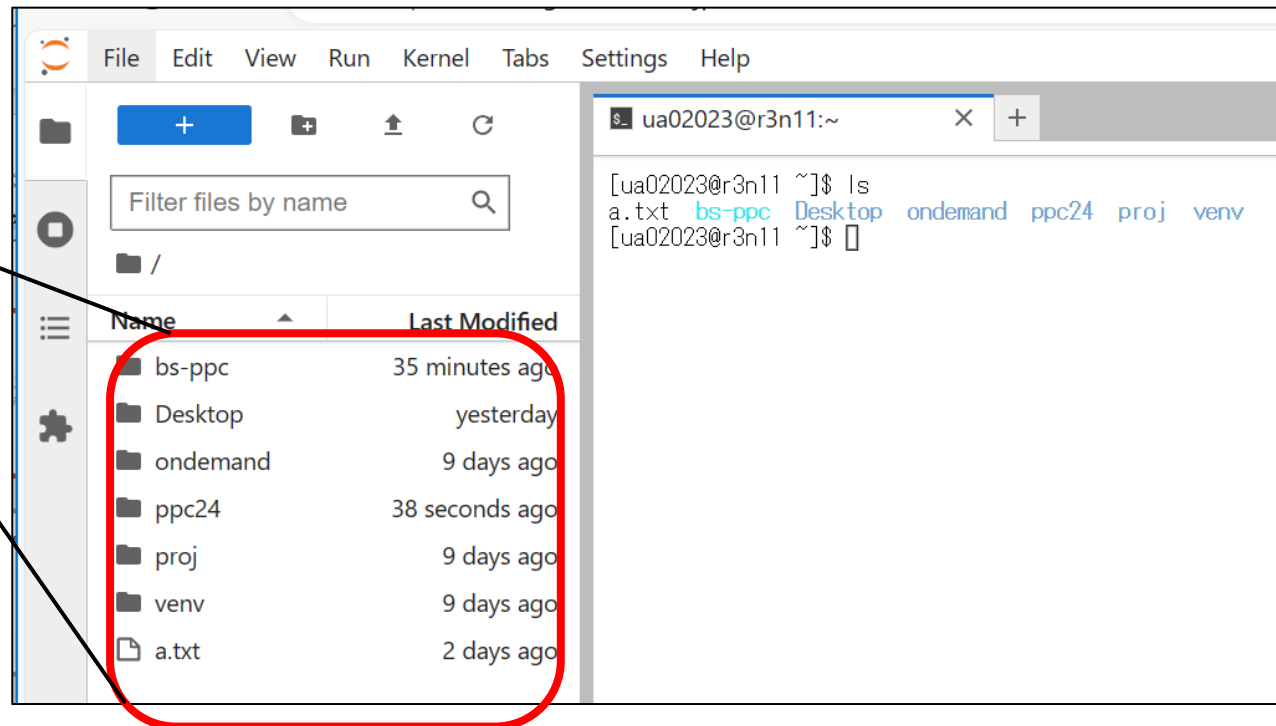
Notes in Web-Only Route

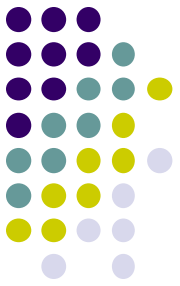


- In Jupyter lab screen, the folder tree shows your home (~/)



~/ (A's home)





Make Copies of Other Samples

Do once for each sample.

“pi” sample

```
cd ~/ppc24  
cp -r /gs/bs/tga-ppcomp/24/pi .
```

“diffusion” sample

```
cd ~/ppc24  
cp -r /gs/bs/tga-ppcomp/24/diffusion .
```

“fib” sample

```
cd ~/ppc24  
cp -r /gs/bs/tga-ppcomp/24/fib .
```

“sort” sample

```
cd ~/ppc24  
cp -r /gs/bs/tga-ppcomp/24/sort .
```

Using Sample Programs (3)

Executing Samples



Before execution, please do cp, cd and make properly for each sample

- mm

```
./mm 2000 2000 2000
```

Options are matrix sizes m, n, k

- pi

```
./pi 1000000000
```

10^8

Option is number of samples n

- diffusion

```
./diffusion 20
```

Option is number of time steps nt

- fib

```
./fib 40
```

Option is sequence index n

- sort

```
./sort 100000000
```

10^7

Option is array length n to be sorted



How Do We Edit C Programs?

There are several ways. The best way is up to you

1. Using editors on Linux

[1a] vim

[1b] emacs

NOTE: emacs is not good on web route, since Ctrl+s does not work well

2. Using editors on your PC

- You need to copy the file into PC, edit on your PC, and copy it to TSUBAME again
 - scp command on your PC, or WinSCP can be used
 - Drag&drop

Web-only route

3. Using Jupyter's editor

Web-only route

“mm” sample: Matrix Multiply



Original version is at </gs/bs/tga-ppcomp/24/mm/>


A: a $(m \times k)$ matrix

B: a $(k \times n)$ matrix


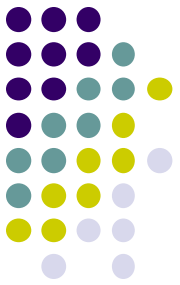
C: a $(m \times n)$ matrix

$C \leftarrow A \ B$

- This sample supports variable matrix sizes
- Execution: `./mm [m] [n] [k]`
 - cf) `./mm 2000 2000 2000`
 - cf) `./mm 1000 3000 1000`



Matrix Multiply Algorithm (1)



- $C_{i,j}$ is defined as the dot product of
- A's i-th row
 - B's j-th column

The algorithm uses triply-nested loop

```
for (i = 0; i < m; i++) {  
  for (j = 0; j < n; j++) {  
    for (l = 0; l < k; l++) {  
       $C_{i,j} += A_{i,l} * B_{l,j};$   
    }  
  }  
}
```

←For each row in C
←For each column in C
←For dot product



Matrix Multiply Algorithm (2)

```
for (i = 0; i < m; i++) {  
  for (j = 0; j < n; j++) {  
    for (l = 0; l < k; l++) {  
      Ci,j += Ai,l * Bl,j;  
    }  
  }  
}
```

← For each row in C
← For each column in C
← For dot product

- The innermost statement is executed for mnk times
- Compute Complexity: $O(mnk)$
 - Computation speed (Flops) is obtained as $2mnk/t$, where t is execution time

The innermost statement includes 2 (floating point) calculations: $*$, $+$

Variable Length Arrays in (Classical) C Language



- `double C[n];` raises an error. How do we do?
- `void *malloc(size_t size);`
⇒ Allocates a memory region of *size* bytes from “heap region”, and returns its head pointer
- When it becomes unnecessary, it should be discarded with `free()` function

A fixed length array

```
double C[5];  
  
... C[i] can be used ...
```

A variable length array

```
double *C;  
C = (double *)malloc(sizeof(double)*n);  
  
... C[i] can be used ...  
  
free(C);
```

array length

How We Do for Multiple Dimensional Arrays



`double C[m][n];` raises an error. How do we do?

Not in a straightforward way. Instead, we do either of:

(1) Use a pointer of pointers

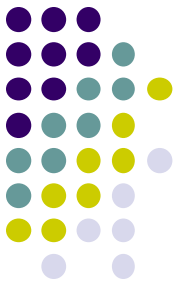
- We *malloc* m 1D arrays for every row (each has n length)
- We *malloc* 1D array of m length to store the above pointers


(2) Use a 1D array with length of $m \times n$

(*mm sample uses this method*)

- To access an array element, we should use `C[i*n+j]` or `C[i+j*m]`, instead of `C[i][j]`

Express a 2D array using a 1D array




“I want
to use ...”

a 2D array $C[m][n]$

m

8	3	7	4	1	2
0	2	1	5	0	3
1	8	6	4	2	1
3	4	8	1	0	2

n

$C[1][3]$

Expressions in C language (Example)

```
double *C; C = malloc(sizeof(double)*m*n);
```

n

8	3	7	4	1	2	0	2	1	5	0	3	8	1	0	2
---	---	---	---	---	---	---	---	---	---	---	---	-------	---	---	---	---

$C[1*n+3]$


In this case, an element $C_{i,j}$ is $C[i*n+j]$

Two Data Formats

Row major format

- More natural for C programmers


$$C_{i,j} \Rightarrow C[i*n+j]$$



Column major format

- BLAS library
- mm sample uses this

$$C_{i,j} \Rightarrow C[i+j*m]$$



- We have more choices for 3D, 4D... arrays

[Q] Does the format affect the execution speed?



Actual Codes in mm Sample

```
for (i = 0; i < m; i++) {  
    for (j = 0; j < n; j++) {  
        for (l = 0; l < k; l++) {  
            Ci,j += Ai,l * Bl,j;  
        } } }  
}
```

IJL order



```
for (j = 0; j < n; j++) {  
    for (l = 0; l < k; l++) {  
        double blj = B[l+j*k];  
        for (i = 0; i < m; i++) {  
            double ail = A[i+l*m];  
            C[i+j*m] += ail*blj;  
        } } }  
}
```

Change (2):
JLI order is used
(a bit faster)

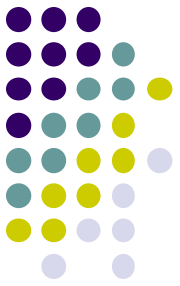
Change (1):
Matrix elements as
1D array elements

Time Measurement in Samples



- `gettimeofday()` function is used
 - It provides wall-clock time, not CPU time
 - Time resolution is better than `clock()`
 - Note: newer `clock_gettime()` is more recommended

```
#include <stdio.h>
#include <sys/time.h>
:
{
    struct timeval st, et;
    long us;
    gettimeofday(&st, NULL); /* Starting time */
    ...Part for measurement...
    gettimeofday(&et, NULL); /* Finishing time */
    us = (et.tv_sec-st.tv_sec)*1000000+
        (et.tv_usec-st.tv_usec);
    /* us is difference between st & et in microseconds */
}
```



If You Have Not Done This Yet

Please do the followings as soon as possible

- Please make your account on TSUBAME
- Please send an e-mail to ppcomp@el.gsic.titech.ac.jp

Subject: [TSUBAME3 ppcomp account](#)

To: ppcomp@el.gsic.titech.ac.jp

Department name:

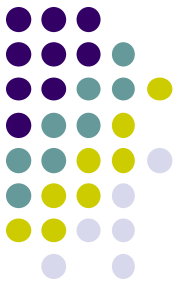
School year:

Name:

Your TSUBAME account name:

Then we will invite you to the TSUBAME group, [please click URL and accept the invitation](#)

その後、TSUBAMEグループへの招待を送ります。[メール中のURLをクリックして参加承諾してください](#)



If You Have Not Done This Yet

Please do 1&2 as soon as possible. We accept the mail later.

1と2をできるだけ早期に行ってください。それより後でも受け付けます

1. (If you are new to TSUBAME) please make your account on TSUBAME
(まだ作ったことがなければ) TSUBAMEアカウントを作成してください
2. Please send an e-mail after account creation
作成してから下記のようなe-mailを送ってください

To: ppcomp@el.gsic.titech.ac.jp

Subject: **TSUBAME4 ppcomp account**

Department name:

Student ID:

Name:

TSUBAME account name:

If You Have Not Done This Yet (cont)



3. You will receive an invitation e-mail to [tga-ppcomp](#) TSUBAME group. Please read it and accept the invitation.

[tga-ppcomp](#) TSUBAMEグループへの招待e-mailが届くはずです。指示に従って招待を受けてください。

Next Class: Introduction to OpenMP



- Shared memory parallel programming API
- Extensions to C/C++, Fortran
- Includes directives & library functions
 - Directives: `#pragma omp ~`

```
int i;  
#pragma omp parallel for  
for (i = 0; i < 100; i++) {  
    a[i] = b[i] + c[i];  
}
```